

menler

Your turning point in the AI era.

AI APTITUDE QUESTION BANK

AI for Students

Aware Level

About This Question Bank

An intermediate-level assessment for students who have some experience with AI tools. Tests how AI models work, capability limits, critical thinking, practical tool judgment, and AI ethics at a deeper level than the Beginner bank.

5

Sets

15

Questions per Set

75

Total Questions

SET GUIDE

Set 1 · How AI Models Actually Work

How AI models actually work — architecture, training, context windows, and outputs

Set 2 · AI Capabilities and Their Real Limits

AI capabilities and real limits — what AI does well and where it reliably fails

Set 3 · Thinking Critically About AI

Thinking critically about AI — evaluating claims, bias, benchmarks, and outputs

Set 4 · AI Tools in Practice — Beyond the Basics

AI tools in practice — applied use, learning, building, and professional judgment

Set 5 · AI Ethics and Society — The Bigger Picture

AI ethics and society — deepfakes, bias, privacy, policy, and societal impact

Each question has one correct answer. Select the best answer from the options provided.

Correct answers and explanations follow each question.

SET

1

How AI Models Actually Work

Set 1 of 5 · 15 Questions · AI for Students — Aware

Q1. A language model generates each word by:

- A. Searching a knowledge base for the most relevant phrase.
- B. Sampling from a probability distribution over its vocabulary.**
- C. Applying logical inference rules to the input question.
- D. Retrieving the closest matching sentence from training data.

Correct Answer: B**Explanation**

Language models compute a probability distribution over all possible next tokens and sample from it. There is no database retrieval, logical inference engine, or sentence matching — just statistical prediction, token by token.

Q2. What does "temperature" control in a language model output?

- A. The maximum number of tokens the model can generate.
- B. The speed at which the model processes your input.
- C. The confidence threshold for including a fact in the response.
- D. The randomness or creativity in the sampled token distribution.**

Correct Answer: D**Explanation**

Temperature scales the probability distribution before sampling. High temperature flattens the distribution (more random, creative outputs); low temperature sharpens it (more predictable, conservative outputs). It is not about speed, length, or confidence.

Q3. Why can a language model produce different answers to the identical prompt in separate sessions?

- A. Randomness in the sampling step produces varied outputs.**
- B. The model updates its weights between each conversation.
- C. The model searches different databases each session.
- D. Previous conversations influence the model's knowledge state.

Correct Answer: A

Explanation

Each generation samples from a probability distribution with randomness ($temperature > 0$). This means the same prompt can yield legitimately different responses. The model does not learn between sessions or access different databases.

Q4. What is a "context window" in a language model?

- A. The time period within which the model can access recent news.
- B. The set of topics the model has been trained to respond to.
- C. The maximum total tokens the model can process at one time.**
- D. The number of follow-up questions allowed in one session.

Correct Answer: C

Explanation

The context window is the maximum amount of text (measured in tokens) the model can take in and process in a single call — including your input, system instructions, and conversation history. It is a technical capacity limit, not a topic or time constraint.

Q5. Which statement about transformer-based AI models is accurate?

- A. They process text sequentially, word by word, like humans read.
- B. They use attention mechanisms to weigh relationships between tokens.**
- C. They store factual knowledge in a searchable lookup table.
- D. They reason from first principles to arrive at conclusions.

Correct Answer: B

Explanation

Transformers use self-attention mechanisms to compute weighted relationships between all tokens in the input simultaneously. They do not process text sequentially, use lookup tables, or reason from first principles.

Q6. Why does Claude sometimes refuse to answer certain questions?

- A. Its training includes guidelines that shape which requests it declines.**
- B. It cannot process questions that exceed its topic training range.
- C. It detects harmful intent using a separate threat classifier.
- D. Its context window is too small for sensitive or complex questions.

Correct Answer: A

Explanation

Claude's behaviour including refusals is shaped by its training — specifically by guidelines and values baked in during the training process. It is not a topic filter, a separate classifier, or a context window limitation.

Q7. What is "fine-tuning" a language model?

- A. Adjusting the model's temperature for better output quality.
- B. Filtering the model's training data to remove harmful content.
- C. Reducing the model's size so it runs faster on a device.
- D. Training an existing model further on a specific dataset or task.**

Correct Answer: D

Explanation

Fine-tuning takes a pre-trained base model and continues training it on a new, domain-specific dataset. This adapts the model's behaviour for specific tasks or domains without starting from scratch.

Q8. A model trained primarily on English text is asked a question in Bengali. The most likely outcome is:

- A. It will refuse the question since it was trained only in English.
- B. It will translate the question to English and respond in English.
- C. It may respond but with lower accuracy and fluency in Bengali.**
- D. It will generate the same quality response as it does in English.

Correct Answer: C

Explanation

Models trained predominantly on English data can often handle other languages but with noticeably lower accuracy, fluency, and cultural nuance — especially for lower-resource languages like Bengali. Performance degrades with training data imbalance.

Q9. What makes "hallucination" a structural property of language models rather than a fixable bug?

- A. Models predict probable text, not verified facts, by design.**
- B. Hallucination is caused by corrupted data in model storage.
- C. It only occurs when the model has insufficient training data.
- D. It can be fully eliminated by using a larger model size.

Correct Answer: A

Explanation

Language models are optimised to generate fluent, probable text — not to retrieve verified facts. This means plausible-sounding incorrect output is an inherent outcome of the architecture, not a bug that can be patched without changing how the model fundamentally works.

Q10. What happens when you reach the end of a model's context window in a long conversation?

- A. The model automatically summarises the conversation and continues.
- B. The model refuses to generate any further responses.
- C. The model switches to a compressed memory mode.
- D. The model can no longer access the earliest parts of the conversation.**

Correct Answer: D

Explanation

When the context window is full, older content "falls off" and becomes invisible to the model. It cannot access information from before the cutoff in that conversation. It does not automatically summarise, refuse, or switch modes.

Q11. Why do larger language models generally perform better on complex tasks?

- A. Larger models are trained on more recent data by necessity.
- B. More parameters allow more complex pattern representations.**
- C. Larger models have access to faster external knowledge bases.
- D. More parameters increase the model's logical reasoning capacity.

Correct Answer: B

Explanation

Larger models have more parameters — more numerical weights — which allows them to represent and generalise from more complex patterns in training data. This is not the same as logical reasoning; it is richer statistical pattern representation.

Q12. What is the key difference between a base model and an instruction-tuned model?

- A. Base models are larger and more capable than instruction-tuned ones.
- B. Instruction-tuned models have access to real-time internet data.
- C. Instruction-tuned models are trained to follow conversational instructions.**
- D. Base models are designed for professional use; tuned models for consumers.

Correct Answer: C

Explanation

A base model is trained on raw text and predicts next tokens without any particular conversational format. Instruction-tuned models are fine-tuned on human instruction-response pairs so they follow instructions and behave more like useful assistants.

Q13. What does RLHF (Reinforcement Learning from Human Feedback) primarily do?

- A. Trains models to retrieve and cite human-written sources accurately.
- B. Reduces hallucination by checking outputs against a fact database.
- C. Speeds up model training by using human-labelled data more efficiently.
- D. Trains models to produce outputs humans rate as more helpful and safe.**

Correct Answer: D

Explanation

RLHF uses human preference ratings to train a reward model, then uses that reward model to fine-tune the language model to produce outputs humans rate as more helpful, harmless, and honest. It does not directly eliminate hallucination or add retrieval.

Q14. Why does providing a system prompt before user input improve Claude's behaviour?

- A. It increases the model's context window for the conversation.
- B. It sets persistent context and behavioural guidelines for the session.**
- C. It reduces the model's response time by pre-loading intent.
- D. It unlocks more accurate model outputs by activating expert mode.

Correct Answer: B

Explanation

A system prompt provides context (persona, task, constraints, tone) that shapes Claude's behaviour throughout the conversation. It is not a technical unlock — it is persistent instructional context that the model attends to when generating each response.

Q15. What is the most accurate description of what a language model "knows"?

- A. Statistical patterns from training text that approximate knowledge.**
- B. A structured knowledge graph of facts stored during training.
- C. Verified encyclopaedic facts curated before deployment.
- D. Real-time information retrieved from internet sources.

Correct Answer: A

Explanation

Language models do not store discrete facts — they learn statistical patterns across vast training text. This "knowledge" is a weighted approximation, not a verified knowledge graph. This is why outputs can seem knowledgeable while containing errors.

SET

2

AI Capabilities and Their Real Limits

Set 2 of 5 · 15 Questions · AI for Students — Aware

Q16. You ask Claude to write a poem, and you ask it to solve a differential equation. Which task is Claude more structurally suited to?

- A. The equation — mathematical reasoning is more reliable.
- B. Both equally — Claude handles all text tasks the same way.
- C. The poem — language generation is its core competency.**
- D. The equation — it can access mathematical libraries.

Correct Answer: C

Explanation

Language generation — including creative writing — is directly what language models are designed to do. Mathematical reasoning requires multi-step logical accuracy, which language models handle less reliably without specialised tools.

Q17. A classmate says "Claude with web search is always accurate because it checks the internet." What is wrong with this?

- A. Claude still generates its synthesis of search results, which can contain errors.**
- B. Nothing is wrong — web search eliminates hallucination entirely.
- C. Web search makes Claude slower but not more accurate overall.
- D. Claude cannot actually browse the internet during web search.

Correct Answer: A

Explanation

Web search retrieves real pages, but Claude still synthesises, interprets, and generates text from those results — and can still make errors in that synthesis. Web search reduces but does not eliminate hallucination risk.

Q18. Which of these tasks is Claude most likely to perform consistently well?

- A. Predicting next week's stock price movements from news.
- B. Diagnosing a rare medical condition from a symptom list.
- C. Translating a complex legal contract to another language perfectly.
- D. Reformatting and restructuring a well-defined piece of text.**

Correct Answer: D

Explanation

Reformatting and restructuring text — changing structure, style, or format within clear constraints — is a language generation task Claude handles reliably. Prediction, rare medical diagnosis, and legal translation all carry high-risk failure modes.

Q19. Why is Claude unreliable for tasks requiring precise numerical calculation?

- A. It was not trained on mathematical data.
- B. It generates number tokens statistically, not arithmetically.**
- C. It does not have access to a calculation engine.
- D. Numerical tasks exceed its context window capacity.

Correct Answer: B

Explanation

Claude predicts tokens — including digits — using statistical patterns, not actual arithmetic operations. It can arrive at wrong answers that look like correct calculations. For precise calculation, a code interpreter or calculator is essential.

Q20. An AI tool claims "zero hallucination." The most critical response is:

- A. Trust it — large companies would not make false capability claims.
- B. Accept it for factual tasks but remain cautious for creative ones.
- C. Ask for evidence — no current language model can guarantee this.**
- D. Verify it by testing the tool on creative writing prompts only.

Correct Answer: C

Explanation

No current large language model architecture fully eliminates hallucination — it is a structural property of text prediction. Any claim of "zero hallucination" should be treated with scepticism and demands rigorous, independent evidence.

Q21. What does Claude do well that a basic search engine does not?

- A. It retrieves the most recent and accurate factual information.
- B. It provides source citations for every claim it makes.
- C. It verifies information before including it in a response.
- D. It synthesises, explains, and reasons across multiple pieces of information.**

Correct Answer: D

Explanation

Search engines retrieve and rank existing pages. Claude synthesises, explains, contextualises, and reasons across information — producing integrated understanding rather than a list of links. But it does not retrieve, cite, or verify by default.

Q22. You give Claude a 10-page document and ask it to find every mention of a specific term. What is the key reliability concern?

- A. Claude may miss instances depending on how it attends to the text.**
- B. Claude cannot process documents longer than 2 pages at a time.
- C. Claude will only search the first and last paragraphs.
- D. Claude requires the document to be in plain text format.

Correct Answer: A

Explanation

Language models use attention mechanisms that may not reliably find every instance of a term across a very long document — especially if the context window is near full. For exhaustive text search, programmatic tools are more reliable.

Q23. What is the most significant limitation of AI for creative writing assistance?

- A. It cannot generate text in a first-person narrative voice.
- B. It produces statistically average outputs that can lack originality.**
- C. Creative writing tasks exceed its context window.
- D. It only generates text in genres heavily represented online.

Correct Answer: B

Explanation

AI generates text that is statistically probable — which tends toward average, familiar outputs that mirror patterns in training data. The result can be technically competent but lacking in the unexpected, original combinations that distinguish excellent creative work.

Q24. A student asks Claude to "read" an uploaded PDF. What is actually happening?

- A. Claude visually reads and interprets each page like a human.
- B. Claude downloads the file and searches it like a browser.
- C. Claude accesses the PDF's metadata and key words only.
- D. Claude processes the extracted text content from the document.**

Correct Answer: D

Explanation

Claude does not "see" a PDF the way a human reads it. When PDF text is extracted and passed to Claude, it processes text content — it cannot read formatting, images within the PDF, or scanned text that has not been converted by OCR.

Q25. Why might Claude perform worse on a highly specialised technical topic than a generalist one?

- A. Claude applies more conservative limits on technical responses.
- B. Technical topics exceed the model's logical processing capability.
- C. Specialised topics have less training data relative to general ones.**
- D. Claude routes technical queries to a reduced-capability sub-model.

Correct Answer: C

Explanation

Language model performance correlates with training data volume for a topic. Highly specialised domains — rare diseases, niche legal frameworks, emerging technologies — have less representation in training data, producing weaker outputs.

Q26. What is the most appropriate use of AI in writing a formal research report?

- A. Generating the complete report and submitting it as written.
- B. Drafting sections, which you revise, verify, and rewrite in your own words.**
- C. Generating citations and sources, which you compile into a bibliography.
- D. Producing the final polished draft after outlining the structure yourself.

Correct Answer: B

Explanation

AI is valuable for generating drafts to react to — not for producing final output. Revising, verifying, and rewriting in your own words develops the analytical skills the exercise is designed to build, and ensures factual accuracy.

Q27. An AI writing tool rates your essay "8/10 for argument strength." How reliable is this score?

- A. It reflects pattern-matching against training examples, not true evaluation.**
- B. It is highly reliable — AI evaluates argument quality objectively.
- C. It is reliable for structure but not for individual word choice.
- D. It is accurate only if the tool was trained on academic essays.

Correct Answer: A

Explanation

AI scoring reflects whether your essay matches patterns associated with high-rated essays in training data — not a genuine logical evaluation of argument strength. Use AI feedback as a signal, not a grade.

Q28. Which scenario most clearly requires a human expert rather than AI?

- A. Drafting the initial structure of a business proposal.
- B. Explaining what a specific legal clause generally means.
- C. Diagnosing and prescribing treatment for a patient's symptoms.**
- D. Summarising a long meeting transcript into key action points.

Correct Answer: C

Explanation

Medical diagnosis and prescription require licensed clinical judgment, liability, legal accountability, and real-time patient assessment. The other tasks are reasonable AI use cases where the output is reviewed and refined by a competent human.

Q29. Why does Claude sometimes decline to answer questions it could technically answer?

- A. The question exceeded the factual scope of its training data.
- B. Safety classifiers running in parallel override its responses.
- C. Declining reduces server load for high-demand queries.
- D. Its training includes values that lead it to prioritise safety over capability.**

Correct Answer: D

Explanation

Claude's refusals reflect values and guidelines built into its training — not external classifiers, topic filters, or load management. The model has internalised principles about which responses are appropriate to generate.

Q30. The most accurate statement about AI and creativity is:

- A. AI is fully creative because it generates content that did not exist before.
- B. AI can produce novel combinations but within statistical boundaries.**
- C. AI cannot be creative at all since it only recombines existing text.
- D. AI creativity is identical to human creativity in practice.

Correct Answer: B

Explanation

AI can generate novel combinations of patterns — which produces outputs that did not exist before — but this recombination is statistically constrained by training data. It is neither fully uncreative nor fully creative in the human sense.

SET

3

Thinking Critically About AI

Set 3 of 5 · 15 Questions · AI for Students — Aware

Q31. You read a study that says "AI outperforms humans at X task." The most important follow-up question is:

- A. "Under what specific conditions and metrics was this evaluated?"**
- B. "Which company funded the research?"
- C. "How long did the AI take compared to the human?"
- D. "Was the same AI used for all human comparison groups?"

Correct Answer: A

Explanation

AI benchmark claims depend critically on the specific task definition, evaluation metric, test conditions, and dataset used. A performance claim without these details tells you almost nothing about real-world capability.

Q32. A chatbot passes the Turing Test with a specific evaluator. This proves:

- A. It is indistinguishable from a human across all contexts.
- B. It possesses general intelligence equivalent to a human.
- C. It can produce human-like text in that particular interaction.**
- D. It has genuine understanding of the questions it answered.

Correct Answer: C

Explanation

Passing the Turing Test in a specific interaction shows the model can produce text that fooled that evaluator in that scenario — nothing more. It does not demonstrate general intelligence, universal human-likeness, or genuine understanding.

Q33. An AI company claims their model scores 90% on a medical knowledge exam. A critical reader asks:

- A. "Why was only 90% achieved and not 100%?"
- B. "Does this performance transfer to real clinical decision-making?"**
- C. "Which competing AI models scored lower on the same exam?"
- D. "Is the exam available for students to practise on?"

Correct Answer: B

Explanation

Exam performance does not automatically transfer to real-world clinical settings, which involve ambiguity, incomplete information, urgency, and consequences. Benchmark performance and deployed performance are not the same thing.

Q34. You see a viral post claiming "AI wrote this and it's perfect — AI is now better than humans at writing." The most accurate assessment is:

- A. "If it looks perfect, it is perfect — quality is objective."
- B. "This claim is certainly false — AI cannot match human writing."
- C. "This claim is certainly true — AI now outperforms humans at text tasks."
- D. "Better at writing" requires defining what dimension of quality is being measured.**

Correct Answer: D

Explanation

Writing quality is multidimensional: technical correctness, originality, insight, emotional resonance, cultural relevance. "Better than humans" requires specifying which dimension, under which conditions, evaluated by whom.

Q35. Your AI-generated essay received positive feedback from a peer. The safest conclusion is:

- A. The essay was effective for this reader in this context.**
- B. The essay is high quality and would perform well in any setting.
- C. AI-generated essays are generally better than human-written ones.
- D. You should use AI to write all future essays based on this result.

Correct Answer: A

Explanation

One positive reader in one context is a very narrow signal. The essay may have limitations — depth, originality, accuracy — that this reader did not notice. A single data point does not generalise to "AI essays are better."

Q36. A student argues that AI cannot be biased because it uses maths. The most precise counter is:

- A. AI maths is subjective and can be coded with intentional bias.
- B. The mathematical patterns were learned from biased human-generated data.**
- C. All mathematical systems contain inherent human biases.
- D. Maths-based AI is only unbiased when trained on balanced datasets.

Correct Answer: B

Explanation

The mathematics of a language model is a neutral optimisation process — but what it is optimised on (training data) reflects human-generated patterns, including historical inequalities and imbalances. Unbiased maths on biased data produces biased outcomes.

Q37. You use AI to research a historical event and it gives a detailed, confident account. The responsible next step is:

- A. Accept it — history is well-documented in AI training data.
- B. Add a disclaimer that the account may contain errors.
- C. Ask the AI to list its sources to verify the account.
- D. Cross-reference with established historical sources before using it.**

Correct Answer: D

Explanation

AI accounts of historical events can contain compressed, distorted, or outright wrong information presented with confidence. Cross-referencing with established historical sources is the correct academic standard — not disclaimers or asking AI for its own sources.

Q38. What is "prompt injection" and why does it matter for students using AI tools?

- A. A technique for making AI prompts shorter and more efficient.
- B. The process of adding context to improve AI response quality.
- C. Malicious instructions embedded in content that hijacks an AI's behaviour.**
- D. A method of injecting real-time data into an AI's training set.

Correct Answer: C

Explanation

Prompt injection is when malicious text embedded in content a student feeds to an AI — a webpage, document, or email — contains hidden instructions that manipulate the AI's behaviour. Students should be aware that content they process with AI may contain such instructions.

Q39. An AI assistant confidently tells you that a specific scientific paper was published in Nature. Before citing it, you should:

- A. Trust it — major journal names are well-known in AI training.
- B. Search for the actual paper in a verified academic database.**
- C. Ask the AI for the DOI number to use as your citation.
- D. Add "(via AI)" after the citation to flag the AI source.

Correct Answer: B

Explanation

AI can fabricate entire papers — author names, titles, journals, years, even abstracts — all of which sound completely real. Never cite a source from AI without finding and reading the actual paper in a verified database.

Q40. Why might AI perform better when you tell it what role to play (e.g., "act as a senior editor")?

- A. It activates patterns from training text written from that perspective.**
- B. It unlocks a dedicated expert module for that professional role.
- C. It increases the model's reasoning capacity for the task.
- D. It bypasses safety guidelines for more expert-level responses.

Correct Answer: A

Explanation

Role prompting activates patterns associated with how that type of expert writes and thinks in training data. There is no dedicated expert module or reasoning boost — it is a prompting technique that leverages how training data is distributed.

Q41. What is the most reliable method for verifying whether an AI output is factually accurate?

- A. Running the same prompt through two different AI tools.
- B. Asking the AI to rate its own confidence level.
- C. Tracing every specific claim to a citable primary source.**
- D. Checking whether the output sounds logical and well-written.

Correct Answer: C

Explanation

Only tracing claims to citable primary sources can actually verify factual accuracy. Cross-checking with another AI uses the same flawed approach. AI self-rated confidence is unreliable — it sounds equally confident whether right or wrong.

Q42. An AI tells you a company raised ₹50 crore in funding last month. What should alert you most?

- A. The large funding amount — AI exaggerates financial figures.
- B. The use of Indian currency — AI is less accurate on Indian data.
- C. The specific number — AI always invents precise financial figures.
- D. The very recent timeframe — AI may not have this data.**

Correct Answer: D

Explanation

"Last month" is the key signal — AI without web search has a knowledge cutoff and cannot reliably report events from the past few months. The specific figure and India context are also worth verifying, but recency is the primary red flag.

Q43. What is the most accurate way to describe AI "understanding" of language?

- A. True semantic comprehension equivalent to human reading.
- B. Sophisticated pattern matching that resembles but differs from human understanding.**
- C. Surface-level word matching with no semantic capability.
- D. Full grammatical understanding with no real-world grounding.

Correct Answer: B

Explanation

AI processes language by computing statistical relationships between patterns — which can produce outputs that resemble understanding. It is not the same as human semantic comprehension (which includes embodied experience and genuine conceptual grounding) but it is more than surface word-matching.

Q44. A student uses an AI to "check if my argument is logically valid." What is the limitation?

- A. AI cannot evaluate the logic of multi-step arguments at all.
- B. AI logical checks are only reliable for arguments under 100 words.
- C. AI assesses plausibility from training patterns, not formal logical validity.**
- D. AI applies legal rather than philosophical logic standards.

Correct Answer: C

Explanation

AI evaluates whether an argument seems like a good argument based on training patterns — it does not perform formal logical validity analysis. An argument that sounds reasonable may pass its check even if it contains a structural logical fallacy.

Q45. What is the core skill that separates students who use AI effectively from those who do not?

- A. Critical judgment in evaluating and directing AI outputs.**
- B. Having access to the most powerful AI tools available.
- C. Being able to write longer, more detailed prompts.
- D. Understanding the technical architecture of the AI model.

Correct Answer: A

Explanation

The core differentiator is critical judgment — the ability to evaluate AI outputs accurately, identify their weaknesses, verify their claims, and direct them toward what is actually needed. Tool access, prompt length, and technical knowledge all matter less.

SET

4

AI Tools in Practice — Beyond the Basics

Set 4 of 5 · 15 Questions · AI for Students — Aware

Q46. Which professional task is Claude most reliably valuable for, right now?

- A. Producing a final, publication-ready report on a complex topic.
- B. Providing legally binding guidance on a regulatory question.
- C. Replacing a qualified editor for a manuscript before submission.
- D. Generating a first draft that a human then revises and verifies.**

Correct Answer: D**Explanation**

AI consistently adds value as a first-draft generator — accelerating the writing process while leaving verification, revision, and final judgment to a human. The other options involve reliance on AI without appropriate oversight.

Q47. You use Claude to generate code that will handle customer data. The first thing you must do after is:

- A. Run it directly in production to see if it performs correctly.
- B. Review the code for logic errors and security vulnerabilities.**
- C. Ask Claude to review its own code to identify issues.
- D. Have another AI tool validate the code before using it.

Correct Answer: B**Explanation**

AI-generated code can contain subtle bugs, logic errors, and security vulnerabilities. Before deploying any code that handles real data, a qualified human developer must review and test it — not another AI tool.

Q48. What is the most effective way to use Claude for learning a difficult concept?

- A. Ask Claude to write a comprehensive summary you can memorise.
- B. Ask Claude for the three most important facts about the concept.
- C. Ask Claude to explain it, then ask follow-up questions on what you didn't understand.**
- D. Ask Claude to generate a test on the concept to take immediately.

Correct Answer: C

Explanation

Active engagement — asking for explanation, then following up on confusion points, then asking for examples — drives genuine understanding. Passive summarisation and immediate testing without first understanding produce shallow learning.

Q49. You need to summarise 50 research papers on a topic. Where does AI most help?

- A. Processing and synthesising patterns across many documents quickly.**
- B. Providing citable academic summaries of each paper.
- C. Replacing the need to read any of the original papers.
- D. Generating an annotated bibliography with verified citations.

Correct Answer: A

Explanation

AI's most genuine advantage in research synthesis is speed of pattern-processing across many documents. But it cannot provide citable summaries (it may misrepresent papers), cannot replace reading originals for critical work, and cannot verify citations.

Q50. A student builds a Claude-powered study assistant for their class. The biggest risk to address first is:

- A. Claude will not understand questions from younger students.
- B. Claude will generate responses that are too long for students.
- C. Claude may respond in English even when asked in other languages.
- D. Claude may confidently give students incorrect information.**

Correct Answer: D

Explanation

For any educational AI deployment, accuracy is the primary risk — students trusting incorrect information that arrives confidently from an AI assistant. This risk must be addressed through human oversight, verification guidance, and appropriate disclaimers.

Q51. What does "RAG" (Retrieval-Augmented Generation) solve in language model applications?

- A. It increases the speed of model responses in production.
- B. It eliminates hallucination by blocking uncertain outputs.
- C. It supplements model knowledge with retrieved external documents.**
- D. It reduces the model size so it can run on local devices.

Correct Answer: C

Explanation

RAG retrieves relevant documents from an external source (a company's database, a website, etc.) and feeds them into the model's context. This allows the model to answer questions about information not in its training data. It reduces but does not eliminate hallucination.

Q52. You are building a product that uses Claude to answer customer questions about your company. The most important safeguard is:

- A. A human review process for responses on sensitive topics.**
- B. Limiting Claude to responses shorter than 100 words.
- C. Using the most expensive Claude model for accuracy.
- D. Asking Claude to only answer questions it is confident about.

Correct Answer: A

Explanation

Human review of AI responses in high-stakes customer-facing contexts is the most important safeguard. Model size and response length do not guarantee accuracy, and asking AI to self-assess confidence is unreliable.

Q53. When is it appropriate to cite Claude as a source in academic work?

- A. When the information it produced cannot be found elsewhere.
- B. When citing the tool itself as the method, not the facts it produced.**
- C. When you used Claude to generate the research question.
- D. Citing Claude is always appropriate for any academic purpose.

Correct Answer: B

Explanation

You cite Claude as a methodological source — acknowledging you used it in your process — not as a factual source for specific claims. The actual facts Claude produced must be verified and cited from their original primary sources.

Q54. A student uses an AI coding assistant that auto-completes their code. The most important habit to maintain is:

- A. Running all suggested code in a separate testing environment.
- B. Only accepting suggestions for syntax corrections, not logic.
- C. Reading and understanding every line of code before accepting it.**
- D. Asking the AI to explain every suggestion before accepting it.

Correct Answer: C

Explanation

AI coding suggestions can be subtly wrong, insecure, or incompatible with your project. Reading and understanding every line before accepting it is the critical habit — not just running it and seeing if it works.

Q55. What is the most honest statement about AI image generation for school projects?

- A. AI images are indistinguishable from professional photography.
- B. AI image generation is always appropriate for school projects.
- C. AI-generated images are copyright-free and safe to use anywhere.
- D. It can produce visuals, but original thinking must come from you.**

Correct Answer: D

Explanation

AI image generation is a powerful creative tool, but using it requires that the concept, direction, and application come from you. AI-generated images may also carry copyright complexities depending on the tool and its terms of service.

Q56. You want to use Claude to practise for a job interview. The most effective approach is:

- A. Ask Claude to conduct a mock interview, then give honest critique.**
- B. Ask Claude for a list of likely questions to memorise.
- C. Ask Claude to write model answers to each question.
- D. Ask Claude to predict which questions your interviewer will ask.

Correct Answer: A

Explanation

Mock interview practice — where Claude plays the interviewer, you respond, and Claude critiques — builds the skill of formulating and delivering answers. Memorising answers or having AI write them does not develop the actual performance skill required.

Q57. An AI tool automatically generates your class schedule. You notice it has scheduled two classes at the same time. This illustrates:

- A. AI scheduling tools should never be used without human input.
- B. AI can make logical errors that humans can easily spot and correct.**
- C. The AI was not trained on scheduling data.
- D. The error occurred because the context window was exceeded.

Correct Answer: B

Explanation

AI can make logical and consistency errors that are obvious to a human reviewing the output. This is precisely why AI outputs require human review — not to reject AI assistance, but to catch the errors AI is prone to making.

Q58. You want Claude to help you prepare for a science debate. Which use adds the most value?

- A. Writing your opening statement so you can memorise it.
- B. Predicting what the opposing team is likely to argue.
- C. Producing a bibliography of sources that support your position.
- D. Generating strong counter-arguments to your position for you to address.**

Correct Answer: D

Explanation

Generating the strongest possible opposing arguments forces you to think through and strengthen your own position. Memorising AI-written content and asking AI to predict opponents do not develop the reasoning and rebuttal skills debate actually requires.

Q59. What is the key difference between AI being a "tool" and AI being an "authority"?

- A. A tool is directed by your judgment; an authority replaces your judgment.**
- B. Tools require technical skill to use; authority does not.
- C. Tools are more accurate; authority is more convenient.
- D. AI is both a tool and an authority depending on the task.

Correct Answer: A

Explanation

The critical distinction is who exercises judgment. When AI is a tool, you remain the decision-maker and use AI to execute. When AI is treated as an authority, you outsource judgment to it — which is inappropriate given its limitations.

Q60. A student says: "I always ask AI to improve my writing, and it always makes it worse by sounding generic." The root cause is:

- A. AI tools are not designed to improve personal writing voice.
- B. Claude's writing style defaults to a formal academic register.
- C. Generic improvement prompts produce generic stylistic output.**
- D. The student's original writing is too informal for AI to process.

Correct Answer: C

Explanation

"Improve my writing" gives AI no constraints on direction. It defaults to statistical averages — formal, hedged, structurally sound but characterless. Effective prompts for writing improvement specify what dimension to improve and what to preserve.

SET

5

AI Ethics and Society — The Bigger Picture

Set 5 of 5 · 15 Questions · AI for Students — Aware

Q61. What is the most significant societal risk of deepfake video technology?

- A. Fabricated video of real people saying things they never said.**
- B. Reducing the quality of professional video production.
- C. Making film and television content easier to produce.
- D. Replacing the need for actors in commercial advertisements.

Correct Answer: A**Explanation**

The most significant risk is the ability to create convincing video of real people — public figures, private individuals — saying or doing things they never did. This weaponises credibility and makes visual evidence fundamentally unreliable.

Q62. In India, AI-driven credit scoring impacts loan access for millions of people. The primary ethical concern is:

- A. AI credit scores are slower to compute than human assessment.
- B. Customers cannot check their AI-generated credit score.
- C. AI credit scoring ignores repayment history as an input.
- D. Embedded bias may systematically disadvantage certain communities.**

Correct Answer: D**Explanation**

If an AI credit system is trained on historical data that reflects systemic disadvantage — by caste, geography, or gender — it may perpetuate and scale those disadvantages. Bias in credit scoring directly affects economic access and opportunity.

Q63. Who is ultimately responsible when an AI system causes harm to a user?

- A. The AI system itself since it made the autonomous decision.
- B. The organisation that deployed the AI in that context.**
- C. The user who chose to trust and use the AI system.
- D. The AI researchers who built the underlying model.

Correct Answer: B

Explanation

Responsibility for AI harm lies with the organisation that deployed the system — they chose the tool, defined its scope, established oversight, and accepted accountability for its outcomes. AI systems cannot bear moral or legal responsibility.

Q64. A student argues that AI should be banned from schools to prevent cheating. The most nuanced response is:

- A. They are right — AI in schools causes more harm than benefit.
- B. AI detection software is sufficient without needing any policy.
- C. Teaching responsible AI use better prepares students for the future.**
- D. Schools should only allow AI tools created for education.

Correct Answer: C

Explanation

Blanket bans remove opportunities to develop AI literacy that is increasingly essential for professional life. Teaching responsible, transparent AI use alongside traditional skills is more future-relevant than prohibition.

Q65. What is the most honest statement about AI and privacy?

- A. Data shared with AI tools may be used in ways users do not expect.**
- B. AI tools are required by law to encrypt and protect all user data.
- C. Privacy risks only exist when using free AI tools, not paid ones.
- D. AI companies cannot access conversations due to technical barriers.

Correct Answer: A

Explanation

Privacy terms for AI tools vary significantly by provider, jurisdiction, and tier. User data may be stored, reviewed, or used in training. Assuming protection without reading the privacy policy — especially for sensitive information — is a risk.

Q66. Why does AI replacing human workers raise concerns beyond just job loss?

- A. AI tools always produce lower-quality work than human workers.
- B. AI replacement is illegal under Indian labour protection laws.
- C. It concentrates economic gains in fewer hands and increases inequality.**
- D. Automation removes jobs that most people do not want anyway.

Correct Answer: C

Explanation

Beyond job loss, AI-driven automation concentrates productivity gains with technology owners while workers bear the economic costs. Without redistributive policy, this can significantly increase inequality at scale — a structural concern beyond individual job loss.

Q67. Which behaviour most clearly demonstrates AI literacy in a student?

- A. Knowing the names and rankings of the top AI tools available.
- B. Being able to explain how transformer architectures work.
- C. Building a personal AI project using APIs and Python code.
- D. Critically evaluating AI outputs and verifying before using them.**

Correct Answer: D

Explanation

AI literacy is demonstrated through practice — specifically the habit of critical evaluation and verification. Knowing tool names, explaining architecture, and building projects are all valuable but secondary to the fundamental judgment skill.

Q68. An AI hiring tool is found to favour candidates from certain universities. The most constructive response is:

- A. Stop using AI in hiring entirely across all organisations.
- B. Audit the model for training data bias and correct it.**
- C. Add a human as a final check who reverses all AI decisions.
- D. Replace the AI tool with one from a different vendor.

Correct Answer: B

Explanation

The constructive response is to audit the model — examine what data it was trained on, identify the source of bias, and correct it. A blanket ban on AI in hiring, randomly reversing decisions, or switching vendors without addressing the root cause all fail to fix the problem.

Q69. What makes AI-generated misinformation more dangerous than traditional misinformation?

- A. It is always more convincing than human-generated content.
- B. It cannot be identified by any current detection tools.
- C. It can be produced at massive scale with minimal effort.**
- D. It spreads faster because AI tools share it automatically.

Correct Answer: C

Explanation

The most significant amplifier is scale — AI allows production of large volumes of convincing misinformation at near-zero cost. Previously, scale required human effort or media infrastructure. AI removes that barrier.

Q70. A company uses AI to personalise news feeds. The most important long-term concern is:

- A. Users see only content that reinforces their existing views.**
- B. Users will be charged for AI-curated news in the future.
- C. Personalised feeds are slower to load than standard feeds.
- D. AI cannot personalise news for users with niche interests.

Correct Answer: A

Explanation

Personalised news feeds optimised for engagement tend to show content that reinforces existing beliefs — creating filter bubbles. Over time, this reduces exposure to challenging perspectives and can deepen polarisation.

Q71. India's draft Digital Personal Data Protection framework addresses AI. The most important implication for students is:

- A. Students must register with the government before using AI.
- B. Personal data used to train AI tools may require your consent.**
- C. All AI tools operating in India must be built in India.
- D. Indian students are legally prohibited from sharing data with foreign AI.

Correct Answer: B

Explanation

Data protection frameworks like India's DPDP aim to give individuals rights over how their personal data is used — including potentially in AI training. Understanding data rights is increasingly important for any AI user.

Q72. What is the most accurate description of AI's current role in scientific research?

- A. A replacement for laboratory experiments and data collection.
- B. A fully autonomous system for generating and validating hypotheses.
- C. A tool only suitable for literature review, not active research.
- D. A powerful tool that accelerates discovery while requiring expert oversight.**

Correct Answer: D

Explanation

AI genuinely accelerates scientific research — pattern recognition, literature synthesis, hypothesis generation, data analysis — but all of these outputs require expert domain knowledge to evaluate, validate, and act on. It is an accelerant, not an autonomous researcher.

Q73. You are a student writing about the ethics of AI surveillance. Claude helps you draft arguments. What is your ethical obligation?

- A. Disclose that Claude co-authored your paper with you.
- B. Only use Claude for arguments you already agree with.
- C. Understand the arguments well enough to defend them as your own.**
- D. Verify that Claude's arguments are legally compliant.

Correct Answer: C

Explanation

Your ethical obligation is intellectual — to understand and be able to stand behind the arguments you submit. If you cannot explain and defend the argument, you have not actually engaged with the intellectual task the assignment was designed for.

Q74. A school introduces an AI ethics module. What is the most important outcome it should produce?

- A. Students who can identify and name all current AI tools.
- B. Students who can reason about AI impacts and make informed decisions.**
- C. Students who can build simple AI models using provided code.
- D. Students who are aware that AI can be used for harmful purposes.

Correct Answer: B

Explanation

The goal of AI ethics education is not information transmission — it is developing reasoning capability. Students who can reason about AI impacts and apply ethical frameworks to new situations are better prepared than those who have simply memorised examples.

Q75. Menler positions itself as India's first Claude-native AI upskilling platform. For an AI-aware student, "Claude-native" most accurately signals:

- A. The platform was built using Claude's API as a back-end service.
- B. All Menler instructors were trained by Anthropic directly.
- C. Menler is the official Claude training partner for Indian institutions.
- D. The curriculum is built around how Claude specifically works and is used.**

Correct Answer: D

Explanation

"Claude-native" means the curriculum, tools, and learning experiences are designed around Claude specifically — not generic AI tools. This allows deep, consistent skill development using one platform rather than shallow exposure to many.

END OF QUESTION BANK

*Your turning
point starts
here.*