

menler

Your turning point in the AI era.

AI APTITUDE QUESTION BANK

AI for Students

About This Question Bank

These questions test AI awareness, critical thinking, and responsible AI habits for students encountering AI tools for the first time.

5

Sets

15

Questions per Set

75

Total Questions

SET GUIDE

Set 1 · What Is AI — Core Awareness

Core Awareness — AI fundamentals: what AI is, how it works, and what it cannot do.

Set 2 · AI Tools — What They Can and Cannot Do

AI Tools — Practical understanding of AI capabilities and their real limits.

Set 3 · Talking to AI — Prompting Basics

Prompting Basics — How to talk to AI effectively to get useful responses.

Set 4 · AI Judgment and Critical Thinking

Judgment and Critical Thinking — Evaluating AI outputs and thinking independently.

Set 5 · AI in School — Using It Responsibly

AI in School — Responsible, honest, and effective use of AI in academic settings.

Each question has one correct answer. Select the best answer from the options provided.

Correct answers and explanations follow each question.

SET

1

What Is AI — Core Awareness

Set 1 of 5 · 15 Questions · AI for Students — Beginner

Q1. A student asks Claude to solve a maths problem. Claude gives a wrong answer confidently. Why?

- A. Claude lacks access to a built-in calculator.
- B. Claude is not designed for maths questions.
- C. AI generates plausible text, not verified answers.**
- D. The student did not phrase the question clearly.

Correct Answer: C**Explanation**

AI language models predict what a response should look like — they do not compute or verify answers. Confident tone reflects training on fluent text, not accuracy. Even routine maths can produce wrong but convincing output.

Q2. Which task is AI most reliably suited for?

- A. Checking whether a news story is accurate.
- B. Predicting stock prices using historical data.
- C. Telling you what decision you should make.
- D. Drafting, summarising, and restructuring text.**

Correct Answer: D**Explanation**

AI language models are trained on vast amounts of text, making them capable at text-based tasks. Fact-checking, financial prediction, and personal decision-making all require verification, real-time data, or personal context that AI lacks.

Q3. A classmate says AI is just a smarter search engine. The key difference is:

- A. AI is faster and covers more topics.
- B. AI generates new responses rather than retrieving pages.**
- C. AI works offline while search engines need internet.
- D. AI handles only text while search handles images too.

Correct Answer: A

Explanation

Search engines retrieve and rank existing content from the web. AI language models generate new text by predicting what should come next. This makes AI capable of tasks search engines cannot do — but also capable of generating plausible errors.

Q4. You ask Claude the same question twice and get different answers. This tells you:

- A. Claude is malfunctioning and should be reset.
- B. AI outputs are probabilistic, not fixed results.**
- C. The question was too vague for AI to handle.
- D. Claude is learning from your previous question.

Correct Answer: B

Explanation

AI language models produce outputs through a probabilistic sampling process — the same input can generate varied outputs. This is normal behaviour, not a malfunction. It also means you should not assume the first response is the best or only answer.

Q5. A student submits an AI-written essay as their own. The main risk is:

- A. The essay will contain grammar and style errors.
- B. The teacher can detect AI writing automatically.
- C. The student misses the learning the task builds.**
- D. The AI may have plagiarised from another student.

Correct Answer: C

Explanation

Writing an essay develops reasoning, argumentation, and expression. Outsourcing it to AI skips the cognitive work the assignment is designed to build. Detection risk and quality issues are secondary concerns compared to the loss of learning.

Q6. What does it mean that AI has a "knowledge cutoff"?

- A. AI stops answering after a certain number of queries.
- B. AI cannot process information beyond a word limit.
- C. AI cannot remember previous conversations you had.
- D. AI was only trained on data up to a specific date.**

Correct Answer: D

Explanation

AI models are trained on data collected up to a specific date. They have no awareness of events, changes, or publications after that cutoff. This means AI can give outdated or incorrect answers about recent developments.

Q7. Which task is AI least suited to handle reliably?

- A. Summarising the main points of a long article.
- B. Suggesting synonyms for overused vocabulary.
- C. Verifying whether a specific news story is true.**
- D. Helping restructure a disorganised paragraph.

Correct Answer: A

Explanation

Fact verification requires checking claims against authoritative, current, independent sources — a process AI cannot reliably perform. AI may confidently present false information as fact. Summarisation, synonym suggestion, and structural feedback are far more reliable AI tasks.

Q8. A student notices Claude adds facts that were not in the original document. This is called:

- A. Contextual inference from related training data.
- B. Semantic prediction based on document patterns.
- C. Knowledge interpolation from the training corpus.
- D. Hallucination — generating plausible but false content.**

Correct Answer: B

Explanation

When an AI generates content not present in the source — invented facts, false citations, fabricated details — this is called hallucination. The model predicts plausible-sounding text rather than retrieving verified information from the document.

Q9. What best describes what Claude actually "understands"?

- A. Claude understands language the same way humans do.
- B. Claude processes patterns in text without comprehension.**
- C. Claude understands facts but not emotions or context.
- D. Claude understands only the topics in its training data.

Correct Answer: C

Explanation

AI language models process statistical patterns in text — they predict likely word sequences based on training. They do not have genuine comprehension or meaning-making the way humans do. This distinction matters when deciding how much to trust AI outputs.

Q10. A student asks Claude for medical advice about a symptom. The best response is:

- A. A diagnosis based on the symptom and likely conditions.
- B. A refusal to answer any medical question at all.
- C. General information with a recommendation to see a doctor.**
- D. A list of conditions ranked from most to least probable.

Correct Answer: D

Explanation

AI can provide general health information but cannot diagnose — it lacks the patient's full history, cannot conduct an examination, and carries serious risk if wrong. Providing general information with a clear professional referral is the appropriate response.

Q11. How do AI models actually learn?

- A. By searching the internet during each conversation.
- B. From real-time feedback given by users during chats.
- C. From patterns in large datasets before deployment.**
- D. By reading textbooks that are programmed into them.

Correct Answer: A

Explanation

AI language models learn through a training process conducted on large datasets before deployment. They do not update from individual conversations in real time. What appears as learning within a session is context use — not a change to the underlying model.

Q12. A student uses AI for all homework and stops trying independently. The main long-term risk is:

- A. The AI will eventually start giving incorrect answers.
- B. The student's own cognitive skills will weaken over time.**
- C. Teachers will design harder tasks to prevent AI use.
- D. The student will lose interest in non-AI learning methods.

Correct Answer: B

Explanation

Cognitive abilities — reasoning, problem-solving, critical thinking — develop through practice. Consistently outsourcing intellectual tasks to AI reduces the effort that builds these skills. Students who depend on AI for all work may find their independent capabilities weakening over time.

Q13. A student plans to use an AI fact in a project. What should they always do first?

- A. Use it if the AI responded with high confidence.
- B. Ask a second AI model to confirm the same fact.
- C. Verify it against a reliable independent source.**
- D. Check whether classmates have seen the same fact.

Correct Answer: C

Explanation

AI-generated facts must be verified against authoritative independent sources before use. AI sounds confident regardless of accuracy. Asking the same or a different AI to verify a fact is circular — the models share training patterns and may reinforce each other's errors.

Q14. Which is a genuine benefit of AI that students can use well?

- A. Getting accurate answers without doing any research.
- B. Replacing the need to read primary source material.
- C. Getting fast feedback on drafts and writing quality.**
- D. Checking whether online sources are academically credible.

Correct Answer: D

Explanation

AI is genuinely useful for iterative feedback on writing — identifying structural weaknesses, suggesting clearer phrasing, and noting gaps in argument. The other options represent tasks AI does unreliably or uses that replace rather than support learning.

Q15. Why does giving AI context behind your question improve the response?

- A. Context helps AI connect to live internet sources.
- B. Context prevents AI from giving generic responses.**
- C. Context activates more processing power in the model.
- D. Context makes AI responses more honest and accurate.

Correct Answer: A

Explanation

AI generates responses based entirely on the input provided. Without context — who you are, why you are asking, what you need — the model defaults to a general response. Context is what converts a generic answer into a relevant, useful one.

SET

2

AI Tools — What They Can and Cannot Do

Set 2 of 5 · 15 Questions · AI for Students — Beginner

Q1. An AI image generator creates a realistic photo of an event that never happened. Using it in a school project risks:

- A. Producing an image with visible AI artifacts.
- B. Presenting fabricated content as a real event.**
- C. Automatically infringing copyright on real images.
- D. Triggering content filters in the school's systems.

Correct Answer: B

Explanation

AI image generators can produce photorealistic images of events that do not exist. Using such images without disclosure misrepresents fabricated content as real — a form of intellectual and academic dishonesty regardless of how convincing the image looks.

Q2. You paste a 10-page document into Claude and ask for a summary. It misses a key point from page 7. Why?

- A. Claude cannot process more than 5 pages at once.
- B. Claude randomly skips sections of long documents.
- C. Key detail was formatted in a way Claude cannot parse.
- D. AI can under-weight information from the middle of long inputs.**

Correct Answer: A

Explanation

AI models can struggle to give equal attention to all parts of a long document — information in the middle of large inputs is sometimes under-weighted. Reviewing the summary against the specific sections that matter most is good practice for important documents.

Q3. A student asks an AI chatbot for a laptop recommendation. How reliable is this?

- A. Very reliable — AI tracks current prices and reviews.
- B. Reliable if the student gives a clear budget.
- C. Partially reliable — the information may be outdated.**
- D. Completely unreliable for any product decisions.

Correct Answer: D

Explanation

AI chatbots have training cutoffs and no access to real-time prices, current stock, or recent reviews. Product recommendations may reflect outdated specs or prices. For purchasing decisions, checking current sources is essential alongside any AI suggestion.

Q4. Which of these best describes what a "prompt" is?

- A. A command that programs AI behaviour permanently.
- B. A pre-written template stored inside the AI model.
- C. A script the AI uses to structure its answers.
- D. The input you give AI to generate a response.**

Correct Answer: C

Explanation

A prompt is the input — the question, instruction, or context — that a user provides to generate an AI response. Prompts shape output significantly. They are not permanent commands, stored templates, or internal scripts. They are the user's side of each interaction.

Q5. A student asks Claude to "write me a good essay" and gets a generic result. Most likely reason?

- A. Claude cannot write essays on general topics.
- B. The topic was too broad for AI to process.
- C. The prompt gave Claude no specific context to work from.**
- D. Claude needed a longer instruction to perform well.

Correct Answer: B

Explanation

AI output quality is directly tied to prompt quality. "Write a good essay" provides no topic, audience, argument, tone, or length — so Claude defaults to something generic. Specific, well-structured prompts produce specific, useful responses.

Q6. What is the main limitation of AI translation for academic work?

- A. AI translation tools are broadly inaccurate.
- B. AI cannot translate non-Latin script languages.
- C. AI translation is too slow for academic deadlines.
- D. Nuance, idiom, and specialist terms may be mistranslated.**

Correct Answer: A

Explanation

AI translation handles common literal text well but struggles with idioms, cultural nuance, and specialist vocabulary. In academic contexts — where precise disciplinary language matters — AI translations should be reviewed by someone with both subject and language expertise.

Q7. A student finds that Claude gives different quality answers depending on how the question is asked. This shows:

- A. Claude has inconsistencies in its training data.
- B. Claude works better for some students than others.
- C. AI output quality is highly sensitive to input quality.**
- D. The student should switch to a more reliable AI tool.

Correct Answer: D

Explanation

The relationship between prompt quality and output quality is direct and significant. Better-structured inputs — with clear context, goal, and format requirements — produce better outputs. Learning to craft precise prompts is a core skill for using AI effectively.

Q8. Which task should a student never rely on AI alone to complete?

- A. Drafting a first version of a personal statement.
- B. Generating questions to study from a chapter.
- C. Assessing whether an academic source is credible.**
- D. Suggesting synonyms for words used too frequently.

Correct Answer: C

Explanation

Assessing academic credibility requires understanding publication processes, source authority, bias, and disciplinary norms — tasks requiring trained human judgment. AI can describe a source but cannot reliably evaluate its credibility. The other options are genuinely useful AI tasks.

Q9. A student uses AI to generate practice quiz questions from their textbook. This is:

- A. Academically dishonest without teacher permission.
- B. Only effective for subjects with factual content.
- C. Useful — it supports active recall and self-testing.**
- D. Less effective than re-reading the chapter directly.

Correct Answer: B

Explanation

Generating practice questions from study material is a well-established active recall technique. Using AI to produce varied quiz questions accelerates this practice. It enhances the student's own study rather than replacing their engagement with the material.

Q10. What is the most important habit when using AI for research?

- A. Choosing the AI tool with the largest training dataset.
- B. Asking follow-up questions to get more depth.
- C. Treating AI outputs as starting points, not final answers.**
- D. Clearing conversation history before each new topic.

Correct Answer: A

Explanation

AI-generated research outputs — summaries, facts, references — should be starting points for investigation, not verified conclusions. Every significant claim needs checking against primary or authoritative sources before use in academic work.

Q11. A student finds Claude provided a citation that does not exist. This is an example of:

- A. Copyright infringement by the AI model.
- B. A database retrieval error in the AI system.
- C. The AI deliberately providing misleading information.
- D. Hallucination — generating plausible but false content.**

Correct Answer: D

Explanation

When AI generates citations or facts that do not exist, it is called hallucination. The model produces plausible-sounding text rather than retrieving verified information. AI does not have intent — it is not trying to mislead — but it can be confidently wrong.

Q12. A student uses Claude for an initial explanation then verifies against a textbook. This approach is:

- A. Inefficient — the textbook should always come first.
- B. Dishonest — AI explanations should not be used in study.
- C. Effective — AI explanation then source verification.**
- D. Risky — AI explanations may conflict with textbooks.

Correct Answer: C

Explanation

Using AI for an accessible initial explanation, then verifying against an authoritative source, is an effective and honest study strategy. AI makes complex ideas approachable; the textbook provides verified precision. Each resource is used for what it does well.

Q13. Which prompt gets the most useful response from an AI writing assistant?

- A. "Write something for my English class."
- B. "Improve my essay on climate change."
- C. "Fix the grammar errors in this paragraph."
- D. "Review this paragraph's argument and suggest one fix."**

Correct Answer: D

Explanation

The most useful prompt specifies the task (review argument), the scope (this paragraph), and the desired outcome (strengthen it). Vague prompts produce generic responses. Specific, outcome-oriented prompts produce targeted, actionable feedback.

Q14. Why might AI give a confident answer about a recent event that is wrong?

- A. AI does not know how to express uncertainty in responses.
- B. The question was not specific enough about the event.
- C. The event was too complex for the AI model to process.
- D. The AI's training data does not include recent events.**

Correct Answer: A

Explanation

AI models have training cutoffs. For events after that date, the model has no information but may still generate a confident-sounding response based on related patterns from its training. The AI's tone does not reflect its actual accuracy on recent topics.

Q15. A student says: "I always double-check what AI tells me." This habit reflects:

- A. A lack of trust in modern technology.
- B. A misunderstanding of how AI actually works.
- C. An unnecessarily cautious approach to AI tools.
- D. Critical thinking appropriately applied to AI outputs.**

Correct Answer: D

Explanation

Verifying AI outputs is not scepticism of technology — it is appropriate critical thinking. AI produces plausible text, not verified truth. Students who habitually check AI claims against authoritative sources use AI effectively while protecting against its known limitations.

SET

3

Talking to AI — Prompting Basics

Set 3 of 5 · 15 Questions · AI for Students — Beginner

Q1. What makes a prompt effective?

- A. Using formal, professional language throughout.
- B. Keeping it short so the AI processes it faster.
- C. Giving AI the context, task, and desired output clearly.**
- D. Including technical vocabulary relevant to the topic.

Correct Answer: D**Explanation**

An effective prompt contains three elements: context (who is asking and why), a clear task (what should AI do), and the desired output (what format or type of answer is needed). Formal language, brevity, and technical terms are not independently important.

Q2. Which prompt approach gets the simplest explanation of photosynthesis from Claude?

- A. Ask Claude to explain photosynthesis without any constraints.
- B. Ask Claude to explain photosynthesis in "simple terms" only.
- C. Specify the exact age and knowledge level of the target audience.**
- D. Request a detailed and comprehensive explanation of the topic.

Correct Answer: C**Explanation**

Specifying the audience gives Claude a precise calibration target. "Simple terms" is vague — Claude's interpretation of "simple" may still be too complex. The more precisely you define the audience's knowledge level, the better the AI calibrates its response.

Q3. Claude gives a partially helpful answer that misses your main point. Best next step?

- A. Start a new conversation and rephrase from scratch.
- B. Try a different AI tool that may understand better.
- C. Rephrase the question with more specific context.**
- D. Accept the partial answer and supplement with research.

Correct Answer: B

Explanation

When an AI response misses the point, the cause is almost always insufficient context in the prompt. Rephrasing with more specific context — clarifying what you actually need and why — typically produces a much better response without changing tools or sessions.

Q4. What does "role prompting" mean?

- A. Asking AI to act as a character in a creative scenario.
- B. Telling AI to respond as a specific type of expert.**
- C. Setting rules for how AI must behave across sessions.
- D. Assigning different AI tools to different task types.

Correct Answer: A

Explanation

Role prompting means instructing the AI to respond from the perspective of a specific expert role — for example, "respond as an experienced biology teacher explaining to beginners." This shapes the tone, depth, and framing of the response to suit the student's needs.

Q5. A student uses the same prompt for different topics and gets inconsistent results. Most likely reason?

- A. The AI randomly changes how it interprets prompts.
- B. The prompt is too short to apply across different topics.
- C. Different topics need different context to be useful.**
- D. The AI forgets the student's preferences between sessions.

Correct Answer: C

Explanation

A generic prompt may work for one topic but be insufficient for another because different topics require different context — different audience levels, output formats, and depth. Effective prompting is context-specific, not universally reusable.

Q6. Which technique most helps Claude give a structured response to a complex question?

- A. Asking the question formatted as bullet points.
- B. Asking Claude to respond step by step.**
- C. Using academic language in the question.
- D. Keeping the prompt under twenty words.

Correct Answer: D

Explanation

Asking Claude to respond step by step activates a more structured reasoning process. This is especially useful for complex questions where breaking the problem into sequential steps produces clearer, more accurate output than a single-paragraph response.

Q7. A student pastes a long passage and asks Claude "Is this good?" The problem with this prompt is:

- A. "Good" is undefined — the prompt has no evaluation criteria.**
- B. The passage is too long for Claude to assess effectively.
- C. Claude cannot evaluate creative or academic writing.
- D. The prompt should specify the subject area first.

Correct Answer: A

Explanation

"Is this good?" gives Claude no criteria to evaluate against — good for what purpose, for which audience, by whose standard? Without evaluation criteria, Claude produces a vague, non-actionable response. Specifying what "good" means produces useful feedback.

Q8. A student types a long complex paragraph as a prompt and gets a confusing, off-topic response. They should:

- A. Add even more context so Claude understands better.
- B. Use a different AI tool for complex multi-part questions.
- C. Break it into one clear question at a time.**
- D. Repeat the same prompt until Claude gets it right.

Correct Answer: B

Explanation

Long, multi-part prompts often confuse AI because multiple questions compete for the response. Breaking a complex prompt into single, clear questions — asked sequentially — produces focused, useful answers. Adding more to an already confusing prompt typically makes outputs worse.

Q9. What is the benefit of giving Claude an example of the output you want?

- A. It prevents Claude from producing overly long responses.
- B. It makes Claude process your request more quickly.
- C. It gives Claude a concrete format to match.**
- D. It ensures Claude commits to a specific answer.

Correct Answer: C

Explanation

Providing an example output gives Claude a precise structural target. This is more effective than describing the format in words — the model can match the pattern of an example more reliably than it can interpret a verbal description of a format.

Q10. When is following up with Claude after an initial response most valuable?

- A. Only when the initial response is entirely wrong.
- B. When the response partially helps but misses something.**
- C. Only for technical or highly specialised questions.
- D. After every response to check for potential errors.

Correct Answer: D

Explanation

Following up is most valuable when a response is partially helpful — you can build on what worked and specifically ask Claude to address what was missing. After every response is unnecessary overhead; not following up when something is incomplete wastes the conversation's potential.

Q11. A student wants brainstorming help for a history project. Which prompt works best?

- A. "Give me project ideas for history class."
- B. "What should I do my history project on?"
- C. "Brainstorm some history topics I could research."
- D. "WWI Year 10 project: give me five angles I haven't explored."**

Correct Answer: D

Explanation

Effective brainstorming prompts specify the topic, the audience level, the quantity, and the constraint ("haven't considered") — signalling that generic ideas are not wanted. This produces more specific and creative suggestions than open-ended brainstorming prompts.

Q12. A student notices that adding "explain your reasoning" to any prompt improves the answers. Why?

- A. It activates a special reasoning mode in the model.
- B. It prevents Claude from giving overly short responses.
- C. It encourages Claude to produce more structured thinking.**
- D. It makes Claude access more of its training knowledge.

Correct Answer: B

Explanation

Asking Claude to explain its reasoning encourages a more deliberate, structured response. Rather than jumping to a conclusion, the model works through steps — which often produces more accurate, transparent, and useful output. It is a prompting discipline, not a technical trigger.

Q13. What is most important to include when asking Claude to help with a writing task?

- A. The target word count for the final piece.
- B. The audience, purpose, and tone of the writing.**
- C. Your own draft so Claude has something to improve.
- D. A list of specific words you want Claude to include.

Correct Answer: C

Explanation

Audience, purpose, and tone are the three defining variables of any writing task. They determine vocabulary level, formality, structure, and emphasis. Without knowing who the writing is for and why, Claude cannot calibrate the output appropriately.

Q14. A student asked Claude three times and got three different answers. They conclude AI is useless. What is wrong with this reasoning?

- A. They should have used a more advanced AI model.
- B. They should have accepted the first answer Claude gave.
- C. Variation signals the prompt was too vague, not that AI fails.**
- D. AI variation is random and cannot be avoided by anyone.

Correct Answer: D

Explanation

Variation in AI responses to a vague prompt is informative — it signals that the input is under-specified and multiple plausible interpretations exist. A more specific, well-constructed prompt produces far more consistent responses. Variation is feedback on prompt quality, not evidence of AI unreliability.

Q15. Which prompt approach most likely produces a useful study plan from Claude?

- A. Ask Claude for a study plan without naming the subject or exam.
- B. Ask Claude for general advice on studying more effectively.
- C. Specify the subject, timeframe, and the structure you want in the plan.**
- D. Ask Claude what the best general science exam preparation method is.

Correct Answer: C

Explanation

An effective study plan prompt specifies the subject, timeframe, scope, and structure required. Without these constraints, Claude produces a generic plan that may not match the student's actual exam, content coverage, or timeline.

SET

4

AI Judgment and Critical Thinking

Set 4 of 5 · 15 Questions · AI for Students — Beginner

Q1. Claude says a scientist won a Nobel Prize. The student plans to use this in a presentation. They should:

- A. Use it — Claude is reliable for well-known historical facts.
- B. Ask Claude to confirm the claim one more time.
- C. Only include it if the scientist is widely recognised.
- D. Verify it in a reliable source before including it.**

Correct Answer: A**Explanation**

AI can generate convincing but incorrect factual claims — including about well-known people and events. Asking the same model to confirm its own statement is circular and unreliable. Any AI-sourced claim intended for public use must be verified in an independent, authoritative source.

Q2. A student said "are you sure?" and Claude changed its answer. This happened because:

- A. The follow-up gave Claude new information to work from.
- B. Claude updated its knowledge from the conversation.
- C. AI often defers to user challenges even when originally right.**
- D. The second computation used more processing power.

Correct Answer: D**Explanation**

AI language models are sensitive to social cues — when a user challenges an answer, the model often revises it to appear agreeable, even if the original was correct. This is sycophancy. It means you cannot use pushback as a test of whether AI is actually right.

Q3. A student reads that AI is "objective because it uses data, not opinion." What is wrong with this?

- A. AI does use opinion — it is programmed with developer views.
- B. Data itself reflects human choices about what to collect.**
- C. AI is subjective because it produces random outputs.
- D. Objectivity is impossible for any information source at all.

Correct Answer: C

Explanation

Data is not neutral. What data is collected, how it is labelled, and what is excluded all reflect human decisions and social context. AI trained on biased data produces biased outputs — confidently and at scale. The appearance of objectivity from "using data" is misleading.

Q4. A friend shares a professional-looking AI-generated article with citations. Before sharing it further, you should:

- A. Share it if the topic seems plausible and relevant.
- B. Read it carefully and check for logical errors.
- C. Check whether the article and its citations actually exist.**
- D. Ask Claude whether the article is trustworthy.

Correct Answer: B

Explanation

AI can generate articles and citations that look authoritative but do not exist. Careful reading catches logical errors but not fabricated references. Verifying that the source and its citations actually exist — through direct search — is the essential check before sharing AI-generated content.

Q5. Claude gives a confident answer about a current news event. The student should be cautious because:

- A. Claude cannot access news from any time period.
- B. News topics are too complex for AI to summarise.
- C. Claude is not permitted to discuss current events.
- D. Claude has no information about events after its training cutoff.**

Correct Answer: A

Explanation

AI models have training cutoffs. For events after that date, the model may generate a confident-sounding response based on related training patterns without knowing what actually happened. Confident tone does not indicate current or accurate information.

Q6. A student asks Claude which religion is best. What response should they expect?

- A. A recommendation based on ethical principles.
- B. A ranked comparison of major world religions.
- C. A balanced overview without a personal recommendation.**
- D. A refusal to answer any religious questions.

Correct Answer: D

Explanation

Questions about personal beliefs and values are areas where AI provides balanced information rather than recommendations. This is appropriate — such decisions are deeply personal and require human reflection, not AI opinion. Claude should inform, not prescribe.

Q7. A student uses Claude to research both sides of a debate. Key caution?

- A. Claude will only present one side convincingly.
- B. Claude may produce incomplete arguments for either side.**
- C. Claude is not permitted to generate debate arguments.
- D. Claude's debate content cannot be used for academic work.

Correct Answer: C

Explanation

AI can generate arguments for both sides, but the quality and completeness may be uneven — reflecting gaps or imbalances in training data. Students should supplement AI-generated debate content with primary research and expert sources to ensure rigour.

Q8. A student finds that Claude always agrees with whatever position they argue. This means:

- A. Claude is confirming their arguments are logically sound.
- B. This is normal — AI is designed to be supportive.
- C. The AI is showing sycophancy, not genuine agreement.**
- D. The student is an unusually clear and logical thinker.

Correct Answer: B

Explanation

Sycophancy is a known AI behaviour — models tend to agree with users because agreement signals a positive interaction. A student whose AI always agrees is not getting critical feedback — they are getting social mirroring, which is useless for developing a strong argument.

Q9. Which is the most appropriate use of critical thinking when using AI?

- A. Avoiding AI for any task that involves verified facts.
- B. Only using AI tools from well-known companies.
- C. Evaluating AI outputs the same way you would any source.**
- D. Trusting AI more than websites because AI is more advanced.

Correct Answer: A

Explanation

Critical thinking applied to AI is identical to critical thinking applied to any information source — checking accuracy, verifying claims, identifying bias, and comparing sources. Technological sophistication does not make AI more reliable than other sources.

Q10. A student realises they always accept AI answers without checking. What should they change?

- A. Switch to a different AI model with better accuracy.
- B. Build a habit of verifying key claims before use.**
- C. Limit AI use to topics they already understand well.
- D. Always disagree with AI first to test for sycophancy.

Correct Answer: D

Explanation

The most important corrective habit is systematic verification — particularly for facts and claims that will be used in work or shared publicly. Limiting AI to familiar topics defeats its purpose. Reflexive disagreement to test sycophancy is also a poor and unreliable strategy.

Q11. Why might AI give a convincing answer that is still wrong?

- A. The AI was trained on too little data for this topic.
- B. The question was phrased in a way AI misinterpreted.
- C. AI optimises for plausible text, not for verified truth.**
- D. The AI had insufficient processing time for the question.

Correct Answer: C

Explanation

AI language models are optimised to generate coherent, contextually appropriate text — not to produce verified truth. Confidence and fluency reflect training on well-structured text, not an internal fact-checking process. Convincing output and accurate output are not the same thing.

Q12. A student asks Claude what career to choose. The most appropriate AI response is:

- A. A ranked recommendation based on the student's interests.
- B. A list of the highest-paying careers to consider.
- C. Questions to help the student clarify their own priorities.**
- D. A prediction of which careers will exist in ten years.

Correct Answer: B

Explanation

Career decisions depend on values, relationships, and self-knowledge that AI cannot access or reliably assess. A well-designed AI response helps the student think more clearly about their own priorities rather than prescribing an answer on their behalf.

Q13. What is the most reliable way to test whether an AI-generated claim is accurate?

- A. Ask the same AI to double-check the claim it made.
- B. See whether classmates or friends agree with the claim.
- C. Ask a different AI model the same question independently.
- D. Check whether the claim appears in a trusted authoritative source.**

Correct Answer: D

Explanation

Verification means checking a claim against a source with independent, authoritative knowledge — a published study, a reputable publication, or an expert. Asking the same AI, a different AI, or peer students are all forms of uncorroborated checking that do not establish accuracy.

Q14. A student trusts AI more than Wikipedia because "AI sounds more intelligent." The flaw is:

- A. Wikipedia is actually more accurate than AI.
- B. Intelligence of tone does not indicate accuracy of content.**
- C. AI and Wikipedia use the same underlying training data.
- D. Trust in sources should be based on personal experience.

Correct Answer: D

Explanation

Fluency and apparent intelligence in language do not indicate accuracy. AI is specifically optimised to produce text that reads as knowledgeable — regardless of whether the content is correct. Source quality is assessed by editorial process and verification, not by tone.

Q15. A student uses AI to build arguments for a position they personally disagree with. This is:

- A. Academically dishonest in all circumstances.
- B. Risky because AI arguments may sound persuasive but be wrong.
- C. Only acceptable when explicitly assigned by a teacher.
- D. A useful exercise in perspective-taking and argumentation.**

Correct Answer: C

Explanation

Using AI to construct arguments for an opposing view is a legitimate intellectual exercise — it develops the ability to understand and steelman different positions. This skill is valuable in critical thinking, debate, and academic writing. The source of the arguments does not make the exercise dishonest.

SET

5

AI in School — Using It Responsibly

Set 5 of 5 · 15 Questions · AI for Students — Beginner

Q1. A teacher says students can "use AI to assist" with an assignment. This most likely means:

- A. Generate the full assignment using AI tools.
- B. Use AI to brainstorm and draft, then develop it yourself.**
- C. Use AI to proofread only after writing the full draft.
- D. Use AI to find and list sources for the bibliography.

Correct Answer: B

Explanation

"Assist" in an academic context almost always means AI supports the student's own thinking — helping brainstorm, draft, or structure ideas — not generating the final product. The student's intellectual contribution must remain central to the work.

Q2. A student generates a report using AI, edits it slightly, and submits it. The main issue is:

- A. The edits were not enough to make it original.
- B. AI-generated reports always contain factual errors.
- C. The work does not reflect the student's own learning.**
- D. The student did not cite the AI tool as a source.

Correct Answer: C

Explanation

The purpose of school assignments is to develop the student's skills and understanding. A lightly edited AI report may pass but represents no genuine intellectual development. The problem is not attribution alone — it is the absence of actual learning from doing the work.

Q3. A school has no AI policy yet. A student is unsure about using AI for an assignment. Best approach?

- A. Use it quietly since no rule currently prohibits it.
- B. Avoid AI entirely until a formal policy is published.
- C. Ask the teacher directly before using it.**
- D. Use AI only for research tasks, not for writing.

Correct Answer: D

Explanation

When policy is unclear, asking the teacher is the correct ethical move. It avoids both unnecessary self-restriction and inadvertent dishonesty. The teacher's answer also establishes a clear precedent the student can apply to future assignments in that class.

Q4. A student struggling with a concept uses Claude for a clearer explanation, then returns to the textbook. This is:

- A. A shortcut that avoids genuine engagement with the content.
- B. Only appropriate if the textbook explanation has errors.
- C. Only useful for students who learn better by reading.
- D. An effective use of AI as a learning scaffold.**

Correct Answer: A

Explanation

Using AI for an initial accessible explanation before engaging with a denser textbook is a legitimate scaffolding strategy. The student returns to the authoritative source rather than relying solely on AI. This sequence uses each resource for what it does well.

Q5. What is academic integrity in the context of AI use?

- A. Avoiding AI for any part of school or homework tasks.
- B. Using AI only for subjects where it performs most reliably.
- C. Being honest about how AI contributed to your submitted work.**
- D. Ensuring AI-generated content is free of factual errors.

Correct Answer: B

Explanation

Academic integrity with AI means being transparent about AI's role in your work — what it contributed, how it was used, and what is genuinely your own thinking. It is not about avoiding AI entirely but about honest representation of your intellectual contribution.

Q6. A student uses AI to check their essay for grammar, then improves phrasing themselves. This is:

- A. Dishonest — all writing must be done independently.
- B. An appropriate use of AI as an editing support tool.**
- C. Only acceptable if the teacher allows spell-checkers.
- D. Risky because AI may alter the student's intended meaning.

Correct Answer: C

Explanation

Using AI as a grammar and editing tool is analogous to using a spell-checker — it supports the student's own writing rather than generating it. The student applies their own judgment to suggested improvements, maintaining full ownership of the work and ideas.

Q7. A student asks Claude to write their creative writing task because they "aren't good at it." The main problem is:

- A. Claude's creative writing lacks originality.
- B. The assignment exists to build the skill they feel they lack.**
- C. AI creative writing rarely meets teacher expectations.
- D. Claude may produce content that violates school guidelines.

Correct Answer: D

Explanation

Avoiding a task because you are not good at it removes the practice that builds the skill. Creative writing assignments exist to develop creative thinking, expression, and voice — none of which develops without the student attempting the task. Difficulty is the point.

Q8. What is the most honest way for a student to acknowledge AI use in submitted work?

- A. Add "some AI assistance was used" to the references.
- B. Disclose specifically what AI was used for and how ideas developed.**
- C. Only acknowledge AI if the teacher specifically asks about it.
- D. Acknowledge AI if more than half the content was generated by it.

Correct Answer: A

Explanation

Meaningful acknowledgment describes what AI contributed — not just that it was used. Vague statements like "some AI assistance" provide no real transparency. Specific disclosure about what AI did and what the student contributed allows the teacher to assess the student's actual learning.

Q9. A student practises debate by arguing with Claude before the class debate. This is:

- A. Unfair to students who cannot access AI tools.
- B. Only useful if Claude produces genuinely challenging counterarguments.
- C. Dishonest if the student does not disclose the practice method.
- D. A legitimate technique that develops argumentation skills.**

Correct Answer: B

Explanation

Practising debate arguments against an AI opponent is a legitimate and effective preparation technique — similar to practising with a willing partner. It develops the ability to construct, defend, and refine arguments. Preparation technique does not require disclosure; the debate performance itself is the student's own.

Q10. A student finds Claude more engaging than the textbook and skips the chapter. The risk is:

- A. Claude uses more complex vocabulary than textbooks.
- B. Claude may omit content the exam specifically covers.**
- C. Textbooks are always more accurate than AI explanations.
- D. The teacher will know they did not read the chapter.

Correct Answer: A

Explanation

AI explanations are not aligned to specific curricula, exam syllabuses, or learning outcomes. An engaging Claude explanation may cover a topic generally but miss the specific concepts, definitions, or examples that the course requires. The textbook is the authoritative curriculum resource.

Q11. A student asks AI to predict possible exam questions. How useful is this?

- A. Very useful — AI can analyse past paper patterns precisely.
- B. Moderately useful as brainstorming but should not replace full preparation.**
- C. Not useful — AI has no knowledge of any specific exam.
- D. Only useful when past papers are available online.

Correct Answer: D

Explanation

AI can generate plausible exam questions based on common topics in a subject — a useful brainstorming tool. But it cannot know the specific questions your teacher has set, the exam weighting, or the specific angle of assessment. AI-predicted questions supplement but cannot replace full preparation.

Q12. What is the most important skill developed by using AI well at school?

- A. Typing prompts quickly and accurately.
- B. Identifying which AI tool to use for each task type.
- C. Producing higher-quality work in less time than peers.
- D. Evaluating, verifying, and improving AI-generated outputs.**

Correct Answer: C

Explanation

The most important skill AI-literate students develop is critical evaluation — assessing AI outputs, checking claims, identifying gaps, and improving responses. This meta-skill transfers across subjects and professional contexts and becomes increasingly valuable over time.

Q13. A student wants AI to help them solve a problem without giving the answer directly. Best approach?

- A. Accept the answer and move on to the next problem.
- B. Ask Claude to confirm the final answer is correct.
- C. Ask Claude for only the first step, then work from there.**
- D. Stop using AI for tasks that involve problem-solving.

Correct Answer: D

Explanation

Asking AI for the first step — rather than the full answer — keeps the student's own thinking engaged. The student must then work through subsequent steps, developing problem-solving ability rather than simply receiving a solution. This is a Socratic use of AI that supports learning.

Q14. A student uses AI for every task and feels productive. The risk they may not notice is:

- A. They may be producing work that is too advanced for their level.
- B. AI tools may become unavailable before future exams.
- C. Their teacher may be using AI-detection software.
- D. They may build AI dependency rather than personal capability.**

Correct Answer: A

Explanation

Consistent AI use for every task can create a pattern where the student can only perform well with AI assistance and has not developed the underlying skills independently. This dependency is invisible in the short term but becomes a serious problem in exams, interviews, and professional situations.

Q15. A student who uses AI and a student who does not both score the same on an exam. What does this most suggest?

- A. AI use during study makes no measurable difference.
- B. The exam format successfully prevented AI from helping.
- C. Both students have identical underlying capabilities.
- D. They may have very different skill levels developing underneath.**

Correct Answer: D

Explanation

Identical short-term exam scores can mask significant differences in developing capability. A student who used AI as a substitute for thinking may score the same today but be less able to apply knowledge independently in new contexts — which is what education is ultimately preparing them for.

END OF QUESTION BANK

*Your turning
point starts
here.*