

menler

Your turning point in the AI era.

AI APTITUDE QUESTION BANK

AI Generalist Entry Diagnostic

About This Question Bank

An entry-level diagnostic for professionals beginning the AI Generalist Fellowship. Tests AI fundamentals, practical working judgment, prompting discipline, workflow thinking, and applied AI decision-making.

5

Sets

15

Questions per Set

75

Total Questions

SET GUIDE

Set 1 · AI Fundamentals You Must Know

AI fundamentals — how models work, key concepts, and what you must know

Set 2 · Working With AI — Practical Judgment

Working with AI — practical judgment, mental models, and team use

Set 3 · Prompting for Professional Output

Prompting for professional output — specificity, templates, iteration

Set 4 · AI Tools and Workflows

AI tools and workflows — designing reliable AI-assisted processes

Set 5 · AI Judgment — Applied Scenarios

AI judgment — applied scenarios, accountability, and edge cases

Each question has one correct answer. Select the best answer from the options provided.

Correct answers and explanations follow each question.

SET

1

AI Fundamentals You Must Know

Set 1 of 5 · 15 Questions · AI Generalist Entry Diagnostic

Q1. You receive an AI output that contains a confidently stated incorrect fact. The correct framing is:

- A. The AI model version you used is outdated — upgrade it.
- B. AI outputs are starting points that require human verification.**
- C. AI is unreliable and should not be used for factual tasks.
- D. You prompted incorrectly — a better prompt would prevent this.

Correct Answer: B

Explanation

Confident errors (hallucinations) are a structural property of how language models work — not a model version issue or prompting failure. The correct framing is that AI produces starting points, and humans are responsible for verification.

Q2. A colleague says "I don't trust AI — it makes things up." The most accurate and constructive response is:

- A. "You're right — AI is unreliable for professional use."
- B. "AI is reliable as long as you use the right tool for the task."
- C. "That concern is outdated — modern AI is much more accurate."
- D. AI has known failure modes but also genuine strengths worth understanding.**

Correct Answer: D

Explanation

The most accurate position acknowledges real limitations without dismissing real capability. Blanket rejection and blanket acceptance are both inaccurate. Understanding the specific failure modes and strengths enables effective, responsible use.

Q3. What distinguishes a generative AI model from a traditional software program?

- A. It produces novel outputs from statistical patterns, not explicit rules.**
- B. It can access the internet to retrieve up-to-date information.
- C. It runs faster than traditional programs on the same hardware.
- D. It makes decisions without any human input or instruction.

Correct Answer: A

Explanation

Traditional programs execute explicit rules written by developers. Generative AI learns statistical patterns from data and generates outputs — text, images, code — that are not explicitly programmed. This makes outputs probabilistic rather than deterministic.

Q4. Which of the following most accurately describes what Claude "knows"?

- A. A verified encyclopaedia of facts maintained by Anthropic.
- B. Real-time information retrieved from the internet on demand.
- C. Statistical patterns from training data that approximate knowledge.**
- D. Expert knowledge contributed by domain specialists.

Correct Answer: C

Explanation

Claude's knowledge is statistical — weighted patterns learned from training data. It is not a verified fact store, real-time retrieval system, or curated expert database. This explains both its breadth and its hallucination tendency.

Q5. A prompt engineer says "prompt quality determines output quality." How accurate is this?

- A. Inaccurate — model size matters far more than prompt quality.
- B. Largely accurate — prompt quality is the highest-leverage input.**
- C. Partially accurate — only for creative tasks, not factual ones.
- D. Inaccurate — all prompts to a good model produce good outputs.

Correct Answer: B

Explanation

Prompt quality is genuinely the highest-leverage variable a user controls. Given the same model, a well-specified prompt with context, role, format, and constraints consistently outperforms a vague one on nearly every task type.

Q6. What is the most important practical difference between Claude and a search engine?

- A. Claude synthesises and generates; a search engine retrieves and ranks.**
- B. Claude is more accurate for factual queries than a search engine.
- C. Claude provides citations while a search engine only shows links.
- D. Claude is faster than a search engine for complex research tasks.

Correct Answer: A

Explanation

Search engines index, retrieve, and rank existing pages. Claude synthesises, interprets, and generates new text from its learned patterns. Neither provides citations by default. For current facts, search engines are more reliable; for synthesis, Claude adds more value.

Q7. You are building a workflow that uses Claude to process customer emails. Which risk must you design for first?

- A. Claude may run too slowly for high-volume email processing.
- B. Claude may respond in a language the customer did not use.
- C. Claude may generate responses that are too long for emails.
- D. Claude may misclassify or misrespond to some customer emails.**

Correct Answer: D

Explanation

Misclassification and incorrect responses are the primary risk in any customer-facing AI system — the consequences include customer harm, trust loss, and operational errors. Speed, language, and length are secondary engineering concerns.

Q8. Why is it important to understand what a model's context window is before building with it?

- A. It determines which languages the model can respond in.
- B. It determines the maximum accuracy of the model's outputs.
- C. It determines how much input and conversation the model can process.**
- D. It determines which API endpoints are available to developers.

Correct Answer: C

Explanation

The context window is the maximum tokens the model can process in one call. Building a workflow without knowing this limit leads to unexpected truncation, lost context, or errors when inputs exceed capacity.

Q9. A manager asks you to "automate this process using AI." The correct first step is:

- A. Map the process and identify which steps AI can actually handle.**
- B. Select the AI tool and start building the automation immediately.
- C. Ask the manager which AI tool they want you to use.
- D. Estimate the cost savings before designing the solution.

Correct Answer: A

Explanation

Effective AI automation starts with process mapping — understanding each step, its inputs, outputs, and failure modes — before selecting tools or building. Jumping to implementation without this creates solutions that fail in unexpected ways.

Q10. What does "temperature = 0" mean in practice when using a language model API?

- A. The model refuses to generate any creative or subjective content.
- B. The model runs at reduced speed to improve accuracy.
- C. The model caps response length at a conservative limit.
- D. The model produces the most probable, deterministic output.**

Correct Answer: D

Explanation

Temperature = 0 means the model always selects the highest-probability next token — producing maximally deterministic, repeatable output. This is useful for structured tasks (classification, extraction) where consistency is more important than variety.

Q11. You need Claude to output structured JSON. The most reliable approach is:

- A. Ask Claude to "output JSON" without additional instructions.
- B. Specify the exact JSON schema in the prompt and use low temperature.**
- C. Use Claude only if it has a dedicated JSON output mode.
- D. Generate the output and manually convert it to JSON afterward.

Correct Answer: B

Explanation

Reliable structured output requires both explicit schema specification (what fields, what types, what structure) and low temperature (to reduce randomness). Without these, output format is unreliable even when Claude attempts JSON.

Q12. What is the most common mistake people make when first using Claude for work tasks?

- A. Over-specifying the task with too much context and detail.
- B. Using Claude for tasks that require real-time information.
- C. Giving vague instructions and expecting Claude to infer the details.**
- D. Relying on Claude before verifying it can handle the task type.

Correct Answer: C

Explanation

The most common failure is vague prompting — treating Claude like a search engine or a smart person who can fill in the gaps. Claude performs best with explicit context, clear role definition, specific format requirements, and precise constraints.

Q13. An AI Generalist's core skill is best described as:

- A. Being able to code AI models and train them on custom data.
- B. Mastering every AI tool available in their professional domain.
- C. Understanding the mathematical foundations of machine learning.
- D. Knowing how to direct AI across varied tasks to produce useful outputs.**

Correct Answer: D

Explanation

An AI Generalist's core skill is judgment in directing AI — knowing what to ask, how to specify it, how to evaluate the output, and when to verify or escalate. This is different from technical ML skills or comprehensive tool mastery.

Q14. You run the same Claude prompt 10 times. 8 of 10 responses are correct and 2 are wrong. For a professional workflow, you should:

- A. Accept the 80% accuracy rate — it is commercially acceptable.
- B. Design the workflow to catch and handle the 2 wrong outputs.**
- C. Switch to a different AI tool with a higher accuracy rate.
- D. Run the prompt 3 times and take the majority answer.

Correct Answer: B

Explanation

Accepting 80% accuracy in a professional workflow means building in the assumption that 20% of outputs will fail — and those failures have consequences. Any deployed AI workflow must be designed to detect, handle, and recover from errors.

Q15. Which statement about AI tools is most useful to hold as a working principle?

- A. AI is a powerful tool that requires human judgment to use safely.**
- B. AI tools are either reliable or unreliable — know which is which.
- C. AI is becoming more reliable and will soon require less oversight.
- D. AI tools are only safe when used for their explicitly intended purpose.

Correct Answer: A

Explanation

The most useful working principle is one that is always true and actionable: AI is powerful but requires human judgment. This principle applies regardless of which tool, which task, or how the technology evolves.

SET

2

Working With AI — Practical Judgment

Set 2 of 5 · 15 Questions · AI Generalist Entry Diagnostic

Q16. You ask Claude to help you prepare for a difficult conversation with a colleague. The best prompt includes:

- A. Only the topic — Claude infers context from professional settings.
- B. The outcome you want — Claude will determine how to get there.
- C. The context, the relationship, the goal, and the constraints.**
- D. The colleague's behaviour — Claude can assess their personality.

Correct Answer: C

Explanation

Effective prompts for interpersonal preparation include all relevant context: what the situation is, who the other person is, what you are trying to achieve, and what constraints exist. Without this, Claude produces generic advice.

Q17. You are using Claude to help plan a complex project. The most valuable use is:

- A. Generating a first-pass plan you then stress-test against your knowledge.**
- B. Producing a final project plan the team can immediately follow.
- C. Replacing the need for a project manager on smaller projects.
- D. Creating the project timeline automatically from your calendar.

Correct Answer: A

Explanation

Claude's planning outputs are starting points — structurally sound, often comprehensive, but built on your prompt rather than deep knowledge of your specific context. Stress-testing against your expert knowledge is where the real value is created.

Q18. What is "prompt chaining" and why is it useful?

- A. Connecting multiple AI tools in sequence to process the same input.
- B. Saving a series of prompts to reuse across different conversations.
- C. Running the same prompt simultaneously across multiple models.
- D. Breaking complex tasks into sequential prompts where each builds on the last.**

Correct Answer: D

Explanation

Prompt chaining decomposes a complex task into manageable steps — each prompt takes the output of the previous one as input. This improves reliability and quality on tasks that are too complex or long for a single prompt to handle reliably.

Q19. A colleague shares a prompt that works well for them. You try it and get poor results. The most likely reason is:

- A. The prompt was overfit to one model and does not transfer.
- B. Your task or context differs from theirs — prompts are context-specific.**
- C. The prompt was shared incorrectly and has formatting errors.
- D. You need to run it several times to get a result matching theirs.

Correct Answer: B

Explanation

Prompts are highly context-dependent. A prompt that works for one person's task, role, and context may produce poor results when the underlying task or context differs — even if it looks identical. Effective prompting requires adapting to your specific situation.

Q20. You are reviewing a long Claude output. Which part should you scrutinise most carefully?

- A. The opening paragraph — AI models front-load errors.
- B. The conclusion — AI models drift from the brief toward the end.
- C. Specific figures, dates, names, and citations — highest hallucination risk.**
- D. Formatting and structure — these are the most common failure points.

Correct Answer: C

Explanation

Specific figures, dates, names, and citations carry the highest hallucination risk — they are precise, verifiable, and AI generates them with the same confidence as correct information. General reasoning is more reliable than specific factual details.

Q21. What does it mean to "engineer a prompt" rather than just "write a prompt"?

- A. Using technical coding skills to pass prompts through an API.
- B. Writing prompts in a structured format that AI tools can parse.
- C. Testing a prompt across multiple AI models before using it.
- D. Deliberately designing each element of the prompt for a specific outcome.**

Correct Answer: D

Explanation

Prompt engineering is a design discipline — deliberately constructing each element (role, context, task, constraints, format, examples) with a specific intended outcome in mind. It is not about technical delivery method or multi-model testing.

Q22. You use Claude to draft a report and it produces 80% of what you need. The right next step is:

- A. Edit, restructure, and complete the remaining 20% yourself.**
- B. Prompt Claude again to add the missing 20% in one step.
- C. Accept the 80% — it is production-ready with minor formatting.
- D. Ask Claude to identify what is missing from its own output.

Correct Answer: A

Explanation

The correct approach is to take Claude's 80% draft as material and complete it yourself — adding your domain knowledge, correcting inaccuracies, restructuring as needed. Asking Claude to patch itself on what it does not know produces more AI output, not better output.

Q23. Which sign indicates someone is using AI as a tool effectively rather than passively?

- A. They ask Claude for help on every task before attempting it themselves.
- B. They evaluate AI output critically before deciding whether to use it.**
- C. They use the most expensive AI model available for all tasks.
- D. They check AI outputs for formatting errors before submitting.

Correct Answer: B

Explanation

Effective tool use is characterised by critical evaluation — engaging judgment before accepting output. Relying on AI before attempting a task, using expensive models, or only checking formatting are patterns of passive use that produce lower quality outcomes.

Q24. You want to use Claude consistently across your team. The most important first step is:

- A. Ensure every team member has a Claude account and access.
- B. Identify which team member will manage Claude interactions.
- C. Select the Claude model tier that fits your budget.
- D. Define what tasks Claude will handle and establish quality standards.**

Correct Answer: D

Explanation

Consistent team AI use requires first defining scope (what tasks) and standards (what good looks like). Without this, different team members use AI differently, producing inconsistent quality and creating coordination failures.

Q25. What is the best mental model for AI in a professional context?

- A. An expert advisor whose recommendations can be trusted directly.
- B. A search tool that retrieves answers from a large knowledge base.
- C. A capable junior colleague who works fast but needs oversight.**
- D. A creative tool useful only for ideation and brainstorming.

Correct Answer: C

Explanation

"Capable junior colleague who needs oversight" is the most accurate working mental model — it acknowledges real capability (fast, broad, reasonable first drafts) while preserving the human responsibility to review, verify, and take ownership of final outputs.

Q26. You discover that a prompt that worked last week produces worse results today. The most likely cause is:

- A. Your internet connection affects Claude's response quality.
- B. The model may have been updated, changing its behaviour.**
- C. The prompt degraded because it was used too many times.
- D. Claude memorised your previous conversations and changed strategy.

Correct Answer: B

Explanation

AI model providers update and fine-tune models regularly. A prompt that relied on specific model behaviour may produce different results after an update. Prompts should be tested when model versions change.

Q27. What is the most appropriate response when Claude says "I'm not sure about this"?

- A. Treat the output as lower-confidence and verify the key claims.**
- B. Ask Claude the same question again for a more confident answer.
- C. Discard the output entirely — uncertainty means the answer is wrong.
- D. Accept the answer — Claude is unusually honest when uncertain.

Correct Answer: A

Explanation

Claude's expressed uncertainty is a signal to treat the output as lower-confidence — not to repeat the prompt or discard the output. Verify the specific uncertain claims independently. Note: Claude is not reliably calibrated on uncertainty and can be confident when wrong.

Q28. You need Claude to maintain a consistent tone across a long document. The best approach is:

- A. Generate the document in one prompt to maintain consistency.
- B. Ask Claude to "maintain consistent tone" throughout.
- C. Provide a style guide or tone example at the start of your prompt.**
- D. Generate each section separately and edit tone manually at the end.

Correct Answer: C

Explanation

Claude maintains tone most reliably when given explicit reference material — a style guide, sample paragraph, or defined tone descriptors — at the start of the prompt. A general instruction to "be consistent" is too vague to reliably produce uniformity across a long document.

Q29. Which approach produces the best AI-assisted research output?

- A. Use AI to find and summarise sources, then compile its summaries.
- B. Use AI exclusively since it has processed more research than any human.
- C. Use AI only to format and structure research you have already done.
- D. Use AI to generate hypotheses and structures, verify facts from primary sources.**

Correct Answer: D

Explanation

The most effective approach combines AI's speed in generating structure and hypotheses with human discipline in verifying facts from primary sources. AI summaries of sources are unreliable; AI-only research bypasses the human judgment that academic research requires.

Q30. A generalist should be able to explain why AI sometimes fails on a task. Which explanation is most accurate?

- A. The model is not large enough for that task type.
- B. The task requires precision or real-time data outside AI's design strengths.**
- C. The task was too complex for any current AI system to handle.
- D. Insufficient training data exists for that specific task.

Correct Answer: B

Explanation

Most professional AI failures occur because the task requires something outside AI's strengths — precise arithmetic, current information, legally binding judgment, or verified sourcing. Understanding this allows a generalist to route tasks appropriately from the start.

SET

3

Prompting for Professional Output

Set 3 of 5 · 15 Questions · AI Generalist Entry Diagnostic

Q31. You want Claude to write a client email that sounds like you. The most effective prompt element is:

- A. A sample of your previous emails as a tone reference.
- B. A list of words you prefer and words to avoid.
- C. An instruction to write in first person and sound natural.
- D. Your name and job title for personalisation.

Correct Answer: A

Explanation

Tone is best conveyed through examples, not descriptions. A sample of your actual previous writing gives Claude far more signal about your voice — sentence length, hedging style, formality, directness — than any description of those qualities.

Q32. You ask Claude to write a 500-word summary. It produces 750 words. The most effective fix is:

- A. "Make this shorter."
- B. "Cut 250 words from this summary."
- C. "Rewrite this as exactly 500 words. Do not add new points — compress the existing ones."
- D. "Summarise this summary in 500 words."

Correct Answer: C

Explanation

The most effective instruction is specific and mechanically clear: target length (500 words), method (compress existing content), and constraint (no new points). Vague length instructions like "make it shorter" rarely hit the target.

Q33. What is the most reliable way to get Claude to produce a structured output format every time?

- A. Use the phrase "always follow this format" in the instructions.
- B. Provide the exact template or schema you want in the prompt.**
- C. Ask Claude to review its output for format compliance.
- D. Generate multiple outputs and select the one with the right format.

Correct Answer: B

Explanation

The most reliable structural control is providing the exact template — showing Claude what the output should look like rather than describing it. Templates are more precise than format descriptions and produce more consistent results.

Q34. You want Claude to generate 10 marketing taglines for your product. You keep getting similar options. The fix is:

- A. Ask for 10 more taglines and select from the combined 20.
- B. Switch to a more creative AI tool built for copywriting tasks.
- C. Run the same prompt 5 times and take the best from each.
- D. Add constraints that force variety — different emotions, audiences, or lengths.**

Correct Answer: D

Explanation

Sameness in creative outputs comes from insufficient constraints. Adding variety constraints — one tagline per target emotion, one per audience segment, one short and punchy — forces Claude to explore different angles rather than converging on the most probable options.

Q35. Which prompt technique most improves the accuracy of factual AI outputs?

- A. Ask Claude to cite evidence or reasoning for each claim it makes.**
- B. Ask Claude to rate its own confidence on a 1-10 scale.
- C. Ask Claude to search the internet for supporting evidence.
- D. Ask Claude to list what it does not know before answering.

Correct Answer: A

Explanation

Asking Claude to provide evidence or reasoning for each claim surfaces the basis of its statements — making it easier to identify which claims are well-grounded and which are assumptions or hallucinations. Self-confidence ratings are unreliable.

Q36. You want to use Claude to prepare an executive summary. The summary should contain no information that is not in the source document. The best instruction is:

- A. "Write an accurate executive summary of this document."
- B. "Summarise only information explicitly present in this document. Do not add external context."**
- C. "Summarise this document without adding opinions."
- D. "Do not hallucinate when summarising this document."

Correct Answer: B

Explanation

Explicit grounding instructions — "only information present in this document" and "do not add external context" — directly constrain Claude's output to the source material. Generic accuracy instructions and hallucination prohibitions are too vague to reliably constrain output.

Q37. What does "few-shot prompting" mean?

- A. Running the same prompt a few times and selecting the best output.
- B. Using a short, concise prompt rather than a long detailed one.
- C. Giving Claude a few tries to produce an acceptable output.
- D. Including examples of desired inputs and outputs in the prompt.**

Correct Answer: D

Explanation

Few-shot prompting means providing example input-output pairs in the prompt — showing Claude what a correct response looks like before giving it the actual task. Examples are generally more reliable than descriptions for defining desired output format and quality.

Q38. A client asks you to produce a document using AI. You should:

- A. Disclose that AI produced it and let the client decide whether to use it.
- B. Only use AI if the client has explicitly approved AI use.
- C. Use AI to draft it, then review, verify, and own the final output.**
- D. Not use AI for client deliverables — only for internal work.

Correct Answer: C

Explanation

Using AI to draft, then reviewing, verifying, and owning the final output is the professional standard. The AI is a production tool; you are the accountable professional. Blanket policies requiring prior approval or prohibiting client work are overly restrictive given how widespread AI use now is.

Q39. You use Claude to write a proposal and it sounds generically professional but lacks your company's specific value proposition. The cause is:

- A. Claude cannot write specific business proposals — only general ones.
- B. The prompt did not include your company's differentiated positioning.**
- C. The proposal genre is too structured for AI to handle effectively.
- D. Claude's training data does not include specialised business writing.

Correct Answer: B

Explanation

Generic output reflects generic input. Claude has no knowledge of your company's specific positioning, differentiators, or language unless you provide it. The fix is to include that positioning in the prompt — not to change tools.

Q40. Which of these represents the most advanced prompting behaviour?

- A. Iterating on output by providing specific, targeted improvement instructions.**
- B. Writing the longest, most detailed possible initial prompt.
- C. Using multiple AI tools on the same task and comparing outputs.
- D. Memorising the best prompts for each task type in your work.

Correct Answer: A

Explanation

The most advanced behaviour is iteration — the ability to evaluate output, identify exactly what needs improvement, and provide targeted instructions to fix it. Length, multi-tool comparison, and memorisation are all lower-leverage than skilled iteration.

Q41. You need Claude to maintain the same tone and approach across 20 different outputs. The most efficient approach is:

- A. Copy-paste the best previous output as context for each new prompt.
- B. Review and manually align tone across all 20 outputs after generation.
- C. Create a reusable system prompt that defines tone, role, and constraints.**
- D. Generate all 20 in one prompt with consistent instructions.

Correct Answer: C

Explanation

A reusable system prompt — defining role, tone, style, and constraints once — is the most efficient approach for consistent multi-output production. It applies the same context to each generation without requiring repetition or manual alignment.

Q42. What is the clearest signal that an AI output requires significant human intervention before use?

- A. It is longer than what you asked for.
- B. It does not match the exact format you requested.
- C. It sounds too polished and formal for the context.
- D. It contains specific figures or claims that are hard to verify quickly.**

Correct Answer: D

Explanation

Specific figures, dates, statistics, or citations that are hard to verify quickly are the clearest signal requiring human intervention — because these are exactly where hallucination risk is highest and where errors have the most consequential impact.

Q43. A new team member asks: "What is the most important thing to know about working with Claude?" The best answer is:

- A. "Use it for creative tasks but avoid it for factual research."
- B. "Claude generates plausible text — verify anything that matters."**
- C. "The more detailed your prompt, the better the output."
- D. "Claude is reliable for most professional tasks if you use it correctly."

Correct Answer: B

Explanation

The single most important principle for a new Claude user is: outputs are plausible-sounding text, not verified facts. This calibrates their trust appropriately and establishes the verification habit that makes all other AI use safe.

Q44. You want Claude to generate ideas that are genuinely different from its first response. The best approach is:

- A. Increase the temperature setting and regenerate the same prompt.
- B. Use a different AI tool which may have different training data.
- C. Ask for ideas that explicitly avoid the approaches in the first response.**
- D. Ask Claude to evaluate and critique its first response before generating more.

Correct Answer: C

Explanation

Explicitly instructing Claude to avoid the approaches in its first response forces it to explore a different part of the idea space. Increasing temperature increases randomness but may produce more of the same type. Critique-then-generate adds a step but does not guarantee different directions.

Q45. Which professional task is least suited to AI assistance right now?

- A. Making a final legal judgment in a contested regulatory dispute.**
- B. Drafting a first version of a legal argument for attorney review.
- C. Summarising case precedents from a set of provided documents.
- D. Generating a list of research questions for a legal brief.

Correct Answer: A

Explanation

Final legal judgment requires licensed professional accountability, deep contextual judgment, and real consequences. All the other tasks produce draft output for human review — an appropriate use case for AI that keeps the professional accountable.

SET

4

AI Tools and Workflows

Set 4 of 5 · 15 Questions · AI Generalist Entry Diagnostic

Q46. What is the correct order of operations when using AI to analyse a dataset?

- A. Upload the data, run AI analysis, and present the outputs directly.
- B. Clean the data with AI first, then define the question, then analyse.
- C. Run exploratory AI analysis first to discover what questions to ask.
- D. Define the question first, then use AI to analyse, then verify key outputs.**

Correct Answer: D**Explanation**

The correct order starts with a defined question — AI analysis without a clear question produces unfocused outputs. After analysis, key outputs should be verified before use. Presenting raw AI analysis directly skips critical verification.

Q47. Which combination of tools is most powerful for a knowledge worker using AI?

- A. Claude plus a second AI tool to cross-check Claude's outputs.
- B. Claude for generation and synthesis, plus a primary source for verification.**
- C. Claude plus a premium plan for guaranteed higher accuracy.
- D. Claude plus an AI writing checker to validate quality before use.

Correct Answer: B**Explanation**

The most powerful combination is generative AI (for speed and synthesis) paired with authoritative primary sources (for verification). Cross-checking with another AI uses the same structural weakness; premium plans do not guarantee accuracy; AI quality checkers are themselves AI.

Q48. You are evaluating whether to use AI for a specific task at work. The most important question is:

- A. "Which AI tool has the highest accuracy rating for this task type?"
- B. "Is this task in the list of approved AI use cases at my company?"
- C. "What happens when AI gets this wrong and how will I catch it?"**
- D. "How much time will AI save me compared to doing it manually?"

Correct Answer: C

Explanation

The most important evaluative question is about failure mode — what the error looks like and how it will be caught. Time savings and approved use cases matter, but they are secondary to understanding what goes wrong and whether you have a way to detect it.

Q49. What does a well-designed AI workflow always include?

- A. A human checkpoint to review and validate AI outputs before they matter.**
- B. A fallback to a different AI tool if the primary one fails.
- C. A manual alternative process in case the AI is unavailable.
- D. A logging system that records all AI inputs and outputs.

Correct Answer: A

Explanation

Every effective AI workflow includes human checkpoints before outputs are used for consequential decisions. Technical redundancy (fallback tools, manual alternatives, logging) is useful but secondary to the fundamental human-in-the-loop design principle.

Q50. You are building an AI-assisted customer service workflow. The most important metric to track is:

- A. The average response time of the AI compared to human agents.
- B. The percentage of queries handled without human intervention.
- C. The cost per query with AI versus without AI.
- D. The rate of customer complaints or escalations from AI responses.**

Correct Answer: D

Explanation

Complaint and escalation rate measures actual customer harm from AI errors — the most important outcome. Response time, automation rate, and cost are efficiency metrics that matter only if quality is maintained.

Q51. Which workflow task is most appropriate to delegate entirely to Claude without human review?

- A. Drafting a response to a customer complaint about a billing error.
- B. Summarising a competitor's product features from a specification.
- C. Formatting a meeting transcript into a consistent structure.**
- D. Generating a project risk assessment for a client proposal.

Correct Answer: C

Explanation

Formatting a transcript into a consistent structure is a low-stakes structural task where errors are immediately visible and easily corrected. The other tasks involve AI making interpretive or evaluative judgments in higher-stakes contexts requiring human review.

Q52. What is the most important reason to document the prompts you use for professional work?

- A. So you can reproduce reliable outputs and improve them over time.**
- B. So you can prove to clients that AI was used professionally.
- C. So you can train your own AI model on effective prompt patterns.
- D. So you comply with your company's AI usage logging policy.

Correct Answer: A

Explanation

Documenting prompts enables reproducibility — you can reliably produce the same quality output again — and improvement over time through iteration. This is the core operational reason; compliance and proof of use are secondary considerations.

Q53. You use Claude to generate 5 customer persona profiles. Before using them, you should:

- A. Ask Claude to check whether the personas are realistic.
- B. Validate them against real customer data or research.**
- C. Add your team's subjective preferences to each persona.
- D. Use the personas directly — AI personas are built from broad patterns.

Correct Answer: B

Explanation

AI-generated personas are built from statistical patterns in training data, not your actual customer data. Before using them professionally, validate them against real research, interviews, or data about your specific customer base.

Q54. What is the key difference between using AI as a "co-pilot" versus using AI as an "autopilot"?

- A. Co-pilot is for creative tasks; autopilot is for repetitive tasks.
- B. Autopilot is more reliable than co-pilot for professional output.
- C. Co-pilot means you remain in control; autopilot means AI makes decisions.**
- D. Co-pilot requires technical skills; autopilot does not.

Correct Answer: C

Explanation

Co-pilot means human judgment remains in the loop — AI assists while the human steers, reviews, and decides. Autopilot means AI operates without meaningful human oversight. In professional contexts, most tasks require co-pilot, not autopilot, given current AI limitations.

Q55. You use AI to generate a client-facing presentation. The most important pre-send check is:

- A. Ensure the formatting and visual design meet brand standards.
- B. Confirm the length and structure match the client's preferences.
- C. Check that the language is appropriately formal for the client.
- D. Verify every factual claim, figure, and recommendation for accuracy.**

Correct Answer: D

Explanation

Factual accuracy is the primary check before any AI output reaches a client. Brand standards, length, structure, and language are important but secondary — inaccurate information reaching a client creates professional and reputational risk.

Q56. What is the most effective way to use AI for decision-making support?

- A. Use AI to map options and tradeoffs, then apply your own judgment.**
- B. Use AI to recommend the best decision based on your inputs.
- C. Use AI to validate your instinct before taking action.
- D. Use AI to identify the option with the highest probability of success.

Correct Answer: A

Explanation

AI is most effective as a structured thinking tool — mapping out options, articulating tradeoffs, and surfacing considerations. The final decision and its accountability must remain with the human, who has context, values, and responsibility AI cannot replicate.

Q57. You are using Claude in a meeting to take live notes and action items. The most important post-meeting step is:

- A. Share the notes immediately — AI is faster than human review.
- B. Review and correct the notes before sharing with attendees.**
- C. Ask Claude to clean up and finalise the notes before sending.
- D. Send the raw notes and ask attendees to correct their own action items.

Correct Answer: B

Explanation

Live AI note-taking is valuable for speed but requires human review before distribution — names, action items, decisions, and context can all be captured incorrectly. Sharing raw AI notes without review risks distributing errors to the whole team.

Q58. Which of these represents the highest-maturity level of AI use in a professional context?

- A. Being faster than colleagues at completing tasks using AI tools.
- B. Using AI for every task in your daily workflow without exception.
- C. Knowing which AI tool to use for each category of work task.
- D. Designing AI-enabled systems that others in the team can use reliably.**

Correct Answer: D

Explanation

The highest maturity level is systemic — not just using AI yourself but designing workflows, prompts, and guardrails that allow others to use AI reliably and safely. Speed, breadth of use, and tool selection are earlier stages of individual AI capability.

Q59. What should a professional do when AI generates a response that confidently contradicts their domain expertise?

- A. Investigate carefully — the AI may be wrong, or they may have a gap.**
- B. Trust their expertise over AI — human domain knowledge is more reliable.
- C. Trust the AI — it has processed more information on the topic.
- D. Run the same query through a different AI to resolve the conflict.

Correct Answer: A

Explanation

When AI confidently contradicts domain expertise, the correct response is careful investigation — not automatic deference to expertise or to AI. The AI may have identified a genuine gap, or it may be hallucinating. Only investigation determines which.

Q60. What is "prompt drift" and why does it matter?

- A. AI models produce lower quality outputs after extended use.
- B. Prompts become less effective when used by many different people.
- C. Output quality degrades over a long conversation as context accumulates.**
- D. AI responses drift toward generic outputs when context is too brief.

Correct Answer: C

Explanation

Prompt drift occurs in long conversations as the context window fills — earlier instructions get pushed further back and receive less attention, causing the model to drift from its original briefing. Refreshing instructions periodically prevents this in extended workflows.

SET

5

AI Judgment — Applied Scenarios

Set 5 of 5 · 15 Questions · AI Generalist Entry Diagnostic

Q61. Your manager forwards an impressive AI-generated report. They ask you to build on it. Your first step is:

- A. Verify the key claims and figures before treating the report as accurate.**
- B. Use Claude to expand and build on the existing AI-generated content.
- C. Accept the report as accurate — your manager reviewed it.
- D. Identify which AI tool generated the report to assess its reliability.

Correct Answer: A

Explanation

AI-generated reports may contain hallucinated facts, distorted analysis, or unverified claims. Before building on them, verify the key inputs. Your manager reviewing it does not guarantee accuracy — they may not have verified either.

Q62. A colleague builds an AI chatbot for your team's internal use without telling anyone. The key concern is:

- A. The chatbot may be slower than commercial AI tools the team uses.
- B. The chatbot will not be as capable as enterprise AI solutions.
- C. Using an internal chatbot may violate AI tool licensing terms.
- D. Sensitive internal data shared with it may not be appropriately secured.**

Correct Answer: D

Explanation

Unvetted internal AI deployments create data security and governance risks — sensitive information shared with the chatbot may be stored, logged, or exposed in ways that violate data policies. Capability and licensing concerns are secondary.

Q63. You receive a task from a client: "Use AI to do X." The client does not know AI limitations. Your responsibility is to:

- A. Tell the client that AI cannot reliably do this task as described.
- B. Complete the task using AI appropriately and deliver verified output.**
- C. Do the task manually since the client doesn't understand AI well.
- D. Ask the client to specify exactly which AI tool they want you to use.

Correct Answer: B

Explanation

Your professional responsibility is to use AI to the extent it genuinely helps, verify the output, and deliver quality work. Educating the client about limitations is appropriate, but not as a reason to avoid the task — bring your professional judgment to bear on how AI is used.

Q64. You use AI to generate 10 ideas for a strategy session. You share them without disclosure. What is the risk?

- A. Someone on the team will identify the ideas as AI-generated.
- B. The ideas will be lower quality than human-generated ones.
- C. The team treats the ideas as human-generated strategic thinking.**
- D. The client will ask which team member generated each idea.

Correct Answer: C

Explanation

Undisclosed AI-generated ideas in strategic discussion are treated as human strategic judgment — which they are not. This distorts the team's confidence in the ideas, bypasses human creative engagement, and creates accountability ambiguity.

Q65. You are working on a time-sensitive task and Claude produces an output with obvious errors. You should:

- A. Fix the errors yourself and complete the task — do not resubmit to Claude.**
- B. Ask Claude to fix its own errors before you review the output.
- C. Start over with a better prompt to avoid getting errors again.
- D. Submit the task with the errors flagged as "pending verification."

Correct Answer: A

Explanation

When you can see the errors, fix them yourself — you know what is wrong and can correct it faster than prompting Claude again and reviewing another output. Asking Claude to fix its own errors is unreliable; starting over wastes time when you already have 80% of what you need.

Q66. A startup asks you to evaluate their AI product. Which question reveals the most about real-world reliability?

- A. "What accuracy rate does it achieve on benchmark tests?"
- B. "Which underlying AI model does it use?"
- C. "What does it do when it gets the answer wrong?"**
- D. "How many customers are currently using it?"

Correct Answer: C

Explanation

How a product handles failure is the most revealing reliability question. Benchmark accuracy measures best-case performance in controlled conditions; the real test is whether the system degrades gracefully, catches errors, and allows recovery when something goes wrong.

Q67. You use Claude to analyse a competitor's strategy from public information. The output is insightful. Before presenting it:

- A. Present it directly — public information analysis is low-risk.
- B. Have Claude revise the analysis for clarity before presenting.
- C. Identify which data points came from Claude's training versus the document.
- D. Verify the key facts about the competitor from primary sources.**

Correct Answer: D

Explanation

Even when AI analysis feels insightful, the specific facts about the competitor — market share figures, product details, strategic moves — must be verified from primary sources. AI can misattribute, confuse, or hallucinate competitor-specific details.

Q68. A junior team member says "Claude told me we should use approach X for this project." The correct response is:

- A. "If Claude recommended it, we should take it seriously."
- B. "Walk me through the reasoning — Claude's suggestions need human judgment."**
- C. "We need to ask another AI tool to validate Claude's recommendation."
- D. "Claude cannot make recommendations on technical project decisions."

Correct Answer: B

Explanation

Claude's suggestions should be inputs to human judgment, not directives. Walking through the reasoning allows the team to evaluate whether the suggestion makes sense in context — which is exactly the human judgment that makes AI useful and safe.

Q69. You use Claude to draft communications for a sensitive HR situation. The most important consideration is:

- A. The communication must be reviewed by the CEO before sending.
- B. You should disclose to the employee that AI drafted the communication.
- C. Legal and procedural accuracy must be verified with HR and legal counsel.**
- D. Claude should not be used for any HR-related communications.

Correct Answer: C

Explanation

HR communications in sensitive situations have legal, procedural, and human implications that AI cannot reliably get right. Verification with HR and legal counsel is essential regardless of how well the draft reads.

Q70. Which of the following best describes an "AI-first" professional approach?

- A. Considering how AI can enhance every task before defaulting to manual work.**
- B. Using AI for every task without exception to maximise efficiency.
- C. Building AI tools before deciding what problem to solve.
- D. Replacing all manual processes with AI-powered alternatives.

Correct Answer: A

Explanation

AI-first means considering AI as an option from the start of every task — asking "how can AI enhance this?" before defaulting to a fully manual approach. It is not using AI for everything regardless of fit, nor is it building tools before problems.

Q71. You delegate a report to Claude and the deadline arrives. The report has errors. Who is accountable?

- A. Claude is — it produced the incorrect content.
- B. You are — you chose to delegate to AI and are responsible for the output.**
- C. The manager is — they set the deadline without allowing review time.
- D. No one — AI errors are unforeseeable and cannot be assigned.

Correct Answer: B

Explanation

Accountability for AI output always rests with the professional who used it. Delegating to Claude does not transfer accountability. The lesson is that AI delegation requires building in review time before the deadline.

Q72. A client insists on a deliverable that AI cannot reliably produce accurately. You should:

- A. Use AI anyway and review carefully to catch all errors.
- B. Decline the deliverable — professional standards prohibit unreliable AI use.
- C. Outsource the deliverable to a specialist who can do it without AI.
- D. Explain the limitation and propose a hybrid approach with AI drafting and human verification.**

Correct Answer: D

Explanation

The professional response is to be transparent about AI limitations and propose a workable hybrid: AI for speed and structure, human expertise for verification and quality assurance. This serves the client better than blind AI use or unnecessary refusal.

Q73. You are asked to train colleagues on AI best practices. The most important lesson to start with is:

- A. "Always use the most advanced AI model available for any task."
- B. "AI tools are most useful for creative and ideation tasks."
- C. "AI outputs are starting points — own the final output yourself."**
- D. "AI can be trusted for most tasks once you have the right prompt."

Correct Answer: C

Explanation

The most foundational lesson — and the one with the most immediate impact on how people use AI safely — is ownership of final output. Everything else (tool selection, prompting technique, task suitability) builds on this foundation.

Q74. What is the best indicator that a professional is ready to lead AI adoption in their organisation?

- A. They have completed an AI certification from a recognised provider.
- B. They can explain both AI's value and its limitations with equal accuracy.**
- C. They have built at least one working AI tool using an API.
- D. They use AI every day across all their professional tasks.

Correct Answer: B

Explanation

Leadership in AI adoption requires nuanced understanding — knowing what AI can genuinely do and where it genuinely fails, and communicating both accurately. Certification, tool-building, and daily use are valuable signals but not sufficient without this balanced judgment.

Q75. What does Menler's AI Generalist Entry Diagnostic assess above all else?

- A. Whether a learner has used AI tools for more than six months.
- B. Whether a learner can build AI tools using Claude's API.
- C. Whether a learner has prior work experience using AI professionally.
- D. Whether a learner's AI judgment and instincts are ready for fellowship training.**

Correct Answer: D

Explanation

The Entry Diagnostic assesses readiness for fellowship training — specifically whether the learner has the judgment, critical instincts, and practical AI understanding needed to benefit from and contribute to the program. Tool experience, tenure, and technical skills are secondary.



Your turning point in the AI era.

END OF QUESTION BANK

*Your turning
point starts
here.*