

menler

Your turning point in the AI era.

AI APTITUDE QUESTION BANK

AI Generalist

About This Question Bank

A complete AI Generalist assessment covering 150 questions across 10 sets. Sets 1–5 cover entry-level fundamentals, practical judgment, prompting, workflows, and applied scenarios. Sets 6–10 advance into strategy, agentic AI, output quality evaluation, and AI leadership.

10

Sets

150

Total Questions

SET GUIDE

Set 1**Entry — AI Fundamentals You Must Know***AI fundamentals — how models work, key concepts, and what you must know***Set 2****Entry — Working With AI: Practical Judgment***Working with AI — practical judgment, mental models, and team use***Set 3****Entry — Prompting for Professional Output***Prompting for professional output — specificity, templates, and iteration***Set 4****Entry — AI Tools and Workflows***AI tools and workflows — designing reliable AI-assisted processes***Set 5****Entry — AI Judgment: Applied Scenarios***AI judgment — applied scenarios, accountability, and edge cases***Set 6****Level Up — Advanced Prompting and Output Quality***Advanced prompting and output quality — specification, control, and iteration***Set 7****Level Up — AI Strategy: Where AI Adds Value***AI strategy — where AI adds value and how to direct adoption***Set 8****Level Up — Agentic AI and Workflow Automation***Agentic AI and workflow automation — agents, tools, and safe design***Set 9****Level Up — Evaluating AI Output Quality***Evaluating AI output quality — fact-checking, calibration, and review***Set 10****Level Up — AI Leadership and Cultural Change***AI leadership and cultural change — teams, governance, and transformation*

Each question has one correct answer. Select the best answer from the options provided.

Correct answers and explanations follow each question.

SET

1

Entry — AI Fundamentals You Must Know

Set 1 of 10 · 15 Questions · AI Generalist

Q1. You receive an AI output that contains a confidently stated incorrect fact. The correct framing is:

- A. The AI model version you used is outdated — upgrade it.
- B. AI outputs are starting points that require human verification.**
- C. AI is unreliable and should not be used for factual tasks.
- D. You prompted incorrectly — a better prompt would prevent this.

Correct Answer: B

Explanation

Confident errors (hallucinations) are a structural property of how language models work — not a model version issue or prompting failure. The correct framing is that AI produces starting points, and humans are responsible for verification.

Q2. A colleague says "I don't trust AI — it makes things up." The most accurate and constructive response is:

- A. "You're right — AI is unreliable for professional use."
- B. "AI is reliable as long as you use the right tool for the task."
- C. "That concern is outdated — modern AI is much more accurate."
- D. AI has known failure modes but also genuine strengths worth understanding.**

Correct Answer: D

Explanation

The most accurate position acknowledges real limitations without dismissing real capability. Blanket rejection and blanket acceptance are both inaccurate. Understanding the specific failure modes and strengths enables effective, responsible use.

Q3. What distinguishes a generative AI model from a traditional software program?

- A. It produces novel outputs from statistical patterns, not explicit rules.**
- B. It can access the internet to retrieve up-to-date information.
- C. It runs faster than traditional programs on the same hardware.
- D. It makes decisions without any human input or instruction.

Correct Answer: A

Explanation

Traditional programs execute explicit rules written by developers. Generative AI learns statistical patterns from data and generates outputs — text, images, code — that are not explicitly programmed. This makes outputs probabilistic rather than deterministic.

Q4. Which of the following most accurately describes what Claude "knows"?

- A. A verified encyclopaedia of facts maintained by Anthropic.
- B. Real-time information retrieved from the internet on demand.
- C. Statistical patterns from training data that approximate knowledge.**
- D. Expert knowledge contributed by domain specialists.

Correct Answer: C

Explanation

Claude's knowledge is statistical — weighted patterns learned from training data. It is not a verified fact store, real-time retrieval system, or curated expert database. This explains both its breadth and its hallucination tendency.

Q5. A prompt engineer says "prompt quality determines output quality." How accurate is this?

- A. Inaccurate — model size matters far more than prompt quality.
- B. Largely accurate — prompt quality is the highest-leverage input.**
- C. Partially accurate — only for creative tasks, not factual ones.
- D. Inaccurate — all prompts to a good model produce good outputs.

Correct Answer: B

Explanation

Prompt quality is genuinely the highest-leverage variable a user controls. Given the same model, a well-specified prompt with context, role, format, and constraints consistently outperforms a vague one on nearly every task type.

Q6. What is the most important practical difference between Claude and a search engine?

A. Claude synthesises and generates; a search engine retrieves and ranks.

B. Claude is more accurate for factual queries than a search engine.

C. Claude provides citations while a search engine only shows links.

D. Claude is faster than a search engine for complex research tasks.

Correct Answer: A

Explanation

Search engines index, retrieve, and rank existing pages. Claude synthesises, interprets, and generates new text from its learned patterns. Neither provides citations by default. For current facts, search engines are more reliable; for synthesis, Claude adds more value.

Q7. You are building a workflow that uses Claude to process customer emails. Which risk must you design for first?

A. Claude may run too slowly for high-volume email processing.

B. Claude may respond in a language the customer did not use.

C. Claude may generate responses that are too long for emails.

D. Claude may misclassify or misrespond to some customer emails.

Correct Answer: D

Explanation

Misclassification and incorrect responses are the primary risk in any customer-facing AI system — the consequences include customer harm, trust loss, and operational errors. Speed, language, and length are secondary engineering concerns.

Q8. Why is it important to understand what a model's context window is before building with it?

A. It determines which languages the model can respond in.

B. It determines the maximum accuracy of the model's outputs.

C. It determines how much input and conversation the model can process.

D. It determines which API endpoints are available to developers.

Correct Answer: C

Explanation

The context window is the maximum tokens the model can process in one call. Building a workflow without knowing this limit leads to unexpected truncation, lost context, or errors when inputs exceed capacity.

Q9. A manager asks you to "automate this process using AI." The correct first step is:

- A. Map the process and identify which steps AI can actually handle.**
- B. Select the AI tool and start building the automation immediately.
- C. Ask the manager which AI tool they want you to use.
- D. Estimate the cost savings before designing the solution.

Correct Answer: A

Explanation

Effective AI automation starts with process mapping — understanding each step, its inputs, outputs, and failure modes — before selecting tools or building. Jumping to implementation without this creates solutions that fail in unexpected ways.

Q10. What does "temperature = 0" mean in practice when using a language model API?

- A. The model refuses to generate any creative or subjective content.
- B. The model runs at reduced speed to improve accuracy.
- C. The model caps response length at a conservative limit.
- D. The model produces the most probable, deterministic output.**

Correct Answer: D

Explanation

Temperature = 0 means the model always selects the highest-probability next token — producing maximally deterministic, repeatable output. This is useful for structured tasks (classification, extraction) where consistency is more important than variety.

Q11. You need Claude to output structured JSON. The most reliable approach is:

- A. Ask Claude to "output JSON" without additional instructions.
- B. Specify the exact JSON schema in the prompt and use low temperature.**
- C. Use Claude only if it has a dedicated JSON output mode.
- D. Generate the output and manually convert it to JSON afterward.

Correct Answer: B

Explanation

Reliable structured output requires both explicit schema specification (what fields, what types, what structure) and low temperature (to reduce randomness). Without these, output format is unreliable even when Claude attempts JSON.

Q12. What is the most common mistake people make when first using Claude for work tasks?

- A. Over-specifying the task with too much context and detail.
- B. Using Claude for tasks that require real-time information.
- C. Giving vague instructions and expecting Claude to infer the details.**
- D. Relying on Claude before verifying it can handle the task type.

Correct Answer: C

Explanation

The most common failure is vague prompting — treating Claude like a search engine or a smart person who can fill in the gaps. Claude performs best with explicit context, clear role definition, specific format requirements, and precise constraints.

Q13. An AI Generalist's core skill is best described as:

- A. Being able to code AI models and train them on custom data.
- B. Mastering every AI tool available in their professional domain.
- C. Understanding the mathematical foundations of machine learning.
- D. Knowing how to direct AI across varied tasks to produce useful outputs.**

Correct Answer: D

Explanation

An AI Generalist's core skill is judgment in directing AI — knowing what to ask, how to specify it, how to evaluate the output, and when to verify or escalate. This is different from technical ML skills or comprehensive tool mastery.

Q14. You run the same Claude prompt 10 times. 8 of 10 responses are correct and 2 are wrong. For a professional workflow, you should:

- A. Accept the 80% accuracy rate — it is commercially acceptable.
- B. Design the workflow to catch and handle the 2 wrong outputs.**
- C. Switch to a different AI tool with a higher accuracy rate.
- D. Run the prompt 3 times and take the majority answer.

Correct Answer: B

Explanation

Accepting 80% accuracy in a professional workflow means building in the assumption that 20% of outputs will fail — and those failures have consequences. Any deployed AI workflow must be designed to detect, handle, and recover from errors.

Q15. Which statement about AI tools is most useful to hold as a working principle?

- A. AI is a powerful tool that requires human judgment to use safely.**
- B. AI tools are either reliable or unreliable — know which is which.
- C. AI is becoming more reliable and will soon require less oversight.
- D. AI tools are only safe when used for their explicitly intended purpose.

Correct Answer: A

Explanation

The most useful working principle is one that is always true and actionable: AI is powerful but requires human judgment. This principle applies regardless of which tool, which task, or how the technology evolves.

SET

2

Entry — Working With AI: Practical Judgment

Set 2 of 10 · 15 Questions · AI Generalist

Q16. You ask Claude to help you prepare for a difficult conversation with a colleague. The best prompt includes:

- A. Only the topic — Claude infers context from professional settings.
- B. The outcome you want — Claude will determine how to get there.
- C. The context, the relationship, the goal, and the constraints.**
- D. The colleague's behaviour — Claude can assess their personality.

Correct Answer: C

Explanation

Effective prompts for interpersonal preparation include all relevant context: what the situation is, who the other person is, what you are trying to achieve, and what constraints exist. Without this, Claude produces generic advice.

Q17. You are using Claude to help plan a complex project. The most valuable use is:

- A. Generating a first-pass plan you then stress-test against your knowledge.**
- B. Producing a final project plan the team can immediately follow.
- C. Replacing the need for a project manager on smaller projects.
- D. Creating the project timeline automatically from your calendar.

Correct Answer: A

Explanation

Claude's planning outputs are starting points — structurally sound, often comprehensive, but built on your prompt rather than deep knowledge of your specific context. Stress-testing against your expert knowledge is where the real value is created.

Q18. What is "prompt chaining" and why is it useful?

- A. Connecting multiple AI tools in sequence to process the same input.
- B. Saving a series of prompts to reuse across different conversations.
- C. Running the same prompt simultaneously across multiple models.
- D. Breaking complex tasks into sequential prompts where each builds on the last.**

Correct Answer: D

Explanation

Prompt chaining decomposes a complex task into manageable steps — each prompt takes the output of the previous one as input. This improves reliability and quality on tasks that are too complex or long for a single prompt to handle reliably.

Q19. A colleague shares a prompt that works well for them. You try it and get poor results. The most likely reason is:

- A. The prompt was overfit to one model and does not transfer.
- B. Your task or context differs from theirs — prompts are context-specific.**
- C. The prompt was shared incorrectly and has formatting errors.
- D. You need to run it several times to get a result matching theirs.

Correct Answer: B

Explanation

Prompts are highly context-dependent. A prompt that works for one person's task, role, and context may produce poor results when the underlying task or context differs — even if it looks identical. Effective prompting requires adapting to your specific situation.

Q20. You are reviewing a long Claude output. Which part should you scrutinise most carefully?

- A. The opening paragraph — AI models front-load errors.
- B. The conclusion — AI models drift from the brief toward the end.
- C. Specific figures, dates, names, and citations — highest hallucination risk.**
- D. Formatting and structure — these are the most common failure points.

Correct Answer: C

Explanation

Specific figures, dates, names, and citations carry the highest hallucination risk — they are precise, verifiable, and AI generates them with the same confidence as correct information. General reasoning is more reliable than specific factual details.

Q21. What does it mean to "engineer a prompt" rather than just "write a prompt"?

- A. Using technical coding skills to pass prompts through an API.
- B. Writing prompts in a structured format that AI tools can parse.
- C. Testing a prompt across multiple AI models before using it.
- D. Deliberately designing each element of the prompt for a specific outcome.**

Correct Answer: D

Explanation

Prompt engineering is a design discipline — deliberately constructing each element (role, context, task, constraints, format, examples) with a specific intended outcome in mind. It is not about technical delivery method or multi-model testing.

Q22. You use Claude to draft a report and it produces 80% of what you need. The right next step is:

- A. Edit, restructure, and complete the remaining 20% yourself.**
- B. Prompt Claude again to add the missing 20% in one step.
- C. Accept the 80% — it is production-ready with minor formatting.
- D. Ask Claude to identify what is missing from its own output.

Correct Answer: A

Explanation

The correct approach is to take Claude's 80% draft as material and complete it yourself — adding your domain knowledge, correcting inaccuracies, restructuring as needed. Asking Claude to patch itself on what it does not know produces more AI output, not better output.

Q23. Which sign indicates someone is using AI as a tool effectively rather than passively?

- A. They ask Claude for help on every task before attempting it themselves.
- B. They evaluate AI output critically before deciding whether to use it.**
- C. They use the most expensive AI model available for all tasks.
- D. They check AI outputs for formatting errors before submitting.

Correct Answer: B

Explanation

Effective tool use is characterised by critical evaluation — engaging judgment before accepting output. Relying on AI before attempting a task, using expensive models, or only checking formatting are patterns of passive use that produce lower quality outcomes.

Q24. You want to use Claude consistently across your team. The most important first step is:

- A. Ensure every team member has a Claude account and access.
- B. Identify which team member will manage Claude interactions.
- C. Select the Claude model tier that fits your budget.
- D. Define what tasks Claude will handle and establish quality standards.**

Correct Answer: D

Explanation

Consistent team AI use requires first defining scope (what tasks) and standards (what good looks like). Without this, different team members use AI differently, producing inconsistent quality and creating coordination failures.

Q25. What is the best mental model for AI in a professional context?

- A. An expert advisor whose recommendations can be trusted directly.
- B. A search tool that retrieves answers from a large knowledge base.
- C. A capable junior colleague who works fast but needs oversight.**
- D. A creative tool useful only for ideation and brainstorming.

Correct Answer: C

Explanation

"Capable junior colleague who needs oversight" is the most accurate working mental model — it acknowledges real capability (fast, broad, reasonable first drafts) while preserving the human responsibility to review, verify, and take ownership of final outputs.

Q26. You discover that a prompt that worked last week produces worse results today. The most likely cause is:

- A. Your internet connection affects Claude's response quality.
- B. The model may have been updated, changing its behaviour.**
- C. The prompt degraded because it was used too many times.
- D. Claude memorised your previous conversations and changed strategy.

Correct Answer: B

Explanation

AI model providers update and fine-tune models regularly. A prompt that relied on specific model behaviour may produce different results after an update. Prompts should be tested when model versions change.

Q27. What is the most appropriate response when Claude says "I'm not sure about this"?

- A. Treat the output as lower-confidence and verify the key claims.**
- B. Ask Claude the same question again for a more confident answer.
- C. Discard the output entirely — uncertainty means the answer is wrong.
- D. Accept the answer — Claude is unusually honest when uncertain.

Correct Answer: A

Explanation

Claude's expressed uncertainty is a signal to treat the output as lower-confidence — not to repeat the prompt or discard the output. Verify the specific uncertain claims independently. Note: Claude is not reliably calibrated on uncertainty and can be confident when wrong.

Q28. You need Claude to maintain a consistent tone across a long document. The best approach is:

- A. Generate the document in one prompt to maintain consistency.
- B. Ask Claude to "maintain consistent tone" throughout.
- C. Provide a style guide or tone example at the start of your prompt.**
- D. Generate each section separately and edit tone manually at the end.

Correct Answer: C

Explanation

Claude maintains tone most reliably when given explicit reference material — a style guide, sample paragraph, or defined tone descriptors — at the start of the prompt. A general instruction to "be consistent" is too vague to reliably produce uniformity across a long document.

Q29. Which approach produces the best AI-assisted research output?

- A. Use AI to find and summarise sources, then compile its summaries.
- B. Use AI exclusively since it has processed more research than any human.
- C. Use AI only to format and structure research you have already done.
- D. Use AI to generate hypotheses and structures, verify facts from primary sources.**

Correct Answer: D

Explanation

The most effective approach combines AI's speed in generating structure and hypotheses with human discipline in verifying facts from primary sources. AI summaries of sources are unreliable; AI-only research bypasses the human judgment that academic research requires.

Q30. A generalist should be able to explain why AI sometimes fails on a task. Which explanation is most accurate?

- A. The model is not large enough for that task type.
- B. The task requires precision or real-time data outside AI's design strengths.**
- C. The task was too complex for any current AI system to handle.
- D. Insufficient training data exists for that specific task.

Correct Answer: B

Explanation

Most professional AI failures occur because the task requires something outside AI's strengths — precise arithmetic, current information, legally binding judgment, or verified sourcing. Understanding this allows a generalist to route tasks appropriately from the start.

SET

3

Entry — Prompting for Professional Output

Set 3 of 10 · 15 Questions · AI Generalist

Q31. You want Claude to write a client email that sounds like you. The most effective prompt element is:

- A. A sample of your previous emails as a tone reference.
- B. A list of words you prefer and words to avoid.
- C. An instruction to write in first person and sound natural.
- D. Your name and job title for personalisation.

Correct Answer: A

Explanation

Tone is best conveyed through examples, not descriptions. A sample of your actual previous writing gives Claude far more signal about your voice — sentence length, hedging style, formality, directness — than any description of those qualities.

Q32. You ask Claude to write a 500-word summary. It produces 750 words. The most effective fix is:

- A. "Make this shorter."
- B. "Cut 250 words from this summary."
- C. "Rewrite this as exactly 500 words. Do not add new points — compress the existing ones."
- D. "Summarise this summary in 500 words."

Correct Answer: C

Explanation

The most effective instruction is specific and mechanically clear: target length (500 words), method (compress existing content), and constraint (no new points). Vague length instructions like "make it shorter" rarely hit the target.

Q33. What is the most reliable way to get Claude to produce a structured output format every time?

- A. Use the phrase "always follow this format" in the instructions.
- B. Provide the exact template or schema you want in the prompt.**
- C. Ask Claude to review its output for format compliance.
- D. Generate multiple outputs and select the one with the right format.

Correct Answer: B

Explanation

The most reliable structural control is providing the exact template — showing Claude what the output should look like rather than describing it. Templates are more precise than format descriptions and produce more consistent results.

Q34. You want Claude to generate 10 marketing taglines for your product. You keep getting similar options. The fix is:

- A. Ask for 10 more taglines and select from the combined 20.
- B. Switch to a more creative AI tool built for copywriting tasks.
- C. Run the same prompt 5 times and take the best from each.
- D. Add constraints that force variety — different emotions, audiences, or lengths.**

Correct Answer: D

Explanation

Sameness in creative outputs comes from insufficient constraints. Adding variety constraints — one tagline per target emotion, one per audience segment, one short and punchy — forces Claude to explore different angles rather than converging on the most probable options.

Q35. Which prompt technique most improves the accuracy of factual AI outputs?

- A. Ask Claude to cite evidence or reasoning for each claim it makes.**
- B. Ask Claude to rate its own confidence on a 1-10 scale.
- C. Ask Claude to search the internet for supporting evidence.
- D. Ask Claude to list what it does not know before answering.

Correct Answer: A

Explanation

Asking Claude to provide evidence or reasoning for each claim surfaces the basis of its statements — making it easier to identify which claims are well-grounded and which are assumptions or hallucinations. Self-confidence ratings are unreliable.

Q36. You want to use Claude to prepare an executive summary. The summary should contain no information that is not in the source document. The best instruction is:

- A. "Write an accurate executive summary of this document."
- B. "Summarise only information explicitly present in this document. Do not add external context."**
- C. "Summarise this document without adding opinions."
- D. "Do not hallucinate when summarising this document."

Correct Answer: B

Explanation

Explicit grounding instructions — "only information present in this document" and "do not add external context" — directly constrain Claude's output to the source material. Generic accuracy instructions and hallucination prohibitions are too vague to reliably constrain output.

Q37. What does "few-shot prompting" mean?

- A. Running the same prompt a few times and selecting the best output.
- B. Using a short, concise prompt rather than a long detailed one.
- C. Giving Claude a few tries to produce an acceptable output.
- D. Including examples of desired inputs and outputs in the prompt.**

Correct Answer: D

Explanation

Few-shot prompting means providing example input-output pairs in the prompt — showing Claude what a correct response looks like before giving it the actual task. Examples are generally more reliable than descriptions for defining desired output format and quality.

Q38. A client asks you to produce a document using AI. You should:

- A. Disclose that AI produced it and let the client decide whether to use it.
- B. Only use AI if the client has explicitly approved AI use.
- C. Use AI to draft it, then review, verify, and own the final output.**
- D. Not use AI for client deliverables — only for internal work.

Correct Answer: C

Explanation

Using AI to draft, then reviewing, verifying, and owning the final output is the professional standard. The AI is a production tool; you are the accountable professional. Blanket policies requiring prior approval or prohibiting client work are overly restrictive given how widespread AI use now is.

Q39. You use Claude to write a proposal and it sounds generically professional but lacks your company's specific value proposition. The cause is:

- A. Claude cannot write specific business proposals — only general ones.
- B. The prompt did not include your company's differentiated positioning.**
- C. The proposal genre is too structured for AI to handle effectively.
- D. Claude's training data does not include specialised business writing.

Correct Answer: B

Explanation

Generic output reflects generic input. Claude has no knowledge of your company's specific positioning, differentiators, or language unless you provide it. The fix is to include that positioning in the prompt — not to change tools.

Q40. Which of these represents the most advanced prompting behaviour?

- A. Iterating on output by providing specific, targeted improvement instructions.**
- B. Writing the longest, most detailed possible initial prompt.
- C. Using multiple AI tools on the same task and comparing outputs.
- D. Memorising the best prompts for each task type in your work.

Correct Answer: A

Explanation

The most advanced behaviour is iteration — the ability to evaluate output, identify exactly what needs improvement, and provide targeted instructions to fix it. Length, multi-tool comparison, and memorisation are all lower-leverage than skilled iteration.

Q41. You need Claude to maintain the same tone and approach across 20 different outputs. The most efficient approach is:

- A. Copy-paste the best previous output as context for each new prompt.
- B. Review and manually align tone across all 20 outputs after generation.
- C. Create a reusable system prompt that defines tone, role, and constraints.**
- D. Generate all 20 in one prompt with consistent instructions.

Correct Answer: C

Explanation

A reusable system prompt — defining role, tone, style, and constraints once — is the most efficient approach for consistent multi-output production. It applies the same context to each generation without requiring repetition or manual alignment.

Q42. What is the clearest signal that an AI output requires significant human intervention before use?

- A. It is longer than what you asked for.
- B. It does not match the exact format you requested.
- C. It sounds too polished and formal for the context.
- D. It contains specific figures or claims that are hard to verify quickly.**

Correct Answer: D

Explanation

Specific figures, dates, statistics, or citations that are hard to verify quickly are the clearest signal requiring human intervention — because these are exactly where hallucination risk is highest and where errors have the most consequential impact.

Q43. A new team member asks: "What is the most important thing to know about working with Claude?" The best answer is:

- A. "Use it for creative tasks but avoid it for factual research."
- B. "Claude generates plausible text — verify anything that matters."**
- C. "The more detailed your prompt, the better the output."
- D. "Claude is reliable for most professional tasks if you use it correctly."

Correct Answer: B

Explanation

The single most important principle for a new Claude user is: outputs are plausible-sounding text, not verified facts. This calibrates their trust appropriately and establishes the verification habit that makes all other AI use safe.

Q44. You want Claude to generate ideas that are genuinely different from its first response. The best approach is:

- A. Increase the temperature setting and regenerate the same prompt.
- B. Use a different AI tool which may have different training data.
- C. Ask for ideas that explicitly avoid the approaches in the first response.**
- D. Ask Claude to evaluate and critique its first response before generating more.

Correct Answer: C

Explanation

Explicitly instructing Claude to avoid the approaches in its first response forces it to explore a different part of the idea space. Increasing temperature increases randomness but may produce more of the same type. Critique-then-generate adds a step but does not guarantee different directions.

Q45. Which professional task is least suited to AI assistance right now?

- A. Making a final legal judgment in a contested regulatory dispute.**
- B. Drafting a first version of a legal argument for attorney review.
- C. Summarising case precedents from a set of provided documents.
- D. Generating a list of research questions for a legal brief.

Correct Answer: A

Explanation

Final legal judgment requires licensed professional accountability, deep contextual judgment, and real consequences. All the other tasks produce draft output for human review — an appropriate use case for AI that keeps the professional accountable.

SET

4

Entry — AI Tools and Workflows

Set 4 of 10 · 15 Questions · AI Generalist

Q46. What is the correct order of operations when using AI to analyse a dataset?

- A. Upload the data, run AI analysis, and present the outputs directly.
- B. Clean the data with AI first, then define the question, then analyse.
- C. Run exploratory AI analysis first to discover what questions to ask.
- D. Define the question first, then use AI to analyse, then verify key outputs.**

Correct Answer: D**Explanation**

The correct order starts with a defined question — AI analysis without a clear question produces unfocused outputs. After analysis, key outputs should be verified before use. Presenting raw AI analysis directly skips critical verification.

Q47. Which combination of tools is most powerful for a knowledge worker using AI?

- A. Claude plus a second AI tool to cross-check Claude's outputs.
- B. Claude for generation and synthesis, plus a primary source for verification.**
- C. Claude plus a premium plan for guaranteed higher accuracy.
- D. Claude plus an AI writing checker to validate quality before use.

Correct Answer: B**Explanation**

The most powerful combination is generative AI (for speed and synthesis) paired with authoritative primary sources (for verification). Cross-checking with another AI uses the same structural weakness; premium plans do not guarantee accuracy; AI quality checkers are themselves AI.

Q48. You are evaluating whether to use AI for a specific task at work. The most important question is:

- A. "Which AI tool has the highest accuracy rating for this task type?"
- B. "Is this task in the list of approved AI use cases at my company?"
- C. "What happens when AI gets this wrong and how will I catch it?"**
- D. "How much time will AI save me compared to doing it manually?"

Correct Answer: C

Explanation

The most important evaluative question is about failure mode — what the error looks like and how it will be caught. Time savings and approved use cases matter, but they are secondary to understanding what goes wrong and whether you have a way to detect it.

Q49. What does a well-designed AI workflow always include?

- A. A human checkpoint to review and validate AI outputs before they matter.**
- B. A fallback to a different AI tool if the primary one fails.
- C. A manual alternative process in case the AI is unavailable.
- D. A logging system that records all AI inputs and outputs.

Correct Answer: A

Explanation

Every effective AI workflow includes human checkpoints before outputs are used for consequential decisions. Technical redundancy (fallback tools, manual alternatives, logging) is useful but secondary to the fundamental human-in-the-loop design principle.

Q50. You are building an AI-assisted customer service workflow. The most important metric to track is:

- A. The average response time of the AI compared to human agents.
- B. The percentage of queries handled without human intervention.
- C. The cost per query with AI versus without AI.
- D. The rate of customer complaints or escalations from AI responses.**

Correct Answer: D

Explanation

Complaint and escalation rate measures actual customer harm from AI errors — the most important outcome. Response time, automation rate, and cost are efficiency metrics that matter only if quality is maintained.

Q51. Which workflow task is most appropriate to delegate entirely to Claude without human review?

- A. Drafting a response to a customer complaint about a billing error.
- B. Summarising a competitor's product features from a specification.
- C. Formatting a meeting transcript into a consistent structure.**
- D. Generating a project risk assessment for a client proposal.

Correct Answer: C

Explanation

Formatting a transcript into a consistent structure is a low-stakes structural task where errors are immediately visible and easily corrected. The other tasks involve AI making interpretive or evaluative judgments in higher-stakes contexts requiring human review.

Q52. What is the most important reason to document the prompts you use for professional work?

- A. So you can reproduce reliable outputs and improve them over time.**
- B. So you can prove to clients that AI was used professionally.
- C. So you can train your own AI model on effective prompt patterns.
- D. So you comply with your company's AI usage logging policy.

Correct Answer: A

Explanation

Documenting prompts enables reproducibility — you can reliably produce the same quality output again — and improvement over time through iteration. This is the core operational reason; compliance and proof of use are secondary considerations.

Q53. You use Claude to generate 5 customer persona profiles. Before using them, you should:

- A. Ask Claude to check whether the personas are realistic.
- B. Validate them against real customer data or research.**
- C. Add your team's subjective preferences to each persona.
- D. Use the personas directly — AI personas are built from broad patterns.

Correct Answer: B

Explanation

AI-generated personas are built from statistical patterns in training data, not your actual customer data. Before using them professionally, validate them against real research, interviews, or data about your specific customer base.

Q54. What is the key difference between using AI as a "co-pilot" versus using AI as an "autopilot"?

- A. Co-pilot is for creative tasks; autopilot is for repetitive tasks.
- B. Autopilot is more reliable than co-pilot for professional output.
- C. Co-pilot means you remain in control; autopilot means AI makes decisions.**
- D. Co-pilot requires technical skills; autopilot does not.

Correct Answer: C

Explanation

Co-pilot means human judgment remains in the loop — AI assists while the human steers, reviews, and decides. Autopilot means AI operates without meaningful human oversight. In professional contexts, most tasks require co-pilot, not autopilot, given current AI limitations.

Q55. You use AI to generate a client-facing presentation. The most important pre-send check is:

- A. Ensure the formatting and visual design meet brand standards.
- B. Confirm the length and structure match the client's preferences.
- C. Check that the language is appropriately formal for the client.
- D. Verify every factual claim, figure, and recommendation for accuracy.**

Correct Answer: D

Explanation

Factual accuracy is the primary check before any AI output reaches a client. Brand standards, length, structure, and language are important but secondary — inaccurate information reaching a client creates professional and reputational risk.

Q56. What is the most effective way to use AI for decision-making support?

- A. Use AI to map options and tradeoffs, then apply your own judgment.**
- B. Use AI to recommend the best decision based on your inputs.
- C. Use AI to validate your instinct before taking action.
- D. Use AI to identify the option with the highest probability of success.

Correct Answer: A

Explanation

AI is most effective as a structured thinking tool — mapping out options, articulating tradeoffs, and surfacing considerations. The final decision and its accountability must remain with the human, who has context, values, and responsibility AI cannot replicate.

Q57. You are using Claude in a meeting to take live notes and action items. The most important post-meeting step is:

- A. Share the notes immediately — AI is faster than human review.
- B. Review and correct the notes before sharing with attendees.**
- C. Ask Claude to clean up and finalise the notes before sending.
- D. Send the raw notes and ask attendees to correct their own action items.

Correct Answer: B

Explanation

Live AI note-taking is valuable for speed but requires human review before distribution — names, action items, decisions, and context can all be captured incorrectly. Sharing raw AI notes without review risks distributing errors to the whole team.

Q58. Which of these represents the highest-maturity level of AI use in a professional context?

- A. Being faster than colleagues at completing tasks using AI tools.
- B. Using AI for every task in your daily workflow without exception.
- C. Knowing which AI tool to use for each category of work task.
- D. Designing AI-enabled systems that others in the team can use reliably.**

Correct Answer: D

Explanation

The highest maturity level is systemic — not just using AI yourself but designing workflows, prompts, and guardrails that allow others to use AI reliably and safely. Speed, breadth of use, and tool selection are earlier stages of individual AI capability.

Q59. What should a professional do when AI generates a response that confidently contradicts their domain expertise?

- A. Investigate carefully — the AI may be wrong, or they may have a gap.**
- B. Trust their expertise over AI — human domain knowledge is more reliable.
- C. Trust the AI — it has processed more information on the topic.
- D. Run the same query through a different AI to resolve the conflict.

Correct Answer: A

Explanation

When AI confidently contradicts domain expertise, the correct response is careful investigation — not automatic deference to expertise or to AI. The AI may have identified a genuine gap, or it may be hallucinating. Only investigation determines which.

Q60. What is "prompt drift" and why does it matter?

- A. AI models produce lower quality outputs after extended use.
- B. Prompts become less effective when used by many different people.
- C. Output quality degrades over a long conversation as context accumulates.**
- D. AI responses drift toward generic outputs when context is too brief.

Correct Answer: C

Explanation

Prompt drift occurs in long conversations as the context window fills — earlier instructions get pushed further back and receive less attention, causing the model to drift from its original briefing. Refreshing instructions periodically prevents this in extended workflows.

SET

5

Entry — AI Judgment: Applied Scenarios

Set 5 of 10 · 15 Questions · AI Generalist

Q61. Your manager forwards an impressive AI-generated report. They ask you to build on it. Your first step is:

- A. Verify the key claims and figures before treating the report as accurate.**
- B. Use Claude to expand and build on the existing AI-generated content.
- C. Accept the report as accurate — your manager reviewed it.
- D. Identify which AI tool generated the report to assess its reliability.

Correct Answer: A

Explanation

AI-generated reports may contain hallucinated facts, distorted analysis, or unverified claims. Before building on them, verify the key inputs. Your manager reviewing it does not guarantee accuracy — they may not have verified either.

Q62. A colleague builds an AI chatbot for your team's internal use without telling anyone. The key concern is:

- A. The chatbot may be slower than commercial AI tools the team uses.
- B. The chatbot will not be as capable as enterprise AI solutions.
- C. Using an internal chatbot may violate AI tool licensing terms.
- D. Sensitive internal data shared with it may not be appropriately secured.**

Correct Answer: D

Explanation

Unvetted internal AI deployments create data security and governance risks — sensitive information shared with the chatbot may be stored, logged, or exposed in ways that violate data policies. Capability and licensing concerns are secondary.

Q63. You receive a task from a client: "Use AI to do X." The client does not know AI limitations. Your responsibility is to:

- A. Tell the client that AI cannot reliably do this task as described.
- B. Complete the task using AI appropriately and deliver verified output.**
- C. Do the task manually since the client doesn't understand AI well.
- D. Ask the client to specify exactly which AI tool they want you to use.

Correct Answer: B

Explanation

Your professional responsibility is to use AI to the extent it genuinely helps, verify the output, and deliver quality work. Educating the client about limitations is appropriate, but not as a reason to avoid the task — bring your professional judgment to bear on how AI is used.

Q64. You use AI to generate 10 ideas for a strategy session. You share them without disclosure. What is the risk?

- A. Someone on the team will identify the ideas as AI-generated.
- B. The ideas will be lower quality than human-generated ones.
- C. The team treats the ideas as human-generated strategic thinking.**
- D. The client will ask which team member generated each idea.

Correct Answer: C

Explanation

Undisclosed AI-generated ideas in strategic discussion are treated as human strategic judgment — which they are not. This distorts the team's confidence in the ideas, bypasses human creative engagement, and creates accountability ambiguity.

Q65. You are working on a time-sensitive task and Claude produces an output with obvious errors. You should:

- A. Fix the errors yourself and complete the task — do not resubmit to Claude.**
- B. Ask Claude to fix its own errors before you review the output.
- C. Start over with a better prompt to avoid getting errors again.
- D. Submit the task with the errors flagged as "pending verification."

Correct Answer: A

Explanation

When you can see the errors, fix them yourself — you know what is wrong and can correct it faster than prompting Claude again and reviewing another output. Asking Claude to fix its own errors is unreliable; starting over wastes time when you already have 80% of what you need.

Q66. A startup asks you to evaluate their AI product. Which question reveals the most about real-world reliability?

- A. "What accuracy rate does it achieve on benchmark tests?"
- B. "Which underlying AI model does it use?"
- C. "What does it do when it gets the answer wrong?"**
- D. "How many customers are currently using it?"

Correct Answer: C

Explanation

How a product handles failure is the most revealing reliability question. Benchmark accuracy measures best-case performance in controlled conditions; the real test is whether the system degrades gracefully, catches errors, and allows recovery when something goes wrong.

Q67. You use Claude to analyse a competitor's strategy from public information. The output is insightful. Before presenting it:

- A. Present it directly — public information analysis is low-risk.
- B. Have Claude revise the analysis for clarity before presenting.
- C. Identify which data points came from Claude's training versus the document.
- D. Verify the key facts about the competitor from primary sources.**

Correct Answer: D

Explanation

Even when AI analysis feels insightful, the specific facts about the competitor — market share figures, product details, strategic moves — must be verified from primary sources. AI can misattribute, confuse, or hallucinate competitor-specific details.

Q68. A junior team member says "Claude told me we should use approach X for this project." The correct response is:

- A. "If Claude recommended it, we should take it seriously."
- B. "Walk me through the reasoning — Claude's suggestions need human judgment."**
- C. "We need to ask another AI tool to validate Claude's recommendation."
- D. "Claude cannot make recommendations on technical project decisions."

Correct Answer: B

Explanation

Claude's suggestions should be inputs to human judgment, not directives. Walking through the reasoning allows the team to evaluate whether the suggestion makes sense in context — which is exactly the human judgment that makes AI useful and safe.

Q69. You use Claude to draft communications for a sensitive HR situation. The most important consideration is:

- A. The communication must be reviewed by the CEO before sending.
- B. You should disclose to the employee that AI drafted the communication.
- C. Legal and procedural accuracy must be verified with HR and legal counsel.**
- D. Claude should not be used for any HR-related communications.

Correct Answer: C

Explanation

HR communications in sensitive situations have legal, procedural, and human implications that AI cannot reliably get right. Verification with HR and legal counsel is essential regardless of how well the draft reads.

Q70. Which of the following best describes an "AI-first" professional approach?

- A. Considering how AI can enhance every task before defaulting to manual work.**
- B. Using AI for every task without exception to maximise efficiency.
- C. Building AI tools before deciding what problem to solve.
- D. Replacing all manual processes with AI-powered alternatives.

Correct Answer: A

Explanation

AI-first means considering AI as an option from the start of every task — asking "how can AI enhance this?" before defaulting to a fully manual approach. It is not using AI for everything regardless of fit, nor is it building tools before problems.

Q71. You delegate a report to Claude and the deadline arrives. The report has errors. Who is accountable?

- A. Claude is — it produced the incorrect content.
- B. You are — you chose to delegate to AI and are responsible for the output.**
- C. The manager is — they set the deadline without allowing review time.
- D. No one — AI errors are unforeseeable and cannot be assigned.

Correct Answer: B

Explanation

Accountability for AI output always rests with the professional who used it. Delegating to Claude does not transfer accountability. The lesson is that AI delegation requires building in review time before the deadline.

Q72. A client insists on a deliverable that AI cannot reliably produce accurately. You should:

- A. Use AI anyway and review carefully to catch all errors.
- B. Decline the deliverable — professional standards prohibit unreliable AI use.
- C. Outsource the deliverable to a specialist who can do it without AI.
- D. Explain the limitation and propose a hybrid approach with AI drafting and human verification.**

Correct Answer: D

Explanation

The professional response is to be transparent about AI limitations and propose a workable hybrid: AI for speed and structure, human expertise for verification and quality assurance. This serves the client better than blind AI use or unnecessary refusal.

Q73. You are asked to train colleagues on AI best practices. The most important lesson to start with is:

- A. "Always use the most advanced AI model available for any task."
- B. "AI tools are most useful for creative and ideation tasks."
- C. "AI outputs are starting points — own the final output yourself."**
- D. "AI can be trusted for most tasks once you have the right prompt."

Correct Answer: C

Explanation

The most foundational lesson — and the one with the most immediate impact on how people use AI safely — is ownership of final output. Everything else (tool selection, prompting technique, task suitability) builds on this foundation.

Q74. What is the best indicator that a professional is ready to lead AI adoption in their organisation?

- A. They have completed an AI certification from a recognised provider.
- B. They can explain both AI's value and its limitations with equal accuracy.**
- C. They have built at least one working AI tool using an API.
- D. They use AI every day across all their professional tasks.

Correct Answer: B

Explanation

Leadership in AI adoption requires nuanced understanding — knowing what AI can genuinely do and where it genuinely fails, and communicating both accurately. Certification, tool-building, and daily use are valuable signals but not sufficient without this balanced judgment.

Q75. What does Menler's AI Generalist Entry Diagnostic assess above all else?

- A. Whether a learner has used AI tools for more than six months.
- B. Whether a learner can build AI tools using Claude's API.
- C. Whether a learner has prior work experience using AI professionally.
- D. Whether a learner's AI judgment and instincts are ready for fellowship training.**

Correct Answer: D

Explanation

The Entry Diagnostic assesses readiness for fellowship training — specifically whether the learner has the judgment, critical instincts, and practical AI understanding needed to benefit from and contribute to the program. Tool experience, tenure, and technical skills are secondary.

SET

6

Level Up — Advanced Prompting and Output Quality

Set 6 of 10 · 15 Questions · AI Generalist

Q1. You want Claude to evaluate a business argument and push back on weak points. The most effective prompt element is:

- A. A request to list the argument's strengths and weaknesses.
- B. An explicit instruction to steelman then critically challenge the argument.**
- C. An instruction to be "honest and critical" in its response.
- D. A request for Claude to play a devil's advocate role.

Correct Answer: B

Explanation

Steelmanning first (presenting the argument as strongly as possible) then critically challenging forces a genuine engagement with both sides. "Honest and critical" is too vague; "devil's advocate" alone may produce superficial pushback rather than rigorous analysis.

Q2. You receive a Claude output that is technically accurate but reads like generic AI text. The most targeted fix is:

- A. Ask Claude to "make it sound more human and natural."
- B. Regenerate with higher temperature to reduce generic patterns.
- C. Use a different AI writing tool better suited to your style.
- D. Provide a voice sample and specify tone dimensions to match.**

Correct Answer: D

Explanation

"Make it more human" is too vague. Providing a voice sample and defining specific tone dimensions (directness level, sentence length, hedging style) gives Claude concrete signal to match against. Temperature affects randomness, not voice — it can increase variation without improving fit.

Q3. What is the most reliable indicator that a multi-step prompt chain is designed correctly?

- A. Each step's output is fully sufficient as the next step's input.**
- B. The chain produces consistent outputs across 10 test runs.
- C. Each step uses a different AI tool for its specific capability.
- D. The total prompt length across all steps is under the context limit.

Correct Answer: A

Explanation

A well-designed prompt chain has clean handoffs — the output of each step is complete, correctly formatted, and sufficient as input to the next step without requiring manual intervention. Consistency and length are secondary indicators.

Q4. You need Claude to produce a high-stakes client analysis. What does "grounding" the prompt mean?

- A. Testing the prompt on a similar lower-stakes task first.
- B. Adding explicit accuracy instructions to the prompt.
- C. Providing primary source documents for Claude to work from.**
- D. Restricting the output to information Claude is confident about.

Correct Answer: C

Explanation

Grounding means providing the source material — documents, data, specific context — that Claude should reason from, rather than allowing it to draw from general training knowledge. This reduces hallucination risk and ties output to verifiable inputs.

Q5. Which prompting behaviour produces the most improvement in Claude output quality over time?

- A. Using longer, more detailed prompts for every task type.
- B. Systematic iteration with specific, targeted feedback on each output.**
- C. Maintaining a library of prompts that worked previously.
- D. Testing each task on three different AI tools before committing.

Correct Answer: B

Explanation

Systematic iteration — evaluating what specifically failed or fell short in each output and providing targeted correction — is the learning loop that compounds. Length, libraries, and multi-tool testing are useful but secondary to this iterative refinement discipline.

Q6. What is "output specification" and why does it matter for professional AI use?

- A. Defining exactly what the desired output looks like before prompting.**
- B. Specifying which AI model should produce the output.
- C. Describing the output quality standard in a system prompt.
- D. Outlining the steps the AI should follow to produce the output.

Correct Answer: A

Explanation

Output specification means defining the target artefact — its format, length, structure, tone, and content boundaries — before writing the prompt. Starting with a clear picture of what you want makes it far easier to write a prompt that gets you there.

Q7. You need Claude to extract specific structured data from unstructured text reliably. The best approach is:

- A. Ask Claude to extract all relevant information it can find.
- B. Ask Claude to summarise the text, then extract from the summary.
- C. Run the extraction 3 times and reconcile differences manually.
- D. Provide the exact extraction schema with labeled field examples.**

Correct Answer: D

Explanation

Reliable extraction requires explicit schema definition with examples showing the exact field names, data types, and format expected. Open-ended extraction produces inconsistent field names and coverage; three-run reconciliation is inefficient when proper specification solves the problem.

Q8. You ask Claude to write a persuasive essay. It produces a balanced, both-sides analysis instead. The cause is:

- A. Persuasive essays are outside Claude's trained capability.
- B. The prompt did not include enough facts to support a one-sided argument.
- C. Claude's training leads it to hedge toward balance unless explicitly told otherwise.**
- D. Claude's safety guidelines prevent one-sided persuasive content.

Correct Answer: C

Explanation

Claude's training reinforces hedged, balanced responses by default. Producing genuinely one-sided persuasive content requires an explicit instruction to commit to one position and not include counterarguments — overriding its default balance tendency.

Q9. A prompt that works perfectly in Claude 3 produces mediocre output in Claude 4. The best response is:

- A. Adapt the prompt — different model versions respond differently.**
- B. Revert to using Claude 3 for this specific task type.
- C. Conclude that Claude 4 is inferior to Claude 3 for this task.
- D. Submit feedback to Anthropic about the regression.

Correct Answer: A

Explanation

Prompts are model-specific — what works optimally for one version may not transfer perfectly to another due to changes in training, instruction following, or default behaviour. Adapting the prompt to the new model's behaviour is the productive response.

Q10. What is the most important output quality dimension to check when Claude produces a complex analysis?

- A. Completeness — whether every possible point has been covered.
- B. Word count — whether the output meets the specified length.
- C. Formatting — whether headers and structure match the request.
- D. Logical consistency — whether conclusions follow from stated premises.**

Correct Answer: D

Explanation

Logical consistency is the most important dimension — whether the argument holds together and conclusions genuinely follow from the evidence and premises. Completeness, length, and formatting are secondary to the integrity of the reasoning itself.

Q11. What does it mean to use "chain of thought" prompting?

- A. Chaining multiple prompts together in a sequential workflow.
- B. Asking the model to show its reasoning step by step before answering.**
- C. Including previous conversation turns as context for the current prompt.
- D. Asking the model to consider multiple perspectives before concluding.

Correct Answer: B

Explanation

Chain of thought (CoT) prompting asks the model to articulate its reasoning process before arriving at its answer. This improves accuracy on complex tasks by encouraging the model to work through steps rather than jumping to a conclusion.

Q12. You notice Claude consistently makes a specific type of error across several outputs. The most effective response is:

- A. Switch to a different model since Claude cannot handle this task.
- B. Add a general instruction to "be careful" about accuracy.
- C. Add a targeted negative example of that error to the prompt.**
- D. Run each output through a separate AI tool to catch the error.

Correct Answer: C

Explanation

Negative examples — showing exactly what the error looks like and labelling it as wrong — are more effective than general accuracy instructions at preventing specific recurring errors. They give Claude a concrete pattern to avoid.

Q13. What is the most effective use of the "system prompt" in a Claude workflow?

- A. Providing the primary user query that Claude should respond to.
- B. Listing all the topics the AI is allowed to discuss in the session.
- C. Summarising previous conversation for continuity across sessions.
- D. Setting persistent role, tone, constraints, and task context for the session.**

Correct Answer: D

Explanation

System prompts are designed to set the persistent operational context — role, persona, constraints, output format, tone — that applies to all subsequent interactions in the session. They are not for user queries, topic whitelisting, or conversation history.

Q14. You want Claude to maintain a consistent analytical framework across a 10-part analysis series. The most efficient approach is:

- A. Copy the framework into each individual prompt as a reminder.
- B. Define the framework in the system prompt and reuse it across all parts.**
- C. Produce all 10 parts in a single prompt to maintain consistency.
- D. Generate part 1, then paste it as context for each subsequent part.

Correct Answer: B

Explanation

A system prompt containing the analytical framework applies it consistently and efficiently across all parts without copy-pasting. Producing all 10 in one prompt risks context window limits and reduced quality on later sections.

Q15. An output fails a critical quality check. Which response demonstrates the highest professional judgment?

- A. Identify the specific failure, fix it manually, and update the prompt to prevent recurrence.**
- B. Regenerate the output with the same prompt and hope for better results.
- C. Lower the quality standard to match what the AI consistently produces.
- D. Document the failure and avoid using AI for similar tasks in the future.

Correct Answer: A

Explanation

Professional response to failure is diagnostic and constructive: identify the specific cause, fix the immediate output, and improve the prompt to prevent the same error recurring. The other options accept or avoid the failure rather than learning from it.

SET

7

Level Up — AI Strategy: Where AI Adds Value

Set 7 of 10 · 15 Questions · AI Generalist

Q16. A company is deciding where to apply AI first. The most strategically sound criterion is:

- A. The most complex tasks to demonstrate AI's maximum capability.
- B. The tasks that executives are most excited about automating.
- C. High-volume, repetitive tasks where errors are catchable and correctable.**
- D. The tasks where AI tools already exist off-the-shelf.

Correct Answer: C

Explanation

The best first AI applications combine volume (where time savings compound), repetition (where reliable prompt patterns emerge), and catchable errors (where quality control is built in). Starting with high-complexity or executive-favoured tasks without these properties is high-risk.

Q17. A startup says "we use AI for everything." What is the most important evaluative follow-up?

- A. "What does your human review process look like for AI outputs?"**
- B. "Which AI tools are you using and how much do they cost?"
- C. "What is your AI accuracy rate across your use cases?"
- D. "How many staff did you replace with AI automation?"

Correct Answer: A

Explanation

The critical question is human review — AI use without meaningful oversight is the failure mode, not the success story. The question probes whether "we use AI for everything" means thoughtful deployment or blind automation.

Q18. What distinguishes a genuine AI strategy from a list of AI tools to use?

- A. A strategy covers more tools and use cases than a simple list.
- B. A strategy is approved by senior leadership; a list is not.
- C. Strategy specifies the technical infrastructure AI will run on.
- D. Strategy defines why AI is used, for what, and how success is measured.**

Correct Answer: D

Explanation

A genuine AI strategy defines purpose (why), scope (which tasks), governance (how), and measurement (success criteria). A list of tools is operational — it tells you what, not why or how to evaluate whether it is working.

Q19. An organisation wants to build internal AI capabilities rather than buy off-the-shelf tools. The most important first question is:

- A. "Which open-source AI models can we fine-tune for our use case?"
- B. "Do we have the data, talent, and processes to do this sustainably?"**
- C. "What is the cost difference between building and buying?"
- D. "What are our competitors building internally?"

Correct Answer: B

Explanation

Build vs buy in AI is ultimately a capability question — do you have the data to train on, the talent to maintain, and the processes to govern? Cost and competitive intelligence are secondary inputs; without this foundational question answered, build decisions often fail.

Q20. What does "AI augmentation" mean in a strategic context?

- A. Adding AI features to an existing product or software platform.
- B. Increasing the size or capability of an existing AI model.
- C. Using AI to enhance human capability rather than replace human roles.**
- D. Using multiple AI models in combination for better outputs.

Correct Answer: C

Explanation

AI augmentation is the strategic choice to deploy AI as a capability multiplier for human workers — making people more effective rather than replacing them. This is distinct from automation (replacing tasks) and from product AI features.

Q21. Which metric most accurately measures the value AI is adding to a knowledge work team?

- A. Total number of AI tool interactions per week.
- B. Reduction in task completion time versus pre-AI baseline.
- C. Employee satisfaction with AI tools after 90 days of use.
- D. Quality-adjusted output per person across AI-assisted tasks.**

Correct Answer: D

Explanation

Quality-adjusted output — accounting for both quantity and quality of work — is the most meaningful measure of AI's value. Speed improvements that produce lower-quality work are not genuine productivity gains. Interaction count and satisfaction are proxies, not outcomes.

Q22. A team has been using AI for three months and productivity has not improved. The most likely cause is:

- A. AI is being used for the wrong tasks or without proper review workflows.**
- B. Three months is too short to see AI productivity improvements.
- C. The team is using AI tools that are not suited to their industry.
- D. AI productivity gains require a minimum team size to be visible.

Correct Answer: A

Explanation

The most common cause of AI productivity failures at the team level is misapplication — either using AI for tasks where it adds little value, or adding AI without the review workflows that ensure quality output. Time, tool choice, and team size are secondary factors.

Q23. What is the correct mental model for an "AI-native" business process?

- A. The process runs entirely on AI with no human involvement.
- B. AI is embedded in how the process works from the start, not added later.**
- C. The process was originally designed by an AI tool.
- D. The process uses only AI tools built specifically for that industry.

Correct Answer: B

Explanation

AI-native means AI is incorporated into the design of the process from the beginning — not retrofitted. This produces better outcomes than bolting AI onto existing manual workflows because the workflow is designed around AI's capabilities and limitations from the outset.

Q24. The most common reason enterprise AI projects fail is:

- A. AI models not being advanced enough for enterprise requirements.
- B. Lack of budget for enterprise AI tool licences.
- C. Resistance from employees who fear job replacement.
- D. Poor data quality, unclear objectives, or insufficient human oversight.**

Correct Answer: D

Explanation

Enterprise AI project failures most commonly trace to foundational issues: poor data quality, undefined success metrics, or AI deployed without the governance and oversight that enterprise use requires. Model capability is rarely the limiting factor; organisational readiness is.

Q25. A client asks: "Should we replace our customer service team with AI?" The most responsible first response is:

- A. "Yes, AI can handle most customer service queries effectively."
- B. "No, customer service requires human empathy AI cannot provide."
- C. "What specific problems in your current customer service are you trying to solve?"**
- D. "It depends on your budget for AI tools versus headcount."

Correct Answer: C

Explanation

The most responsible response starts with the problem, not the technology. AI replacement decisions should be driven by what specific problems need solving — not by cost reduction or technology enthusiasm. The right solution depends entirely on the diagnosis.

Q26. What is the highest-value AI application for a solo entrepreneur?

- A. Using AI to replace the need for professional advisors.
- B. Automating high-volume, time-consuming tasks to free strategic time.**
- C. Building a personal AI model trained on their specific expertise.
- D. Using AI to generate social media content at scale.

Correct Answer: B

Explanation

For a solo operator, the highest-value AI use is freeing strategic time — identifying the repetitive, high-volume tasks that consume disproportionate time and using AI to handle them. AI advisors and personal models are high-effort; social content is a specific use, not the highest-value one universally.

Q27. What is the best indicator that an organisation's AI adoption is maturing?

- A. AI use is governed by defined policies and measured against outcomes.**
- B. The organisation has deployed AI tools across every department.
- C. AI was mentioned in the annual report and investor presentations.
- D. The organisation employs dedicated AI specialists in every team.

Correct Answer: A

Explanation

AI maturity is best indicated by governance and measurement — defined policies for appropriate use, quality standards, and outcome tracking. Broad deployment, communications, and headcount are all compatible with immature, uncontrolled AI adoption.

Q28. A company trains all employees on AI tools in a one-day workshop. What is most likely missing?

- A. A certificate that validates employee AI proficiency.
- B. An executive sponsor who mandates AI use after training.
- C. Ongoing practice, feedback, and use cases specific to each role.**
- D. A longer training duration — one day is insufficient for any learning.

Correct Answer: C

Explanation

A one-day AI workshop provides awareness but not capability. Genuine AI proficiency develops through role-specific practice, feedback on real work output, and iteration over time. The absence of ongoing, contextual learning is what makes one-time workshops fail.

Q29. Which capability is most scarce and valuable in an AI-enabled workforce?

- A. Technical ability to build and fine-tune AI models.
- B. Speed of AI tool adoption across different platforms.
- C. Ability to write long, detailed prompts for complex tasks.
- D. The judgment to direct AI toward the right problems with the right constraints.**

Correct Answer: D

Explanation

Direction and judgment — knowing what problem to solve, how to frame it for AI, how to evaluate the output, and when to escalate — is the scarcest and most compounding capability. Technical building, adoption speed, and prompt length are all secondary to this meta-skill.

Q30. An AI Generalist's ultimate responsibility in an organisation is:

- A. Training all employees to use AI tools effectively.
- B. Ensuring AI is directed toward genuine value and governed responsibly.**
- C. Building the technical AI infrastructure for the organisation.
- D. Staying updated on the latest AI tools and recommending new ones.

Correct Answer: B

Explanation

An AI Generalist's core responsibility is judgment — directing AI toward tasks where it genuinely adds value, and ensuring its use is governed, reviewed, and aligned with organisational standards. Training, infrastructure, and tool research are supporting activities.

SET

8

Level Up — Agentic AI and Workflow Automation

Set 8 of 10 · 15 Questions · AI Generalist

Q31. What makes an "agentic" AI different from a standard AI chatbot?

- A. It can take multi-step actions autonomously to complete a goal.**
- B. It has higher intelligence and can handle more complex questions.
- C. It learns from interactions and improves over time.
- D. It connects to the internet and retrieves live information.

Correct Answer: A**Explanation**

Agentic AI uses tools, executes multi-step plans, and takes actions in the world (browsing, coding, filing, API calls) to achieve a goal autonomously. A chatbot generates a response; an agent executes a process.

Q32. What is the most important design principle for an agentic AI workflow?

- A. Maximising the number of steps the agent can complete without human input.
- B. Using the most powerful available model as the agent's core.
- C. Defining clear boundaries for what actions the agent can take autonomously.**
- D. Ensuring the agent can retry failed steps automatically.

Correct Answer: C**Explanation**

Agentic systems must have clearly defined action boundaries — what they can and cannot do without human approval. Without boundaries, agents can take consequential irreversible actions the user did not intend. Autonomy scope is the primary design constraint.

Q33. You build an AI agent that books meetings on your behalf. What safety mechanism is most critical?

- A. Limiting the agent to reading calendar data, not writing to it.
- B. Requiring human confirmation before any booking is finalised.**
- C. Using only open calendar slots to avoid conflicts.
- D. Setting a maximum of 3 bookings per day to limit errors.

Correct Answer: B

Explanation

Human confirmation before any irreversible action — such as booking a meeting that appears in others' calendars — is the most critical safety mechanism. Limiting reads, using open slots, and capping volume are secondary constraints that do not prevent wrong bookings.

Q34. What is "tool use" in the context of Claude agents?

- A. The ability to use a wider range of vocabulary and expression styles.
- B. Access to multiple AI models depending on the task type.
- C. The ability to display formatted output including tables and code blocks.
- D. The ability to call external functions or APIs as part of generating a response.**

Correct Answer: D

Explanation

Tool use (function calling) allows Claude to call defined external functions — web search, code execution, file reading, API calls — as part of producing a response. This is what enables agentic behaviour beyond pure text generation.

Q35. You are designing an AI agent to process and categorise incoming documents. The most important reliability feature is:

- A. A logging system and human review queue for low-confidence categorisations.**
- B. Using the highest-accuracy AI model for all categorisations.
- C. A retry mechanism that attempts each categorisation 3 times.
- D. A testing protocol that evaluates the agent on 100 sample documents.

Correct Answer: A

Explanation

Any autonomous classification system needs a mechanism to flag and escalate uncertain cases for human review. Model accuracy, retries, and initial testing all matter, but a live system without a human review queue for edge cases will eventually make confident wrong decisions.

Q36. What is the key risk of allowing an AI agent to send emails autonomously?

- A. Email sending API costs may accumulate without visibility.
- B. It may send emails the user did not intend or with incorrect content.**
- C. Recipients may respond directly to the AI without knowing it.
- D. Email deliverability may be affected by AI-generated content.

Correct Answer: B

Explanation

Autonomous email sending means the agent can commit to external communications on behalf of the user — including wrong information, wrong recipients, or tone failures — without any human review. This is an irreversible, externally visible action and the primary risk.

Q37. What does "human in the loop" mean in an agentic AI system?

- A. A human writes the initial prompt that the agent acts on.
- B. A human monitors the agent's performance via a dashboard.
- C. A human is available to answer the agent's questions.
- D. A human reviews and approves agent decisions at defined checkpoints.**

Correct Answer: D

Explanation

"Human in the loop" means designed checkpoints where a human reviews, approves, or redirects agent decisions before they are executed — not just initial setup or passive monitoring. For high-stakes actions, these checkpoints are non-negotiable.

Q38. You want to use Claude to automate a multi-step research workflow. Which task should remain human-controlled?

- A. Compiling search results from multiple sources into a summary.
- B. Reformatting research notes into a consistent template.
- C. Deciding which sources are authoritative and whether to trust the findings.**
- D. Identifying which sources were published in the past 12 months.

Correct Answer: C

Explanation

Source trust evaluation and final judgment on research findings require domain expertise, critical thinking, and accountability that AI cannot reliably supply. Compilation, formatting, and date-filtering are structural tasks well-suited to automation.

Q39. What is the most accurate description of "prompt injection" risk in agentic systems?

- A. Too many prompts in a session cause the agent to lose context.
- B. Malicious content in external sources manipulates the agent's behaviour.**
- C. A long prompt reduces the agent's ability to follow instructions.
- D. Users inject conflicting instructions that confuse the agent.

Correct Answer: B

Explanation

In agentic systems that read external content — web pages, emails, documents — malicious instructions embedded in that content can override the agent's original instructions. This is prompt injection: content becomes commands, hijacking the agent's actions.

Q40. Which human task is hardest to automate with current AI agents?

- A. Exercising judgment in ambiguous, high-stakes situations.**
- B. Reading and summarising multiple long documents.
- C. Scheduling meetings across complex calendar constraints.
- D. Categorising incoming requests by topic and urgency.

Correct Answer: A

Explanation

Judgment in genuinely ambiguous situations — where context, values, stakes, and relationships all matter — remains the hardest task for AI agents. Reading, scheduling, and categorisation are all tasks that current AI can handle reliably with appropriate setup.

Q41. You build an agent that automates customer onboarding. After deployment, what is the most important ongoing activity?

- A. Training the agent on new customer data to improve over time.
- B. Expanding the agent's scope to handle more onboarding steps.
- C. Monitoring outputs for errors and updating the agent as processes change.**
- D. Reducing human oversight as the agent's performance proves reliable.

Correct Answer: C

Explanation

Deployed agents require ongoing monitoring — errors surface over time, processes change, and edge cases accumulate. Reducing human oversight as the agent "proves reliable" is the exact opposite of good governance; errors often emerge after initial performance looks stable.

Q42. A well-designed AI workflow for document review should:

- A. Route all documents to AI first, then humans only if AI rejects them.
- B. Achieve 100% AI coverage to eliminate the need for human review.
- C. Randomise 10% of documents to human review as a quality check.
- D. Flag uncertain cases for human review rather than making all decisions autonomously.**

Correct Answer: D

Explanation

Uncertainty-based routing — where AI handles confident cases and flags ambiguous ones for humans — is the correct design. Routing only AI-rejected documents means humans only see documents AI has already decided are problems. Random sampling is too blunt; 100% AI coverage is too risky.

Q43. What is the most important constraint to define before deploying any autonomous AI agent?

- A. The maximum number of API calls the agent may make per hour.
- B. What irreversible actions the agent is and is not permitted to take.**
- C. Which AI model the agent should use for each decision type.
- D. The acceptable error rate for the agent's primary function.

Correct Answer: B

Explanation

Irreversible actions — sending messages, making bookings, modifying records, spending money — are the most critical constraint to define. These are the actions that cannot be undone if the agent makes an error, and they require the most conservative permission design.

Q44. You notice that an AI agent has started taking an action you did not design it to take. The first response is:

- A. Monitor the unauthorised action to see if it is producing good outcomes.
- B. Update the agent's instructions to include the new action formally.
- C. Stop the agent, investigate the trigger, and restore intended boundaries.**
- D. Report the unexpected behaviour to the AI model provider.

Correct Answer: C

Explanation

Unexpected autonomous behaviour in an agent must be stopped and investigated immediately — not monitored or formalised without understanding the cause. The agent has deviated from its intended design, which is a safety incident regardless of whether the action seems beneficial.

Q45. Which statement best describes the current state of AI agents in professional settings?

- A. Powerful for well-defined tasks but requiring careful human oversight.**
- B. Ready to operate autonomously across most professional workflows.
- C. Still experimental and not yet reliable enough for production use.
- D. Limited to creative tasks and not yet applicable to operational work.

Correct Answer: A

Explanation

AI agents in professional settings are genuinely powerful for well-scoped, well-defined tasks — but they still require meaningful human oversight to catch errors, handle edge cases, and make judgment calls. Neither "ready for autonomous operation" nor "still experimental" is accurate.

SET

9

Level Up — Evaluating AI Output Quality

Set 9 of 10 · 15 Questions · AI Generalist

Q46. You receive an AI output that is well-written and confidently stated. The most important quality check is:

- A. Whether the output follows the format you requested.
- B. Whether the writing style matches your target audience.
- C. Whether the output is within the requested length range.
- D. Whether the specific claims are accurate and verifiable.**

Correct Answer: D

Explanation

Writing quality and confidence are orthogonal to factual accuracy — AI produces well-written hallucinations regularly. The most important check is whether specific claims can be traced and verified, not how good the output looks on the surface.

Q47. What is the most reliable method for evaluating an AI-generated strategy document?

- A. Run the document through a second AI tool for comparison.
- B. Apply your own domain expertise to challenge each key recommendation.**
- C. Ask Claude to critique its own strategy document.
- D. Share it with peers and use their reactions as quality signal.

Correct Answer: B

Explanation

Domain expertise applied critically — asking "does this recommendation make sense given what I know about this domain, this market, and this organisation?" — is the most reliable quality filter. Second AI tools share the same structural limitations; self-critique is unreliable; peer reactions vary by expertise.

Q48. An AI output is factually correct but logically inconsistent. Which statement is true?

- A. Logical consistency and factual accuracy always align in AI outputs.
- B. A logically inconsistent output is always also factually incorrect.
- C. Individual facts can be accurate while the overall argument is flawed.**
- D. Logical consistency is not relevant for factual AI outputs.

Correct Answer: C

Explanation

AI can produce individual facts that are accurate while constructing an argument with logical gaps, contradictions, or invalid inferences. Factual accuracy and logical consistency are independent dimensions that must be evaluated separately.

Q49. You are evaluating two AI outputs on the same task. Output A is shorter and direct; Output B is longer and thorough. How do you choose?

- A. Evaluate which better achieves the task objective, regardless of length.**
- B. Choose A — shorter outputs reflect more confident and accurate AI.
- C. Choose B — more thorough coverage indicates higher effort.
- D. Choose the output with fewer hedging phrases and qualifications.

Correct Answer: A

Explanation

Output quality is determined by task fit, not length. Longer outputs can be padding; shorter outputs can be incisive. Evaluate whether each output actually achieves the stated objective — then choose based on that, not surface signals.

Q50. What is "hallucination detection" in practical professional terms?

- A. A software tool that flags AI outputs with low confidence scores.
- B. An AI feature that marks statements the model is uncertain about.
- C. The practice of asking Claude to check its own outputs for errors.
- D. The process of verifying specific claims against primary sources.**

Correct Answer: D

Explanation

In practical professional terms, hallucination detection means developing the habit of identifying specific verifiable claims in AI output and tracing each one to a primary source. It is a human practice, not a software feature — AI self-assessment of confidence is unreliable.

Q51. Which type of AI error is most dangerous in a professional context?

- A. A vague, hedged answer that provides no actionable information.
- B. An answer that misunderstands the question and goes off-topic.
- C. A confident, specific, wrong fact that is hard to detect without expertise.**
- D. A response that is too long and includes irrelevant information.

Correct Answer: C

Explanation

The most dangerous errors are confident wrong facts about specific things — a fabricated statistic, wrong date, or incorrect attribution — because they are hardest to detect without domain expertise and most likely to be used as-is in downstream work.

Q52. You are building an AI quality scoring rubric for your team. Which dimension should be weighted highest?

- A. Factual accuracy of specific, verifiable claims.**
- B. Fluency and grammatical correctness of the output.
- C. Adherence to the specified format and structure.
- D. Completeness and coverage of relevant topics.

Correct Answer: A

Explanation

Factual accuracy of verifiable claims is the highest-stakes dimension — wrong facts with good grammar and perfect formatting are the worst outcome. Fluency, format, and completeness matter but are correctable without expertise; factual errors may require domain knowledge to detect.

Q53. An AI output includes a reference to a study with specific findings. Your first action is:

- A. Add a footnote noting the source is AI-generated.
- B. Search for the actual study and verify it exists and says what AI claims.**
- C. Proceed if the findings align with your prior knowledge.
- D. Ask Claude for the study's DOI or URL for citation.

Correct Answer: B

Explanation

AI-fabricated academic references are common — plausible-sounding author names, journal names, and findings that do not exist. Always search for the actual study independently. Asking Claude for the DOI produces another potential hallucination.

Q54. What is the most useful signal that an AI output is at risk of being low quality?

- A. The output is shorter than expected for the task.
- B. The output uses hedging language like "generally" or "may."
- C. The task required precise, specific, or recent information.**
- D. Claude took longer than usual to generate the response.

Correct Answer: C

Explanation

Tasks requiring precise, specific, or recent information are structurally higher risk — these are the conditions under which hallucination and error most commonly occur. Hedging language is actually a positive signal of appropriate uncertainty; length and speed are not quality indicators.

Q55. A team member argues that if AI output sounds professional, it must be good enough to use. The most precise counter is:

- A. AI writing quality varies too much to be relied on for professional work.
- B. Professional-sounding AI output still needs to be reformatted for clients.
- C. Professional quality is subjective and AI cannot achieve it reliably.
- D. Professional writing quality and factual accuracy are independent dimensions.**

Correct Answer: D

Explanation

The most precise correction is that fluency and accuracy are independent. AI produces fluent, professional-sounding text whether the content is accurate or not. Professional quality of writing is not a proxy for professional quality of information.

Q56. You ask Claude to rate the quality of its own output on a 1-10 scale. How useful is this rating?

- A. Very limited — AI self-assessment is not reliably calibrated.**
- B. Very useful — it identifies outputs the AI is uncertain about.
- C. Useful as a baseline — compare ratings across multiple outputs.
- D. Fully reliable — Claude knows its own capability boundaries well.

Correct Answer: A

Explanation

AI self-assessment of quality is not reliably calibrated — Claude can rate a hallucinated output highly because it cannot detect its own errors. Use self-assessment as a weak signal at best, never as a primary quality gate.

Q57. What is the correct response when you spot a factual error in an AI output mid-task?

- A. Re-run the entire prompt to get a fresh error-free output.
- B. Correct the error yourself and continue — do not re-prompt for a full rewrite.**
- C. Ask Claude to find and fix its own factual errors.
- D. Flag the error and submit the task anyway for time efficiency.

Correct Answer: B

Explanation

When you spot a specific error and know the correct information, fix it directly and continue. Re-running the prompt produces a new output that may have different errors; asking Claude to self-correct is unreliable; submitting with known errors is not acceptable.

Q58. Which output review habit has the highest return on time invested?

- A. Reading every word of every AI output before using it.
- B. Running outputs through grammar and spell-check tools.
- C. Comparing AI output to what you expected before reading it.
- D. Focusing review on specific figures, citations, and time-sensitive claims.**

Correct Answer: D

Explanation

Focusing review on specific figures, citations, and time-sensitive claims is the highest-leverage habit — these are the highest risk elements and the ones most likely to cause professional harm if incorrect. Reading every word of everything is exhaustive but less targeted.

Q59. What does "output calibration" mean in professional AI use?

- A. Developing an accurate sense of where AI is reliable versus where it fails.**
- B. Adjusting the AI model's settings to improve output quality.
- C. Testing AI outputs against a calibration dataset for accuracy.
- D. Calibrating response length to match the task requirements.

Correct Answer: A

Explanation

Output calibration is the professional judgment skill of knowing, across task types, where AI tends to perform well and where it tends to fail — so you can adjust your review depth and verification effort accordingly. It develops through systematic experience.

Q60. A manager receives two reports on the same topic: one from a human analyst and one from AI. The manager cannot distinguish them. What does this prove?

- A. It proves AI has reached human-level analytical capability.
- B. It proves the human analyst should be reviewed for underperformance.
- C. It proves only that AI can produce human-quality prose, not that the content is accurate.**
- D. It proves AI is suitable to replace analysts on this type of task.

Correct Answer: C

Explanation

Indistinguishability at the prose level proves only that AI can write like a human analyst. It says nothing about the accuracy of the analysis, the validity of the data sources, or whether the conclusions are sound. Quality of writing is not quality of thinking.

SET

10

Level Up — AI Leadership and Cultural Change

Set 10 of 10 · 15 Questions · AI Generalist

Q61. A team member says "I feel stupid using AI — it does things faster than I can." The most effective response from a team leader is:

- A. "AI handles speed; your judgment determines if the output is actually right."
- B. "You should use AI more so you feel more comfortable with it."
- C. "AI is a tool, not a measure of intelligence — do not compare yourself."
- D. "Everyone feels this way initially — the feeling passes with more use."

Correct Answer: A

Explanation

The most accurate and empowering reframe is to distinguish speed (AI's strength) from judgment (the human's irreplaceable contribution). This positions the team member as the critical quality check rather than a slower competitor.

Q62. What is the most important organisational condition for successful AI adoption?

- A. Executive mandate requiring all employees to use AI tools daily.
- B. A large budget for premium AI tool access across the team.
- C. Dedicated AI specialists who manage tools on behalf of others.
- D. **Psychological safety to experiment, fail, and share learning openly.**

Correct Answer: D

Explanation

Psychological safety — where people can experiment, make mistakes, and share what they learned without fear of judgment — is the cultural condition that enables genuine learning and adoption. Mandates, budgets, and specialists support adoption but cannot create genuine engagement.

Q63. You want to build an AI-fluent culture in your organisation. The most effective first action is:

- A. Train everyone on AI fundamentals before allowing any tool use.
- B. Create visible quick wins by using AI on real tasks with the team.**
- C. Hire a Chief AI Officer to lead the cultural transformation.
- D. Survey employees on their current AI tool usage and attitudes.

Correct Answer: B

Explanation

Culture change in AI adoption is driven by visible results on real work — not training sessions or surveys. Quick wins demonstrate value, create advocates, and lower the psychological barrier to experimentation. Abstract training without practical use rarely sticks.

Q64. A team member uses AI extensively and produces more output than their colleagues. What is the most important performance question to ask?

- A. "Are they using AI tools that are approved by the organisation?"
- B. "Are they documenting their AI prompts for the team to learn from?"
- C. "Is the quality and accuracy of their AI-assisted output meeting the standard?"**
- D. "Are they becoming less capable at the underlying skills themselves?"

Correct Answer: C

Explanation

Volume is not the metric — quality is. High output volume from AI use is only valuable if quality and accuracy are maintained. A team member who produces twice as much work with 60% accuracy is producing more errors, not more value.

Q65. An AI-sceptical colleague says: "AI will make everyone's thinking lazy and shallow." The most nuanced response is:

- A. "That's a genuine risk that depends on how AI is used — passive use atrophies skills; active use can sharpen thinking."**
- B. "That concern is valid — AI is best kept out of professional workflows."
- C. "That concern is outdated — research shows AI improves critical thinking."
- D. "AI only makes thinking lazy for people who already have that tendency."

Correct Answer: A

Explanation

The nuanced position acknowledges the real risk (passive AI use does reduce the thinking effort required) while pointing to the countervailing possibility (active, critical use of AI can actually sharpen judgment by forcing you to evaluate and direct output).

Q66. What is the most effective way to transfer AI knowledge within a team?

- A. Sending team members to an external AI certification programme.
- B. Holding a monthly AI tool demo and exploration session.
- C. Sharing specific prompts, examples, and workflows that produced good results.**
- D. Requiring team members to document every AI interaction.

Correct Answer: C

Explanation

Knowledge transfer in AI is most effective through concrete, contextual examples — actual prompts that worked, specific outputs that failed, and workflow designs that proved reliable. Certifications, demos, and documentation requirements are less direct.

Q67. You are leading an AI adoption initiative. What does failure look like most commonly?

- A. Employees adopt AI faster than the infrastructure can support.
- B. AI produces outputs that require too much human review.
- C. Senior leadership does not approve the AI tools selected.
- D. AI is used superficially without changing underlying workflows.**

Correct Answer: D

Explanation

The most common failure mode is surface adoption — people use AI tools but in ways that do not change how work is fundamentally done or improve outcomes. AI becomes a checkbox rather than a genuine capability upgrade.

Q68. A team member asks: "Should I tell clients I used AI to do this work?" The most accurate guidance is:

- A. "Always disclose — transparency with clients is non-negotiable."
- B. "Follow your profession's disclosure norms and your client's stated preferences."**
- C. "Never disclose — clients don't need to know your production tools."
- D. "Disclose only if the client specifically asks about your process."

Correct Answer: B

Explanation

AI disclosure in professional services is not universally standardised — it depends on professional codes of conduct, industry norms, explicit client agreements, and relationship context. The correct guidance is to navigate these factors, not apply a blanket rule.

Q69. What is the most meaningful way to measure a team's AI maturity?

- A. How many AI tools they use across different task categories.
- B. How quickly they adopted new AI tools after they became available.
- C. Whether they can identify and correct AI failures in their own workflows.**
- D. Whether they have formal AI policies documented in writing.

Correct Answer: C

Explanation

AI maturity is best measured by the ability to critically engage with AI — catching errors, improving processes, and governing use responsibly. Breadth of tool use, adoption speed, and documentation are all compatible with low-maturity AI adoption.

Q70. A colleague presents a highly compelling AI-generated analysis in a strategy meeting. No one questions it. What is the risk?

- A. The group may act on hallucinated or unverified data.**
- B. The colleague may be over-relying on AI for strategy work.
- C. The group will expect AI-generated analysis in all future meetings.
- D. The AI tool used may not be appropriate for strategic analysis.

Correct Answer: A

Explanation

The most concrete risk is acting on hallucinated or unverified data. A compelling AI output that goes unchallenged may contain fabricated statistics, incorrect market data, or flawed logic that no one tested — with strategic decisions made on that faulty basis.

Q71. What does it mean to "govern AI use" in an organisational context?

- A. Controlling who in the organisation has access to AI tools.
- B. Defining which tasks AI is used for, how outputs are reviewed, and what standards apply.**
- C. Monitoring AI tool usage volumes and costs across teams.
- D. Ensuring all AI use is approved by senior leadership in advance.

Correct Answer: B

Explanation

AI governance means defining the scope of appropriate use, quality standards for outputs, review processes, and accountability — not just access control or monitoring. It is a policy and process framework, not a technical gating mechanism.

Q72. How should an AI Generalist respond when asked to evaluate whether a task is appropriate for AI?

- A. Check whether the task type is on the company's approved AI use list.
- B. Ask which AI tool the requester has in mind for the task.
- C. Estimate the time saved versus manually completing the task.
- D. Assess what can go wrong, how it will be caught, and whether the stakes justify the risk.**

Correct Answer: D

Explanation

Appropriate use evaluation requires a risk-based analysis: what are the failure modes, are they catchable, and do the stakes justify the current level of AI capability and oversight? Approved lists, tool choice, and time savings are secondary considerations.

Q73. Which statement about AI and professional expertise is most accurate?

- A. AI reduces the need for expertise by providing expert-level outputs.
- B. AI and expertise are interchangeable for most professional tasks.
- C. AI amplifies expertise — experts get more from AI than novices do.**
- D. AI makes expertise less valuable because it democratises information.

Correct Answer: C

Explanation

Expert users get more from AI because they know what to ask, how to evaluate the output, when the AI is wrong, and how to direct it toward what matters. Novices get generic outputs; experts get targeted, high-quality drafts they can critically evaluate.

Q74. What is the most honest advice to give a student entering the workforce in an AI-transformed environment?

- A. "Learn to code AI systems — technical skills are the only durable advantage."
- B. "Develop strong judgment skills — AI handles speed; you add direction and quality control."**
- C. "Use AI constantly to maximise productivity in your early career."
- D. "Focus on skills AI cannot do rather than on working alongside it."

Correct Answer: B

Explanation

The most honest and durable advice is to invest in judgment — the ability to direct AI, evaluate its output, and take responsibility for final quality. This is the skill that compounds over time and cannot be automated. Technical skills and AI use are valuable but not sufficient on their own.

Q75. Menler's AI Generalist Level Up bank tests you at the mid-fellowship stage. What does "level up" indicate?

- A. You have mastered every AI tool used in the Generalist Fellowship.
- B. You are ready to teach AI fundamentals to students below your level.
- C. You have completed the technical AI engineering prerequisites.
- D. You are applying AI judgment to complex, ambiguous professional situations.**

Correct Answer: D

Explanation

"Level Up" in Menler's framework means progressing from foundational awareness to applied professional judgment — using AI strategically, governing it responsibly, and directing it toward complex real-world situations. It is not about tool mastery, teaching readiness, or technical prerequisites.

END OF QUESTION BANK

*Your turning
point starts
here.*