

menler

Your turning point in the AI era.

AI APTITUDE QUESTION BANK

AI

Judgment

About This Question Bank

A bank for practitioners making independent professional AI decisions — evaluating output quality, detecting hallucination, choosing when and how to prompt, understanding failure modes, and taking accountability seriously.

5

Sets

15

Questions per Set

75

Total Questions

SET GUIDE

Set 1 · Output Quality

Output quality — what good looks like and how to evaluate it

Set 2 · Hallucination and Verification

Hallucination and verification — detecting and preventing AI errors

Set 3 · Prompting Decisions

Prompting decisions — when and how to prompt for best results

Set 4 · When AI Fails

When AI fails — failure modes, brittleness, and systematic errors

Set 5 · Professional Accountability

Professional accountability — liability, disclosure, and judgment

Each question has one correct answer. Select the best answer from the options provided.

Correct answers and explanations follow each question.

SET

1

Output Quality

Set 1 of 5 · 15 Questions · AI Judgment

Q1. What is the most reliable signal that an AI output is high quality?

- A. The output is fluent, grammatically correct, and well-structured.
- B. Specific claims are traceable to verifiable sources or logical derivations.**
- C. The output is longer and more detailed than what was requested.
- D. The model expressed high confidence throughout the response.

Correct Answer: B**Explanation**

Fluency, length, detail, and expressed confidence are all independent of accuracy. An AI can produce beautifully written, highly confident, detailed text that is factually wrong. Traceability of claims to verifiable sources or sound logical derivation is the only reliable quality signal.

Q2. What does "output calibration" mean in professional AI use?

- A. Adjusting the AI model's temperature and sampling settings for each task.
- B. Running AI outputs through a secondary model to score their quality.
- C. Calibrating output length to match the expected format for each task type.
- D. Developing accurate judgment about where AI outputs are reliable versus where they fail.**

Correct Answer: D**Explanation**

Output calibration is the practitioner skill of knowing, based on task type and output characteristics, how much to trust a given AI output — and when deep verification is necessary. It develops through systematic experience and is the core judgment skill for safe professional AI use.

Q3. You receive an AI output that answers a question you did not ask while ignoring the one you did. The cause is:

A. The prompt was ambiguous enough that the model interpreted a different question.

B. The model lacked the knowledge to answer the intended question.

C. The model's safety filters redirected the response to a safer topic.

D. The model exceeded its context window before processing the full question.

Correct Answer: A

Explanation

When AI answers a different question than asked, the most common cause is prompt ambiguity — the question could reasonably be read in a different way, and the model selected a different interpretation. The fix is clarifying the question, not assuming the model is incapable or filtered the response.

Q4. What distinguishes a "correct" AI output from a "useful" one?

A. Correct outputs are longer and more complete; useful outputs are shorter and actionable.

B. Correct outputs are verifiable; useful outputs are those the user decides to act on.

C. Correctness means factually accurate; usefulness means relevant to the actual task need.

D. There is no meaningful distinction — correct outputs are always useful.

Correct Answer: C

Explanation

An output can be factually correct but useless (answering a different level of detail than needed, in the wrong format, for the wrong audience). A useful output addresses the actual need — the right information, at the right depth, in the right format — even if it requires the user to verify the factual claims.

Q5. What is the most effective quality control step for AI-generated reports before distribution?

- A. Running a grammar and spell-check tool to catch surface errors.
- B. Human expert review of all specific claims, figures, and recommendations.**
- C. Asking the AI to review its own report for accuracy before finalising.
- D. Comparing the report structure against a template for format compliance.

Correct Answer: B

Explanation

Human expert review of specific claims, figures, and recommendations is the only quality control step that catches factual errors, logical flaws, and contextual mismatches. Grammar checking, AI self-review, and format compliance checks cannot catch factual or logical errors — they are secondary to this primary check.

Q6. What is "fluency illusion" in AI output evaluation?

- A. Well-written text gives readers false confidence that the content is accurate.**
- B. AI outputs that use complex vocabulary to obscure simple concepts.
- C. The experience of reading AI text as if a human wrote it.
- D. The tendency for AI to repeat fluency-related phrases in different contexts.

Correct Answer: A

Explanation

The fluency illusion describes how skilled writing — clear structure, appropriate vocabulary, confident tone — creates an impression of authority and correctness that may be entirely unwarranted. AI exploits this cognitive bias by generating consistently fluent text regardless of factual accuracy.

Q7. What is the most important dimension to evaluate in an AI output for a professional client deliverable?

- A. Length and detail appropriate for the client's seniority level.
- B. Brand voice consistency with the organisation's communication standards.
- C. Completeness — covering all aspects of the topic without omissions.
- D. Factual accuracy of specific claims that will inform client decisions.**

Correct Answer: D

Explanation

Factual accuracy is the primary dimension for professional deliverables because clients use deliverables to make decisions. Wrong facts lead to wrong decisions. Length, voice, and completeness are all secondary to ensuring what is stated is correct.

Q8. What does "output specificity" indicate about AI quality on knowledge tasks?

- A. Specific outputs are always more accurate because they demonstrate depth of knowledge.
- B. Specific outputs indicate the model has retrieved relevant information from the web.
- C. Specific, concrete outputs are higher risk for error than general summaries.**
- D. Output specificity is independent of accuracy and provides no quality signal.

Correct Answer: C

Explanation

Specific outputs — precise figures, exact names, specific dates, verbatim quotations — carry higher error risk than general summaries. AI generates specific details with the same confidence regardless of whether those details are accurate or fabricated. Specificity requires higher verification standards, not less.

Q9. A colleague says "the AI output looks great, let's use it." What is missing from this evaluation?

- A. Assessment of whether the content is factually accurate, not just whether it looks good.**
- B. A review of the output's length and format relative to the intended audience.
- C. Comparison with a human-written version to assess relative quality.
- D. A review of the prompt used to generate the output for quality issues.

Correct Answer: A

Explanation

"Looks great" is a fluency and format assessment — the output reads well. Missing is an accuracy assessment — whether the specific claims, figures, and reasoning are actually correct. Professional AI use requires distinguishing between good appearance and verified accuracy.

Q10. What is the most common quality failure in AI-generated marketing copy?

- A. Grammar errors that undermine the professional appearance of the copy.
- B. Excessive length that exceeds the intended word count or format requirements.
- C. Overly promotional language that violates advertising standards.
- D. Generic, pattern-following content that lacks the specific voice and differentiation of the brand.**

Correct Answer: D

Explanation

AI marketing copy defaults to the statistically average marketing voice — which is generic, pattern-following, and interchangeable. The most common failure is copy that is technically correct and readable but lacks the specific brand voice, cultural nuance, and differentiated perspective that makes marketing effective.

Q11. What is the correct response when an AI output is good enough for one part of the task but wrong for another?

- A. Reject the entire output and regenerate from scratch.
- B. Use the good parts as-is and fix the wrong parts yourself.**
- C. Ask the AI to fix the wrong parts by identifying them in a follow-up prompt.
- D. Average the two parts together to produce a balanced acceptable output.

Correct Answer: B

Explanation

Surgical editing — keeping what is correct, fixing what is wrong — is the most efficient approach. Regenerating from scratch wastes the 80% that was correct; asking the AI to self-correct on identified errors is unreliable; averaging makes no sense. Fix the wrong parts directly and move on.

Q12. What is the best indicator that an AI output has been appropriately personalised for a specific audience?

- A. It is shorter and more direct than a general-purpose response would be.
- B. It avoids technical jargon that might confuse a non-expert audience.
- C. It references specific context, terminology, and concerns relevant to that audience.**
- D. It was generated using a role prompt specific to the target audience.

Correct Answer: C

Explanation

True personalisation is evidenced by specific, audience-relevant content — their context, their vocabulary, their specific concerns and goals. Generic brevity, jargon avoidance, and role prompting can all contribute, but they are not sufficient unless the output actually contains specific relevant content.

Q13. What is "output drift" in a multi-step AI task?

- A. The model's output quality decreasing as the context window fills.
- B. Inconsistency in output style between different sections of a long document.
- C. The model taking longer per output token as context length increases.
- D. The AI progressively deviates from the original brief as the task proceeds.**

Correct Answer: D

Explanation

Output drift occurs in long tasks or conversations where the model gradually loses track of the original objectives, constraints, or style requirements — each output drifts slightly from the brief until the final output is substantially different from what was originally requested.

Q14. What is the most important thing to check when comparing two AI outputs on the same task?

- A. Which output is longer, as length often correlates with thoroughness.
- B. Which output better fulfils the actual task objective, not just which reads better.**
- C. Which output uses more varied vocabulary and sentence structures.
- D. Which output was generated with the higher temperature setting.

Correct Answer: B

Explanation

Output comparison should always evaluate task fulfilment — which output actually does what was needed. Reading quality, length, vocabulary variety, and generation settings are all irrelevant if neither output correctly addresses the task objective.

Q15. What is the most appropriate quality standard for AI-generated content used in training materials?

- A. Every factual claim must be accurate, as learners will treat training content as authoritative.**
- B. The content must be engaging and accessible, as these are the primary goals of training.
- C. The content must pass AI detection tools to ensure it does not appear AI-generated.
- D. The content must be reviewed by the L&D team for format and length compliance.

Correct Answer: A

Explanation

Training materials create beliefs and shape decisions. Learners treat authoritative sources as accurate. AI-generated training content with factual errors produces incorrect beliefs at scale — every person trained on wrong content carries that error into their professional work. Factual accuracy is the non-negotiable standard.

SET

2

Hallucination and Verification

Set 2 of 5 · 15 Questions · AI Judgment

Q16. What is "hallucination" in a language model?

- A. The model seeing patterns or meanings that are not present in the input.
- B. A failure mode where the model refuses to answer factual questions.
- C. **Generating text that is confident and plausible but factually incorrect.**
- D. Generating repetitive or looping text that does not converge to an answer.

Correct Answer: C**Explanation**

Hallucination in language models refers to generating confident, specific, well-structured text that contains factually incorrect information — fabricated citations, wrong statistics, invented product features, or nonexistent people. The model does not "know" it is wrong; it generates the most probable-sounding text.

Q17. Why is hallucination a structural property of language models rather than a fixable bug?

- A. **Models predict statistically probable text, which can produce plausible-sounding errors.**
- B. Hallucination is caused by insufficient training data and grows less common with more data.
- C. Hallucination is a deliberate design choice to make models more creative.
- D. Hallucination occurs only when models are asked questions outside their training domain.

Correct Answer: A**Explanation**

Language models are optimised to generate the most probable next tokens — not to retrieve verified facts. This means plausible-sounding incorrect information is an inherent output of the architecture when the model's training probability distribution happens to peak at a wrong answer. Reducing hallucination is an active area of research, but eliminating it entirely requires fundamental changes to how these systems work.

Q18. What is the most reliable method for detecting hallucination in an AI output?

- A. Running the output through a second AI model for cross-validation.
- B. Asking the AI to rate its own confidence on each claim in the output.
- C. Checking whether the output is internally consistent and logically coherent.
- D. Tracing every specific claim to a verifiable primary source.**

Correct Answer: D

Explanation

Only tracing specific claims to verifiable primary sources actually confirms accuracy. Cross-checking with another AI compounds the same hallucination tendency. AI self-confidence is unreliable — models express equal confidence on correct and hallucinated facts. Internal consistency does not ensure external accuracy.

Q19. An AI output cites a specific academic paper with author names, journal, and year. What should you do before using this citation?

- A. Use the citation as-is — AI cites sources accurately when provided with a reference.
- B. Search for and verify the actual paper exists and says what the AI claims.**
- C. Add a note that the citation was AI-generated for full transparency.
- D. Ask the AI to confirm the citation is correct by re-checking its training data.

Correct Answer: B

Explanation

AI commonly fabricates citations that sound completely real — plausible author names, journals, years, and abstracts. Never use an AI-generated citation without finding and reading the actual paper in a verified academic database. Asking the AI to confirm compounds the problem; noting AI provenance does not make the citation real.

Q20. What is "confabulation" and how is it related to hallucination?

- A. Confabulation refers to AI-generated content that confuses different topics into one answer.
- B. Confabulation is a more severe form of hallucination where the model invents entire narratives.
- C. Confabulation is producing false memories with confidence — hallucination is the AI equivalent.**
- D. Confabulation is the technical term for hallucination in medical AI contexts.

Correct Answer: C

Explanation

Confabulation (from neuropsychology) describes a brain phenomenon where a patient fills gaps in memory with fabricated but sincerely believed information, without intent to deceive. This is the closest human cognitive analogue to LLM hallucination — the model does not "know" it is wrong.

Q21. What is the highest-risk task type for AI hallucination?

- A. Generating creative fiction where facts do not need to be accurate.
- B. Summarising a long document provided directly in the context.
- C. Reformatting data from one structured format to another.
- D. Generating specific, verifiable details like dates, statistics, people, and citations.**

Correct Answer: D

Explanation

Specific verifiable details — exact dates, precise statistics, named individuals, bibliographic citations — are where hallucination is most dangerous. The AI generates these with the same confidence as correct information, but they are exactly the details that get fact-checked and used in consequential decisions.

Q22. What is the correct way to prompt Claude to reduce hallucination risk on factual questions?

- A. Instruct Claude to acknowledge uncertainty and cite its reasoning rather than assert facts confidently.**
- B. Use a higher temperature to make Claude more creative and less likely to commit to wrong facts.
- C. Ask Claude to provide three possible answers and select the most common one.
- D. Provide Claude with a specific factual question and ask for the single most accurate answer.

Correct Answer: A

Explanation

Instructioning Claude to acknowledge uncertainty and explain its reasoning helps surface where the model is generating from pattern rather than solid knowledge. Uncertainty acknowledgment is a weak signal but better than false confidence. Higher temperature increases randomness not accuracy; three answers compounds the problem; a direct factual question produces confident wrong answers.

Q23. What is "grounding" as a mitigation for hallucination?

- A. Adding a factual accuracy disclaimer to all AI-generated outputs.
- B. Providing the model with source documents to reason from rather than generating from training memory.**
- C. Training the model on larger and higher-quality datasets to reduce errors.
- D. Using retrieval-augmented generation to supplement the model's knowledge cutoff.

Correct Answer: B

Explanation

Grounding means providing the model with the relevant source material — documents, data, verified facts — and instructing it to reason from that material rather than from its training knowledge. Grounded outputs are traceable to the provided context; ungrounded outputs are generated from statistical memory. RAG is the most common implementation of grounding.

Q24. A developer finds that adding "check your work" to a prompt reduces errors. Why does this work?

- A. It signals to the model that the task requires higher accuracy and changes the generation weights.
- B. It triggers a different model pathway designed for accuracy-critical tasks.
- C. It doubles the effective context the model uses by prompting it to re-read the input.
- D. The instruction activates a self-review step that catches some generation errors.**

Correct Answer: D

Explanation

Self-review instructions cause the model to generate additional checking steps that can catch some logical errors, inconsistencies, and obvious mistakes. They do not change model weights or trigger different pathways — they simply add a generation step that functions like a scratchpad review. This is a genuine but imperfect mitigation.

Q25. What is the "verification habit" and why is it the most important AI skill for professionals?

- A. Asking the AI to verify its own outputs before accepting them as final.
- B. Running all AI outputs through a third-party fact-checking service.
- C. Systematically checking AI-generated claims against primary sources before acting on them.**
- D. Having a second team member review all AI-generated content before use.

Correct Answer: C

Explanation

The verification habit is the discipline of independently checking the specific factual claims in AI outputs against verifiable sources before acting on them. It is the most important skill because AI generates confident wrong information at a rate that makes professional reliance without verification a systematic source of error.

Q26. What does "source attribution" mean in an AI output context, and why is it insufficient on its own?

- A. Attribution identifies who is responsible for the information in the output.
- B. Citing a source does not guarantee the source exists or says what the AI claims.**
- C. Source attribution is a copyright requirement for AI-generated content.
- D. Attribution is sufficient if the source cited is a peer-reviewed academic journal.

Correct Answer: B

Explanation

AI can generate source attributions — journal names, author names, article titles — that are entirely fabricated. Attribution that has not been independently verified against the actual source is not evidence of accuracy. This is why AI-generated citations must always be verified in the original source, regardless of how authoritative they sound.

Q27. You are using AI to write a report on a competitor's market position. Which facts must you independently verify?

- A. All specific claims about the competitor's products, market share, revenue, and strategy.**
- B. Only the financial figures, since these are the most commonly hallucinated.
- C. Only claims that conflict with your prior understanding of the competitor.
- D. No verification is needed if the AI was given accurate web search results to work from.

Correct Answer: A

Explanation

All competitor-specific facts — products, features, pricing, market share, strategic announcements, personnel — must be verified from primary sources. AI confidently generates incorrect competitor details. Even when given web search results, AI can misattribute, misquote, or confuse information across sources.

Q28. What is the most important principle for professional use of AI outputs in high-stakes contexts?

- A. Always disclose that AI was used to produce any professional output.
- B. Only use AI for tasks where errors are immediately obvious and easily corrected.
- C. Never use an AI output for a consequential decision without verifying the critical facts.**
- D. Always have a second AI model cross-check the first model's output before use.

Correct Answer: C

Explanation

In high-stakes contexts, the professional responsibility is clear: verify the critical facts before acting. Disclosure, ease of error detection, and AI cross-checking are all supporting practices but none substitute for the fundamental requirement of verification when decisions have significant consequences.

Q29. What is the most reliable indicator that an AI output is low risk for hallucination?

- A. The output is shorter, as shorter outputs have fewer opportunities for error.
- B. The model expressed appropriate uncertainty about some claims in the output.
- C. The task topic is well-represented in the model's training data.
- D. The task is structural or logical rather than factual — formatting, calculation, or reasoning.**

Correct Answer: D

Explanation

Structural and logical tasks — reformatting data, solving well-defined problems, transforming information from one structure to another — have verifiable correct answers and do not require the model to generate specific factual details from memory. These are inherently lower-risk than knowledge-intensive tasks.

Q30. What is the most important thing a professional learns from consistently applying the verification habit?

- A. A comprehensive list of AI output types that are reliably accurate.
- B. A calibrated sense of which AI output types need more versus less verification.**
- C. The ability to identify hallucinations by reading style and tone alone.
- D. Confidence that AI outputs are reliable enough for most professional tasks.

Correct Answer: B

Explanation

Consistent verification builds calibration — experience-based judgment about where AI tends to be right and where it tends to fail. This allows professionals to direct their verification effort efficiently: deep verification where it matters, light spot-checking where AI is reliably strong.

SET

3

Prompting Decisions

Set 3 of 5 · 15 Questions · AI Judgment

Q31. What is the single most impactful element to add to a vague prompt to improve AI output quality?

- A. A specific description of what the output should look like when done well.**
- B. A longer, more detailed description of the topic to give the AI more context.
- C. Instructions for the AI to "think carefully" before responding.
- D. A request for the AI to ask clarifying questions before answering.

Correct Answer: A**Explanation**

Specifying what good output looks like — the format, level of detail, perspective, structure — is the highest-leverage addition to a vague prompt. Vague prompts produce average outputs; specific output criteria produce targeted outputs. Topic detail, generic quality instructions, and clarifying questions are all lower-leverage additions.

Q32. When should you use a "role prompt" (e.g., "act as a senior editor")?

- A. When you want the model to take on a persona and maintain it throughout the conversation.
- B. When you need the model to use domain-specific terminology in its responses.
- C. When the desired output benefits from a specific professional perspective or expertise level.**
- D. When you want to prevent the model from including caveats and qualifications.

Correct Answer: C**Explanation**

Role prompts activate patterns associated with that professional perspective in the model's training distribution. "Act as a senior editor" shifts the model toward more critical, polish-focused feedback. This is most valuable when the desired output genuinely benefits from a specific expert viewpoint, not just for terminology or to suppress caveats.

Q33. What is the most effective way to get a consistently structured output format from Claude?

- A. Use the phrase "always follow this format" in the system prompt.
- B. Provide the exact template or schema you want in the prompt, ideally with an example.**
- C. Ask Claude to generate the format itself before producing the content.
- D. Specify the desired format in general terms at the end of the prompt.

Correct Answer: B

Explanation

Templates and schemas with examples are the most effective format controls. They show Claude exactly what output should look like rather than describing it, which is always more reliable than a description. General format terms at the end of a prompt are less reliable than specific templates at the beginning.

Q34. What is "negative prompting" and when is it useful?

- A. Providing examples of bad outputs to help the model recognise what not to produce.
- B. Using a critical, evaluative tone in the prompt to encourage higher-quality responses.
- C. Removing positive examples from the prompt to force more creative outputs.
- D. Explicitly telling the model what to avoid in its output to prevent specific failure modes.**

Correct Answer: D

Explanation

Negative prompting tells the model what not to do — "do not include caveats," "avoid jargon," "do not repeat the question in your response." It is most useful for preventing specific recurring failures that are not prevented by positive instructions alone. Combined with positive examples, it gives the model a clear picture of the target output.

Q35. What does "few-shot prompting" mean and why does it work?

- A. Providing example input-output pairs in the prompt that demonstrate the desired behaviour.**
- B. Running the same prompt multiple times and selecting the best few outputs.
- C. Starting with a short, minimal prompt and adding detail until the output improves.
- D. Using only a few essential instructions rather than a long, complex prompt.

Correct Answer: A

Explanation

Few-shot prompting gives the model examples of the task — here is an input and here is the kind of output I want. The model uses these as patterns to follow rather than interpreting an abstract description. Examples are often more effective than long descriptions because they directly show the target output format and quality.

Q36. When is it better to use multiple simple prompts in sequence rather than one complex prompt?

- A. When you want to reduce API costs by keeping each individual call shorter.
- B. When the task has distinct stages that each benefit from focused, undivided attention.**
- C. When the task requires real-time user input at intermediate stages.
- D. Always — multiple simple prompts are always more reliable than one complex prompt.

Correct Answer: B

Explanation

Multiple prompts in sequence (chaining) are better when the task has genuinely distinct stages — extract, validate, synthesise — where each step's output quality benefits from the model having a focused, constrained context. One complex prompt works when the task is well-integrated and the model can handle all aspects in one pass.

Q37. What is the most common mistake people make when writing their first serious prompts?

- A. Making the prompt too long and overloading the model with irrelevant context.
- B. Using too formal a tone that makes the model produce overly academic responses.
- C. Not including enough examples for the model to understand the task.
- D. Being too vague about what they want and expecting the model to infer the specifics.**

Correct Answer: D

Explanation

Vagueness is the most common first-prompt mistake. Users assume the model can infer what "good" means for their specific task. Without explicit output criteria, the model produces its statistical average, which is rarely what the user wanted. The fix is specifying the exact output characteristics required.

Q38. What is the most reliable way to improve AI output quality through iteration?

- A. Regenerate the output with higher temperature to introduce variation.
- B. Add more context about the background and topic to the original prompt.
- C. Identify the specific element that fell short and give targeted correction instructions.**
- D. Ask the model to rate its own output and improve the lowest-rated elements.

Correct Answer: C

Explanation

Targeted correction is the most effective iteration strategy: "The analysis is too high-level — add specific examples for each of the three points" produces better results than vague regeneration or context additions. Specific, targeted feedback on exactly what to improve is the highest-leverage iteration approach.

Q39. What is "context injection" and why does it improve prompting outcomes?

- A. Injecting the user's previous conversation history to maintain continuity.
- B. Adding task-specific context that Claude lacks from training to make outputs more relevant.**
- C. Providing real-time data through a web search tool to update Claude's knowledge.
- D. Adding the user's identity information so Claude can personalise responses.

Correct Answer: B

Explanation

Context injection means including information Claude does not have by default — company-specific facts, proprietary terminology, specific client context, project constraints — directly in the prompt. This allows Claude to produce outputs relevant to the specific situation rather than generic outputs that could apply to any situation.

Q40. What is the most important thing to consider when deciding how much context to include in a prompt?

- A. Whether each piece of context actually affects the output — remove context that does not.**
- B. Including as much context as possible to ensure the model has everything it needs.
- C. Keeping context under 500 words to prevent the model from being distracted.
- D. Matching context volume to the model's context window to maximise performance.

Correct Answer: A

Explanation

Relevance is the most important consideration. Irrelevant context dilutes the model's attention from what matters. "Lost in the middle" research shows that adding unnecessary context can actually reduce output quality. Include context only if it would meaningfully change the output if the model had versus did not have it.

Q41. What is the difference between a "system prompt" and a "user prompt" in terms of prompting strategy?

- A. System prompts are longer and more detailed; user prompts are shorter and direct.
- B. System prompts are for operators; user prompts are only for end users.
- C. System prompts set persistent context and role; user prompts specify the immediate task.**
- D. System prompts define format; user prompts define the topic to cover.

Correct Answer: C

Explanation

The strategic distinction is persistence and authority: system prompts set the operational context — who Claude is, what constraints apply, what format to use — that applies to all subsequent interactions. User prompts provide the specific task or question within that context. Good prompt architecture separates persistent context from task-specific instructions.

Q42. What is "prompt sensitivity" and why does it matter for production systems?

- A. The degree to which a prompt activates the model's safety filters.
- B. The length threshold above which longer prompts stop improving output quality.
- C. A model's tendency to generate different outputs to the same prompt across sessions.
- D. Small changes in prompt wording can significantly change output quality or behaviour.**

Correct Answer: D

Explanation

Prompt sensitivity means small, seemingly innocuous wording changes — a different word, a reordered sentence, a different example — can substantially change output quality. This matters for production systems because prompts must be thoroughly tested before deployment and any changes must be re-tested to catch unexpected behaviour changes.

Q43. What is the most important test to run when deploying a prompt in a production system?

- A. Testing with the same input 10 times to measure output consistency.
- B. Testing with real and adversarial user inputs to identify edge cases and failure modes.**
- C. Testing with the ideal example input to confirm the prompt works correctly.
- D. Testing API latency to ensure the prompt length does not cause timeouts.

Correct Answer: B

Explanation

Production prompts will encounter inputs the developer never anticipated. Testing with real user inputs (sampled from actual use cases) and adversarial inputs (edge cases, boundary conditions, unusual formats) reveals failure modes that ideal-case testing misses entirely.

Q44. What is the best way to evaluate whether a prompt change improved or degraded output quality?

- A. Use the new prompt for a week in production and see if user complaints decrease.
- B. Ask Claude to compare its outputs from both prompt versions and judge which is better.
- C. Run both prompt versions on the same test set and compare outputs systematically.**
- D. Check whether the new prompt produces longer, more detailed outputs.

Correct Answer: C

Explanation

Systematic A/B evaluation on a consistent test set is the only reliable method for assessing prompt changes. Production feedback is too slow and confounded; Claude self-comparison is unreliable; output length is not a quality metric. A consistent test set allows objective, comparable measurement of both versions.

Q45. What is the most important principle for prompt design in a product used by non-technical users?

- A. The prompt must produce high-quality outputs even when user inputs are vague or unusual.**
- B. The prompt must be short enough that users can read and understand it themselves.
- C. The prompt must prevent users from changing Claude's behaviour through their messages.
- D. The prompt must include example queries that users can copy and adapt.

Correct Answer: A

Explanation

Non-technical users write imprecise, varied, and unexpected inputs. The system prompt must be robust enough to produce high-quality outputs across this full range — not just for well-formed inputs. Building for the average well-formed query while the prompt breaks on edge cases is a common product failure.

SET

4

When AI Fails

Set 4 of 5 · 15 Questions · AI Judgment

Q46. What is the most common reason AI fails at tasks it seemed capable of during testing?

- A. The model degrades when used at high volume due to server load.
- B. Users communicate less clearly in production than in controlled testing.
- C. The model's capabilities diminish as the conversation grows longer.
- D. Production inputs are more diverse and edge-case-rich than test inputs.**

Correct Answer: D**Explanation**

Testing typically uses representative, well-formed inputs that the developer anticipated. Production exposes the system to the full diversity of real user behaviour — unusual phrasing, edge cases, domain-specific terminology, incomplete information — which tests never fully cover. Distribution shift between test and production is the most common cause of deployment failure.

Q47. What is the most dangerous failure mode in AI-assisted decision support?

- A. The AI produces too many options, creating decision paralysis for the user.
- B. The AI confidently produces wrong recommendations that the decision-maker accepts without review.**
- C. The AI refuses to make a recommendation on high-stakes decisions.
- D. The AI makes different recommendations on the same situation across sessions.

Correct Answer: B**Explanation**

Confident wrong recommendations that are accepted without review is the worst failure mode — it combines AI's tendency to hallucinate confidence with a human tendency to defer to confident-sounding authority. The result is wrong decisions made on authoritative-seeming AI output without the verification that would have caught the error.

Q48. What is "specification gaming" as an AI failure mode?

- A. The AI interprets an ambiguous specification too literally and produces wrong output.
- B. The AI fails to follow a specification it was given in the system prompt.
- C. The AI achieves the specified metric while violating the underlying intent.**
- D. The AI produces outputs that technically meet the specification but are of poor quality.

Correct Answer: C

Explanation

Specification gaming (reward hacking) occurs when the AI finds a way to satisfy the measured objective without achieving the actual goal. A summarisation AI told to minimise length might produce a one-sentence summary that technically satisfies the constraint while losing all meaningful content. The measure is optimised; the intent is subverted.

Q49. What is "automation bias" and how does it affect AI-assisted professional work?

- A. The tendency to trust and accept AI outputs without applying sufficient critical review.**
- B. The tendency to prefer automated processes over manual ones regardless of quality.
- C. The AI model's tendency to produce outputs consistent with its training biases.
- D. A systematic error introduced when AI automates tasks with inherent bias.

Correct Answer: A

Explanation

Automation bias is a human cognitive phenomenon where people over-rely on and under-challenge automated systems. In AI-assisted professional work, automation bias causes people to accept AI outputs with less scrutiny than they would apply to human-generated work, even when verification is essential.

Q50. What is the most appropriate response when an AI makes a factual error in a professional context?

- A. Discard all AI use on this task type to prevent future errors.
- B. Ask the AI to regenerate the response with a request to be more accurate.
- C. Report the error to the AI provider as a model defect to be fixed.
- D. Correct the specific error and assess whether it indicates a pattern requiring systematic changes.**

Correct Answer: D

Explanation

The immediate response is to correct the specific error. The strategic response is to assess whether the error reveals a systematic pattern — a task type where AI consistently fails — that should change prompting strategy, verification protocols, or whether to use AI at all for that task type.

Q51. What is "brittleness" in AI systems and when does it most commonly appear?

- A. When the AI system is fragile and frequently produces error messages or crashes.
- B. When the model performs well in testing but deteriorates quickly after deployment.
- C. When small changes in input produce unexpectedly large changes in output quality.**
- D. When the model's outputs are correct but structured incorrectly for downstream use.

Correct Answer: C

Explanation

AI system brittleness means the system is sensitive to small input variations — slight rephrasing, different ordering, unusual punctuation — that cause output quality to drop dramatically. It commonly appears at the edges of the training distribution and at boundary conditions the system was not tested for.

Q52. What is "task decomposition failure" in an AI workflow?

- A. The AI misidentifies how to break a complex task into manageable subtasks.**
- B. The workflow breaks down because the context window cannot hold all subtasks.
- C. The AI fails to complete tasks because it is given too many subtasks at once.
- D. A failure where one agent's task decomposition conflicts with another agent's.

Correct Answer: A

Explanation

Task decomposition failure occurs when the AI creates a flawed plan — wrong subtasks, wrong order, missing steps, or incorrect dependencies — that cannot produce the correct final output even if each subtask is executed perfectly. It is a planning error, not an execution error.

Q53. What is the most important diagnostic question when an AI workflow produces wrong outputs?

- A. "Did the AI model produce lower-quality outputs than in our initial testing?"
- B. "Which specific step in the workflow produced the error that propagated to the output?"**
- C. "Did the user input contain errors that the AI then replicated in the output?"
- D. "Was the context window too small for the volume of information processed?"

Correct Answer: B

Explanation

Diagnosing AI workflow failures requires tracing to the specific step where the error originated. Without step-level diagnosis, fixes are guesses. The error might be in extraction (wrong data), reasoning (wrong inference), or generation (wrong format) — each requiring a different fix.

Q54. What is "overconfident failure" as a unique risk in AI systems?

- A. The AI overestimates how complex a task is and over-engineers its response.
- B. The AI refuses tasks within its capability because it underestimates its own ability.
- C. AI expresses high confidence on outputs that are actually wrong or uncertain.**
- D. A failure where confident AI use leads organisations to reduce human expertise.

Correct Answer: C

Explanation

AI systems produce confident-sounding text even when they are effectively guessing. This creates a specific risk category where wrong outputs are harder to detect than human errors, which often involve hedging language or obvious uncertainty. Confident wrong answers pass initial review more easily than unsure wrong answers.

Q55. What is the most reliable way to identify systematic AI failure patterns in production?

- A. Surveying users about whether they were satisfied with AI outputs.
- B. Comparing production output error rates to model benchmark accuracy scores.
- C. Running periodic A/B tests with human-generated outputs as the control condition.
- D. Tracking error types and correlating them with input characteristics and task types.**

Correct Answer: D

Explanation

Systematic pattern tracking — logging errors with metadata about input type, topic, format, and user context — allows identification of conditions under which AI consistently fails. User satisfaction surveys are too aggregated; benchmark comparisons are too distant from production; A/B tests catch some but not systematic errors across all input types.

Q56. What is the most important human capability that must not be lost even with advanced AI assistance?

- A. The ability to critically evaluate outputs and make final judgments independently.**
- B. The ability to perform tasks manually without AI assistance when systems fail.
- C. The ability to write effective prompts to extract value from AI systems.
- D. The ability to explain AI behaviour to non-technical stakeholders.

Correct Answer: A

Explanation

Critical evaluation and independent judgment are the irreplaceable human capabilities in AI-assisted work. If these atrophy — if people stop being able to assess whether an AI output is right or wrong — then AI errors propagate unchecked. All other skills (manual fallback, prompting, communication) are secondary to this core judgment capability.

Q57. What is the most common reason AI underperforms on specialised professional tasks?

- A. The model applies excessive safety filtering on professional domain content.
- B. Training data underrepresents the specific domain, reducing pattern quality for those tasks.**
- C. Professional tasks require real-time data that exceeds the model's knowledge cutoff.
- D. The model cannot distinguish between professional and general user queries.

Correct Answer: B

Explanation

Specialised professional domains (rare medical subspecialties, niche regulatory frameworks, proprietary methodologies) have less representation in training data. The model's pattern quality for these tasks is weaker than for general topics. This manifests as outputs that use the right vocabulary but miss domain-specific nuances or make domain-specific errors.

Q58. What is "model collapse" risk in a multi-model pipeline?

- A. Multiple model calls on the same data produce contradictory outputs that cannot be resolved.
- B. The pipeline fails when any single model call errors out, halting the entire workflow.
- C. Individual models produce lower quality outputs when called from a pipeline than directly.
- D. Errors propagate and compound through the pipeline, producing increasingly wrong outputs.**

Correct Answer: D

Explanation

In multi-model pipelines, early model errors become the inputs for subsequent models. Each model processes the flawed input as if it were correct, often compounding the error. The final output can be very wrong while each individual model call seemed plausible given its input. This is the pipeline equivalent of cascading failure.

Q59. What is the most important pre-deployment test for an AI system that will make consequential recommendations?

- A. Testing on adversarial inputs designed to find recommendations that are dangerously wrong.**
- B. Testing the system's recommendation quality against human expert baseline scores.
- C. Testing the system's latency under the expected production request volume.
- D. Testing whether the system correctly declines requests outside its competence scope.

Correct Answer: A

Explanation

Adversarial testing for dangerously wrong recommendations is most critical before deploying decision-support AI. The system will inevitably encounter edge cases in production; pre-deployment adversarial testing discovers the failure modes before real users do, when the cost of errors is high.

Q60. What is the most important thing an organisation should do after an AI system causes a significant error?

- A. Disable AI use for that task type until a better model is available.
- B. Disclose the error to all affected stakeholders and offer remediation.
- C. Conduct a root cause analysis and implement systematic prevention measures.**
- D. Report the failure to the AI model provider for model improvement.

Correct Answer: C

Explanation

Root cause analysis and systematic prevention are the most important responses. Understanding exactly what failed and why allows the organisation to fix the specific problem — whether it is a prompting issue, a missing verification step, or a fundamental task mismatch — and prevent recurrence across all similar deployments.

SET

5

Professional Accountability

Set 5 of 5 · 15 Questions · AI Judgment

Q61. Who is accountable for the consequences of an AI-generated output used in a professional decision?

- A. The professional who used the AI output and chose to act on it.
- B. The AI provider who created the model that generated the output.
- C. The organisation that deployed the AI tool used in the workflow.
- D. Accountability is shared between the AI provider and the user in proportion to their role.

Correct Answer: A**Explanation**

Professional accountability for the consequences of AI outputs rests with the professional who used the output and chose to act on it. AI providers disclaim liability for misuse; organisational deployment is a contributing factor. But the professional who made the decision — relying on unverified AI output — bears the primary accountability.

Q62. What is the most important disclosure consideration when submitting AI-assisted professional work?

- A. Always disclosing AI use in all professional work, regardless of context.
- B. Disclosing only when AI generated more than 50% of the final output.
- C. Not disclosing AI use since it is now standard practice and assumed.
- D. Whether the context requires disclosure and what level of disclosure is appropriate.

Correct Answer: D**Explanation**

Disclosure requirements vary by context: academic submission (strict policies), client deliverable (relationship and contract dependent), regulatory filing (jurisdiction and sector dependent). The professional must understand the applicable requirements for their specific context — not apply a blanket always-disclose or never-disclose policy.

Q63. A professional submits AI-generated work to a client as their own analysis. The client makes a wrong decision based on AI hallucinations in the work. What is the professional's primary responsibility?

- A. To notify the AI provider of the hallucination so the model can be improved.
- B. To own the consequences — they verified the AI output insufficiently before submission.**
- C. To share responsibility with the AI tool since it generated the wrong content.
- D. To disclose to the client that AI was used so they can make their own assessment.

Correct Answer: B

Explanation

The professional took accountability for the work by submitting it as their analysis. Insufficient verification of AI outputs before acting on them is a professional failure, not a technology failure. Disclosure after harm has been done does not discharge the accountability for insufficient verification before submission.

Q64. What is the most important principle for maintaining professional integrity when using AI?

- A. Disclose AI use consistently to all clients and stakeholders.
- B. Only use AI for tasks where errors are immediately obvious and inconsequential.
- C. Own every output fully — verify, understand, and be able to defend every claim.**
- D. Use AI only for tasks explicitly approved by your professional body or employer.

Correct Answer: C

Explanation

Professional integrity with AI means full ownership of every output — you have verified the key claims, you understand the reasoning, and you can defend every assertion as if you had produced it yourself. Using AI as a production tool while maintaining this standard of ownership is the correct professional posture.

Q65. What is "deskilling risk" in the context of professional AI use?

- A. Relying on AI so heavily that core professional skills atrophy from disuse.**
- B. The risk of hiring less-skilled professionals because AI reduces skill requirements.
- C. The risk of AI errors causing professionals to make decisions below their normal skill level.
- D. The demotion of skilled professionals who resist AI adoption in their organisations.

Correct Answer: A

Explanation

Deskilling risk refers to the cognitive atrophy that occurs when professionals delegate tasks so completely to AI that they lose the underlying skill. A lawyer who lets AI draft everything may lose the ability to draft independently; an analyst who lets AI interpret data may lose analytical judgment. This creates vulnerability when AI fails.

Q66. What is the correct response when a client asks you to use AI to do something you are not confident it can do reliably?

- A. Use AI anyway and review the output carefully to catch errors.
- B. Decline the task entirely to avoid professional risk.
- C. Be transparent about AI's limitations for this task and propose an approach with appropriate human judgment.**
- D. Charge more to compensate for the additional verification time required.

Correct Answer: C

Explanation

Transparency about AI's limitations for the specific task — and proposing an approach that appropriately supplements AI with human judgment where needed — is the professional response. It sets accurate expectations, proposes a viable solution, and maintains the professional relationship without either blind AI use or unnecessary refusal.

Q67. What is the most important thing to understand about AI and professional liability?

- A. AI providers carry partial liability for model errors that cause client harm.
- B. Professional liability is reduced when AI is used, since errors are the tool's fault.
- C. Professional liability applies only when AI is used without human review.
- D. AI use does not transfer professional liability — the professional remains fully accountable.**

Correct Answer: D

Explanation

Professional liability — the legal and ethical accountability for one's professional work — does not transfer to an AI tool. A lawyer, doctor, engineer, or accountant who uses AI remains fully liable for the outputs they provide to clients, regardless of whether AI generated them. This is the foundational liability principle for professional AI use.

Q68. What is "professional judgment" in the context of AI-assisted work?

- A. A licensed professional's ability to produce work without any AI assistance.
- B. The human expert's capacity to direct AI appropriately and evaluate its outputs critically.**
- C. The certification that a professional has received AI literacy training.
- D. The right of a professional to choose whether or not to use AI tools.

Correct Answer: B

Explanation

Professional judgment in AI-assisted work is the capacity to: know which tasks benefit from AI, direct AI with appropriate specifications, critically evaluate AI outputs for quality and accuracy, and take accountability for the final work product. It is the meta-skill that makes all other AI capabilities professionally safe and valuable.

Q69. What is the most important way AI changes professional development requirements?

- A. Professionals can now develop expertise more quickly because AI speeds learning.
- B. Domain expertise becomes less important as AI can supply domain knowledge on demand.
- C. Professionals must develop critical evaluation skills alongside domain expertise.**
- D. Professional development should focus on AI skills rather than core domain skills.

Correct Answer: C

Explanation

AI as a production tool means professionals must now develop AI judgment skills — how to direct AI, how to evaluate outputs, when to verify, when to override — alongside (not instead of) their core domain expertise. Domain expertise is more important than ever as the basis for AI output evaluation.

Q70. What is the correct professional response when AI produces an output that contradicts your domain expertise?

- A. Investigate carefully — the AI may have found something you missed, or it may be wrong.**
- B. Trust your expertise — it is always more reliable than AI pattern-matching.
- C. Trust the AI — it has processed more information on the topic than you have.
- D. Present both the AI output and your expert view and let the client decide.

Correct Answer: A

Explanation

Contradiction between AI output and domain expertise is a signal worth investigating — not automatically resolving in favour of either. The AI may have identified a genuine gap; you may have knowledge the AI lacks; the AI may be hallucinating. Investigation determines which, rather than reflexive deference to either.

Q71. What is the most important governance requirement for AI use in a regulated profession?

- A. Regulatory approval for each AI tool used in professional practice.
- B. Documented processes showing how AI outputs are verified before professional use.**
- C. Client consent before any AI tool is used in delivering professional services.
- D. A dedicated AI compliance officer who reviews all AI-generated professional work.

Correct Answer: B

Explanation

Regulatory governance for professional AI use requires documented verification processes — showing that professional judgment was applied, that AI outputs were checked, and that the professional owns the final work product. Approval, consent, and compliance officers are supporting elements; documented verification processes are the evidentiary foundation.

Q72. What is "AI washing" in a professional context and why is it problematic?

- A. Adding unnecessary AI features to products to appear more technologically advanced.
- B. Over-claiming AI capabilities to win business that the AI cannot actually deliver.
- C. Using AI branding on content to attract attention and signal modernity.
- D. Claiming AI-generated work as purely human-generated to meet expectations or avoid disclosure.**

Correct Answer: D

Explanation

AI washing in professional contexts means presenting AI-generated work as original human analysis — to meet client expectations of human professional work, avoid disclosure obligations, or exploit the reputational value of human expertise. This is a form of professional deception with ethical and potentially legal consequences.

Q73. A junior team member asks whether they should use AI for everything to maximise productivity. The most accurate guidance is:

- A. Use AI for everything — productivity gains always outweigh any quality concerns.
- B. Avoid AI until you have mastered the underlying skills independently.
- C. Use AI where it genuinely helps and exercise judgment to develop your own skills alongside AI.**
- D. Use AI only for tasks that have been approved by the team lead for AI use.

Correct Answer: C

Explanation

The most balanced and accurate guidance is to use AI where it genuinely accelerates work, but to actively exercise and develop your own capabilities alongside AI use. Complete AI delegation without human engagement atrophies the skills that make you a professional capable of evaluating AI output — and employable when AI changes.

Q74. What is the most important quality a professional needs to use AI responsibly in the long term?

- A. Technical understanding of how AI models work at the architecture level.
- B. Intellectual honesty about what AI can and cannot do, and when its outputs should not be trusted.**
- C. Comprehensive awareness of all available AI tools and their respective capabilities.
- D. A consistent record of disclosing AI use to all professional stakeholders.

Correct Answer: B

Explanation

Intellectual honesty — the honest assessment of AI capability and limitation, and the courage to say "I don't trust this AI output" when appropriate — is the foundational quality for responsible long-term AI use. Technical knowledge, tool awareness, and disclosure are all valuable but secondary to this core intellectual honesty.

Q75. Menler's AI Judgment bank tests practitioners at what level of professional maturity?

- A. Those learning about AI for the first time and exploring its basic capabilities.
- B. Those building AI systems and models who need to evaluate their own outputs.
- C. Those managing AI teams and evaluating which AI tools to invest in.
- D. Those who must make independent, accountable decisions about AI use in their professional work.**

Correct Answer: D

Explanation

The AI Judgment bank tests the maturity needed to use AI responsibly as a professional — making independent decisions about when to trust AI, when to verify, when to override, and taking full professional accountability for AI-assisted work. It is for practitioners making real professional decisions, not beginners, builders, or managers.

END OF QUESTION BANK

*Your turning
point starts
here.*