

menler

Your turning point in the AI era.

AI APTITUDE QUESTION BANK

AI-Native Execution

About This Question Bank

These questions cover the practise of building and operating with AI as the primary execution layer: vibe coding, rapid prototyping, MVP construction, AI-assisted development, workflow acceleration, and AI-first operations.

5

Sets

15

Questions per Set

75

Total Questions

SET GUIDE

Set 1 · Vibe Coding and Rapid Prototyping

Vibe coding and rapid prototyping — natural language to code, iteration discipline

Set 2 · MVP Building with AI

MVP building with AI — go-to-market assets, demand validation, scaling patterns

Set 3 · AI-Assisted Development

AI-assisted development — code review, debugging, test generation, documentation

Set 4 · Workflow Acceleration

Workflow acceleration — delegation patterns, human oversight, compounding practises

Set 5 · AI-First Operations

AI-first operations — org design, escalation paths, quality monitoring, governance

Each question has one correct answer. Select the best answer from the options provided.

Correct answers and explanations follow each question.

SET

1

Vibe Coding and Rapid Prototyping

Set 1 of 5 · 15 Questions · AI-Native Execution

Q1. A founder wants to build a working web app prototype in 48 hours without a development team. The best AI-native approach is:

- A. Hire a contractor to write the spec first
- B. Use a natural language-to-code tool to generate a working prototype from a plain English spec, then iterate**
- C. Wait until the team can properly architect the backend
- D. Build in Figma only and call it a prototype

Correct Answer: B

Explanation

AI code generation tools allow non-engineers and engineers alike to go from idea to working prototype in hours. Speed of learning beats architectural perfection at this stage.

Q2. What is vibe coding as a practise in AI-native product development?

- A. Writing code while listening to music to improve focus
- B. Using AI-generated music as a backend for audio apps
- C. A style of pair programming where both engineers code simultaneously
- D. Using conversational AI to iterate on code through natural language instructions, accepting imperfect output and steering toward the goal**

Correct Answer: D

Explanation

Vibe coding treats AI as a co-pilot where the human steers with intent and accepts rough output that gets refined. It optimises for speed of exploration over code quality.

Q3. A non-engineer founder is building a Chrome extension using Claude. They get a working build on the third iteration. What is the most important practise to maintain momentum?

- A. Test each iteration against the specific use case before requesting the next change**
- B. Accept all AI output without testing to move faster
- C. Switch to a different AI tool for each iteration
- D. Hire an engineer to review every AI output

Correct Answer: A

Explanation

In AI-assisted prototyping, tight test-iterate loops catch errors early and give precise feedback to the AI for the next prompt. Untested iterations compound errors.

Q4. When building a rapid AI prototype, which type of input gives an AI coding tool the most useful starting point?

- A. A list of all desired features
- B. A database schema
- C. A concrete user story describing what a specific user does, not what the system does**
- D. An architecture diagram

Correct Answer: C

Explanation

User stories ground the AI in behaviour rather than structure. They produce code oriented toward outcomes, not implementation choices the AI may get wrong.

Q5. A team uses AI to generate a prototype in 2 days. A stakeholder asks them to clean it up before showing investors. The AI-native approach to this request is:

- A. Rewrite the entire codebase from scratch
- B. Ask the AI to refactor for readability and add comments, then review the key interaction paths manually**
- C. Decline the request because prototypes should not be cleaned up
- D. Hand off to a senior engineer to rewrite

Correct Answer: B

Explanation

AI can clean up its own output quickly. The human role is to verify that refactoring did not break functionality - not to rewrite everything manually.

Q6. What is the primary risk of using AI-generated code in production without review?

- A. Security vulnerabilities and logic errors that the AI introduced but did not flag**
- B. The code will be too slow
- C. The code will not run on all browsers
- D. AI licensing restrictions prevent production use

Correct Answer: A

Explanation

AI code generators can introduce subtle bugs, insecure patterns, and incorrect logic without warning. Production deployment requires human review of security-critical paths.

Q7. A solo founder uses AI to build a working SaaS MVP. The first feature they should prioritise shipping is:

- A. The admin dashboard
- B. The settings page
- C. The onboarding email sequence
- D. The one feature that allows a real user to complete the core job-to-be-done end to end**

Correct Answer: D

Explanation

The minimum viable product must do the core job. Peripheral features add no learning until the core workflow is validated with real users.

Q8. When should a rapid AI prototype be discarded rather than refactored into a production codebase?

- A. After every prototype, always
- B. Only if the prototype did not work
- C. When the prototype was built with speed-first decisions that create architectural debt incompatible with production requirements**
- D. When it takes longer than one week to build

Correct Answer: C

Explanation

Prototypes built for speed carry shortcuts that do not scale. The decision to discard vs refactor depends on how deep the architectural debt runs.

Q9. A developer uses an AI tool to generate a database schema from a natural language description. The most important post-generation step is:

- A. Validating that the schema reflects actual data relationships and will support the required queries**
- B. Accepting the schema immediately to save time
- C. Asking the AI to generate the schema again with different words
- D. Converting the schema to a different database engine

Correct Answer: A

Explanation

AI-generated schemas often miss edge cases or model relationships incorrectly. Structural validation before building on top of a schema prevents costly rewrites.

Q10. What does context window as a scratchpad mean in AI-native development practise?

- A. Storing code snippets in a text file
- B. Using the terminal as the primary development environment
- C. Keeping AI outputs in a separate document for later review
- D. Using the conversation history with the AI to maintain project state, decisions, and constraints across a session**

Correct Answer: D

Explanation

The AI context window serves as working memory for the session. Skilled AI-native developers front-load context such as specs, constraints, and prior decisions to get consistent outputs.

Q11. A team prototyping with AI notices the AI frequently forgets earlier constraints as the conversation grows. The best mitigation is:

- A. Starting a new conversation for every change
- B. Re-stating key constraints at the start of each new prompt and using a system prompt or pinned context**
- C. Reducing the length of prompts
- D. Switching to a different AI tool

Correct Answer: B

Explanation

Context window limitations cause AI to drift from earlier instructions. Re-stating constraints and using persistent system prompts maintains coherence across a session.

Q12. Which describes the correct sequencing in AI-native rapid prototyping?

- A. Idea to architecture document to engineering sprint to AI review to launch
- B. Idea to design in Figma to engineer the design to add AI later
- C. Idea to spec in natural language to AI generates working code to test with real users to iterate**
- D. Idea to hire a team to write a PRD to AI assists with documentation

Correct Answer: C

Explanation

AI-native prototyping collapses the traditional build cycle. Natural language specs feed directly into working code, compressing the idea-to-user-feedback loop.

Q13. A developer wants to add authentication to an AI-generated app. The safest approach is:

- A. Ask the AI to generate a custom authentication system
- B. Skip authentication for the MVP
- C. Copy auth code from a GitHub repository without review
- D. Use a battle-tested auth library or service such as Auth0 or Clerk rather than asking the AI to generate custom auth code**

Correct Answer: D

Explanation

Authentication is a security-critical component where custom AI-generated code frequently introduces vulnerabilities. Established services handle this more safely.

Q14. What is the value of maintaining a decision log when prototyping with AI?

- A. It satisfies compliance requirements
- B. It captures why choices were made, enabling the AI to stay consistent and the team to onboard faster**
- C. It improves AI response speed
- D. It is required for deploying to production

Correct Answer: B

Explanation

A decision log is a human-readable record of architectural and design choices. It feeds back into AI context and preserves rationale as the project grows.

Q15. A non-technical founder builds an MVP with AI assistance and gets their first paying customer. The next highest-leverage AI use is:

- A. Using AI to automate the manual parts of onboarding and fulfilment before hiring anyone**
- B. Using AI to redesign the landing page
- C. Using AI to write investor pitch decks
- D. Using AI to generate more code features immediately

Correct Answer: A

Explanation

After first revenue, AI-native operators use AI to scale operations before headcount. Automating manual steps extends runway and proves the business model further.

SET

2

MVP Building with AI

Set 2 of 5 · 15 Questions · AI-Native Execution

Q16. A startup uses AI to generate their first landing page, waitlist form, and email sequence. What does this demonstrate about AI-native execution?

- A. AI can fully replace a marketing team
- B. AI is only useful after product-market fit
- C. AI compresses the cost and time of go-to-market assets, enabling solo founders to test demand before building**
- D. AI-generated landing pages convert better than human-written ones

Correct Answer: C

Explanation

AI-native founders use AI to validate demand with minimal investment. Landing page, waitlist, and email sequence can be shipped in hours, enabling pre-build demand testing.

Q17. When building an MVP with AI assistance, which deliverable should be AI-generated last?

- A. The core data model, which should be designed by a human to avoid structural errors that compound**
- B. The landing page copy
- C. The email templates
- D. The FAQ section

Correct Answer: A

Explanation

Data models are foundational. AI-generated schemas may look correct but miss relational integrity or scalability concerns that become expensive to fix later.

Q18. A team of two uses AI tools to ship a product in 3 weeks that would have taken a 10-person team 6 months traditionally. The primary leverage factor is:

- A. The team worked 24 hours a day
- B. The product had fewer features than normal
- C. The team used a no-code tool instead of AI
- D. AI handling execution tasks that previously required specialist headcount across design, code, copy, and QA**

Correct Answer: D

Explanation

AI-native teams are force multiplied because AI performs execution work across multiple functions simultaneously. Headcount is no longer the bottleneck for shipping.

Q19. A PM is building an MVP and must decide what to do manually vs what to delegate to AI. The correct filter is:

- A. Delegate everything that takes more than one hour to AI
- B. Delegate execution tasks to AI; retain human judgment for strategy, validation, and decisions with irreversible consequences**
- C. Keep all code tasks human and delegate only writing tasks to AI
- D. Use AI for external tasks and humans for internal tasks

Correct Answer: B

Explanation

AI excels at execution but lacks contextual judgment. Humans should own strategic choices, quality validation, and decisions that cannot be undone easily.

Q20. An AI-assisted MVP is live and getting traction. The founder wants to scale. The AI-native scaling principle is:

- A. Hire first and then use AI to make hires more productive
- B. Stop using AI and build a proper engineering team
- C. Identify the highest-friction manual operations and automate them with AI before hiring for those roles**
- D. Replace all AI tooling with custom-built internal systems

Correct Answer: C

Explanation

AI-native scaling defers headcount until automation reaches its limit. The pattern is: automate, then saturate, then hire, then automate again.

Q21. What is the primary role of AI in the build, measure, learn lean startup loop for an AI-native team?

- A. Replacing the measure phase with AI analytics
- B. Automating the learn phase by having AI interpret all data
- C. Making the loop unnecessary by getting it right the first time
- D. Compressing the build phase from weeks to days, enabling faster iteration through the full loop**

Correct Answer: D

Explanation

AI compresses build time, which accelerates the entire loop. More loops per unit time means faster learning - the core advantage of AI-native execution.

Q22. A founder uses AI to build an MVP in 2 weeks. A VC asks whether this is defensible. The most honest AI-native answer is:

- A. The build speed is not the moat - the data, distribution, and user insight accumulated through rapid iteration are**
- B. Yes, because AI tools are expensive and competitors cannot afford them
- C. Yes, because the code is proprietary
- D. No, and the product should not be shown to investors yet

Correct Answer: A

Explanation

AI-native build speed lowers barriers for everyone. The defensible advantage is the learning and asset accumulation such as data, users, and brand that speed enables.

Q23. A non-technical founder is building a B2B SaaS MVP with AI. The first technical decision they must make correctly is:

- A. Which programming language to use
- B. Which hosting and database platform will allow them to iterate without refactoring the stack later**
- C. Which AI model to integrate first
- D. How many engineers to hire

Correct Answer: B

Explanation

Platform choices are hard to reverse. Selecting a flexible, developer-friendly stack prevents early architectural lock-in even if AI generates all the code.

Q24. A team builds an MVP in 3 weeks and launches to 100 beta users. After 2 weeks, 80% have churned. The AI-native response is:

- A. Shut down and rebuild from scratch
- B. Wait for more data before making changes
- C. Add more features to retain remaining users
- D. Use AI to rapidly analyse qualitative feedback, identify the core failure, and iterate within days**

Correct Answer: D

Explanation

AI-native teams respond to failure with speed. AI can synthesise feedback, identify patterns, and support rapid re-prioritisation - compressing the response cycle.

Q25. A PM is building a B2C app MVP. Which AI-generated asset directly reduces time-to-first-user the most?

- A. A detailed technical specification document
- B. A competitive analysis report
- C. A working landing page with a clear value proposition and email capture**
- D. A 12-month roadmap

Correct Answer: C

Explanation

First users come from demand - not from documentation or planning. A working landing page enables immediate demand testing and user acquisition before the product is built.

Q26. A founder instructs Claude to build a payment integration. What must they verify before going live?

- A. That the code was written in under 10 minutes
- B. That the integration correctly handles edge cases like failed payments, refunds, and webhook retries**
- C. That the AI chose the cheapest payment processor
- D. That the code uses the latest programming language version

Correct Answer: B

Explanation

Payment integrations fail silently on edge cases. AI-generated code handles the happy path well but often misses error handling - human verification here is critical.

Q27. Which of the following best describes an AI-native team relationship with documentation?

- A. Documentation is AI-generated alongside the code and kept in sync via the same AI workflow**
- B. Documentation is written manually after the product ships
- C. Documentation is not needed because the AI knows what the code does
- D. Documentation is only written for enterprise customers

Correct Answer: A

Explanation

AI-native teams generate docs in the same workflow as code. This removes the documentation lag that plagues traditional teams and keeps context available for future AI sessions.

Q28. A startup wants to test three different product concepts. The AI-native approach is:

- A. Spend six months researching which concept is best before building
- B. Build the most technically complex concept first
- C. Build all three as lightweight prototypes using AI in parallel, then test each with users before committing to one**
- D. Test only one concept at a time to keep focus

Correct Answer: C

Explanation

AI compresses prototype cost enough to test multiple concepts simultaneously. Parallel validation beats sequential testing when the cost of building is near-zero.

Q29. A PM is using AI to generate user interview scripts for MVP validation. The most important human contribution to this process is:

- A. Editing the grammar of the AI-generated script
- B. Translating the script into different languages
- C. Recording the AI-generated script as an audio file
- D. Selecting which questions to prioritise based on what assumptions are most critical to test**

Correct Answer: D

Explanation

AI can generate comprehensive interview scripts. Human judgment determines which assumptions most need testing - that prioritisation is the highest-value contribution.

Q30. Which metric most directly validates an AI-native MVP after the first 30 days?

- A. Total number of AI prompts used to build the product
- B. Whether users return to complete a second session without being prompted**
- C. Number of features shipped
- D. Landing page conversion rate

Correct Answer: B

Explanation

Unprompted return usage is the earliest signal of genuine product-market fit. It shows the product created enough value to motivate behaviour change.

SET

3

AI-Assisted Development

Set 3 of 5 · 15 Questions · AI-Native Execution

Q31. A developer uses an AI coding assistant to write a function. The most effective review practise is:

- A. Testing the function against the full range of expected inputs including edge cases before merging**
- B. Reading the AI explanation of what the code does and accepting it
- C. Running the code once to see if it produces output
- D. Asking the AI if the code is correct

Correct Answer: A

Explanation

AI coding assistants are confident but not always correct. Edge case testing is the human job - AI explanations are not a substitute for execution-level verification.

Q32. When using an AI to debug code, the most effective prompt structure is:

- A. Describing the bug in general terms and asking the AI to find it
- B. Asking the AI to rewrite the entire file
- C. Pasting the exact error message, the relevant code block, and the expected vs actual behavior**
- D. Running all tests and sending the output without context

Correct Answer: C

Explanation

Specific, structured bug reports give the AI the signals it needs to identify root causes. Vague descriptions produce generic suggestions that rarely resolve the issue.

Q33. An engineer uses AI to generate test cases. What is the primary limitation of AI-generated tests?

- A. AI-generated tests are always too slow
- B. AI tends to generate tests that validate the code it wrote rather than testing genuinely adversarial inputs**
- C. AI cannot write tests in Python
- D. AI-generated tests require separate licensing

Correct Answer: B

Explanation

AI test generation is biased toward confirming its own code. Human testers or adversarial prompting is needed to surface the edge cases AI will not naturally include.

Q34. A team integrates an AI coding assistant into their CI/CD pipeline to auto-generate code review comments. The key governance requirement is:

- A. The AI must review all pull requests without human override
- B. The AI must use the same style guide as the team
- C. The AI must be trained on the team codebase first
- D. Human engineers must retain final approval authority before any suggested change is merged**

Correct Answer: D

Explanation

AI in code review is advisory. Merging AI-suggested changes without human sign-off removes accountability and risks introducing bugs at scale.

Q35. A developer is working with an AI assistant on a large codebase. The most important context to provide at the start of each session is:

- A. The relevant module architecture, the specific task being worked on, and any constraints or conventions in use**
- B. The full codebase pasted into the prompt
- C. A list of all features the product has
- D. The team org chart

Correct Answer: A

Explanation

AI coding assistants produce better outputs with focused, relevant context. Full codebase dumping wastes context window; targeted context scopes the task correctly.

Q36. A team uses AI to migrate a legacy codebase from one framework to another. What is the highest-risk phase of this process?

- A. Converting syntax from one language to another
- B. Validating that migrated code preserves all business logic and edge case behaviour from the original**
- C. Renaming variables and functions
- D. Updating import statements

Correct Answer: B

Explanation

AI migrations handle syntax conversion well but can silently alter business logic. The human-critical phase is verifying that all behaviours are preserved, not just that the code runs.

Q37. An AI generates a 200-line function to implement a feature. A senior engineer says it should be broken into 5 smaller functions. The correct response is:

- A. Accept the 200-line function because the AI wrote it
- B. Reject AI assistance entirely because the output was wrong
- C. Manually refactor without using AI
- D. Ask the AI to refactor the function with a clear instruction about desired modularity and maximum function length**

Correct Answer: D

Explanation

AI refactors on instruction. The human role is to specify the quality standard; the AI executes it. This is faster than manual refactoring and preserves AI leverage.

Q38. A developer uses AI to generate a SQL query. Before running it on production data, the most critical check is:

- A. Checking the query length
- B. Confirming the AI generated the query in under 5 seconds
- C. Verifying the query against a schema diagram to confirm it targets the correct tables and relationships**
- D. Ensuring the query uses uppercase keywords

Correct Answer: C

Explanation

AI-generated SQL frequently targets plausible but incorrect tables or columns. A schema check before execution prevents silent data errors or accidental data loss.

Q39. A PM asks an AI to generate API documentation from the codebase. The most important post-generation step is:

- A. Publishing the documentation immediately
- B. Having an engineer verify that generated descriptions match actual API behaviour, especially for error responses**
- C. Translating the documentation into other languages
- D. Asking the AI to make the documentation shorter

Correct Answer: B

Explanation

AI-generated API docs can describe what an API should do, not what it actually does. Error response documentation is especially prone to inaccuracy and needs verification.

Q40. Which practise most increases the long-term productivity of a team using AI-assisted development?

- A. Maintaining a shared prompt library for common tasks so the team benefits from each other prompt iterations**
- B. Each engineer developing their own private AI workflow
- C. Using only one AI tool to avoid confusion
- D. Limiting AI assistance to junior engineers

Correct Answer: A

Explanation

Prompt quality compounds. A shared library lets the whole team benefit from individual prompt improvements and creates institutional knowledge around AI-assisted workflows.

Q41. A developer is using an AI assistant to implement a feature described in a ticket. The ticket is ambiguous. The correct first step is:

- A. Asking the AI to interpret the ambiguous ticket and implement its interpretation
- B. Implementing both interpretations and asking the PM to choose
- C. Clarifying the ambiguity with the PM before asking the AI to implement, to avoid building the wrong thing correctly**
- D. Closing the ticket and asking for a rewrite

Correct Answer: C

Explanation

AI will execute ambiguous instructions confidently. Clarifying human intent before AI execution prevents well-executed implementations of the wrong requirement.

Q42. An AI coding assistant suggests using a deprecated library. The correct response is:

- A. Using the deprecated library because the AI chose it
- B. Asking the AI if the library is deprecated
- C. Switching to a different AI coding tool
- D. Rejecting the suggestion and specifying the current library version before regenerating**

Correct Answer: D

Explanation

AI training data has a knowledge cutoff. Library recommendations can reflect outdated choices. Specifying current versions in the prompt prevents deprecated or insecure dependencies.

Q43. A company deploys an AI coding assistant to its engineering team. The most important policy to establish is:

- A. AI-generated code may be merged directly without review to maximise speed
- B. All AI-generated code must be reviewed and understood by the engineer before merging**
- C. AI must only be used for test files, not production code
- D. Engineers must disclose to the AI that they are professionals

Correct Answer: B

Explanation

Code ownership requires comprehension. Engineers who merge code they do not understand cannot debug, extend, or be accountable for it - regardless of who wrote it.

Q44. A team uses AI to write code and AI to write tests for that code. Which external check best validates both?

- A. Running AI-generated unit tests only
- B. Asking the AI to review both the code and tests
- C. End-to-end testing with real user flows and production-like data**
- D. Checking that the code compiles without errors

Correct Answer: C

Explanation

AI-generated code and AI-generated tests can be mutually reinforcing blind spots. End-to-end testing with real user data catches failures that unit tests miss.

Q45. Which describes the correct use of AI in a code review process?

- A. AI surfaces potential issues and style violations as a first pass; human reviewers evaluate logic, architecture, and context**
- B. AI replaces human code reviewers entirely
- C. AI only reviews documentation, not code
- D. AI approves pull requests after running tests

Correct Answer: A

Explanation

AI code review is fast and pattern-based but lacks contextual judgment. The correct model pairs AI broad pattern detection with human evaluation of intent and architecture.

SET

4

Workflow Acceleration

Set 4 of 5 · 15 Questions · AI-Native Execution

Q46. A founder uses AI to compress a task that previously took 3 hours to 20 minutes. The most important follow-up action is:

- A. Using AI for even more tasks to save more time
- B. Documenting the time saved for investor reporting
- C. Telling team members to do the same task manually to keep skills sharp
- D. Reinvesting the saved time into higher-leverage strategic work that AI cannot do**

Correct Answer: D

Explanation

AI-native execution multiplies output by freeing human time for judgment-intensive work. The compound benefit comes from where saved time is reinvested, not from the saving itself.

Q47. A content team uses AI to produce first drafts. The correct role for human editors in this workflow is:

- A. Rewriting all AI drafts entirely
- B. Fact-checking, adding unique insight, and ensuring brand voice - not rewriting AI drafts from scratch**
- C. Publishing AI drafts without review
- D. Using AI only for titles, not full drafts

Correct Answer: B

Explanation

Human editors in an AI-native content workflow are curators and elevators, not re-writers. Their value is judgment and specificity - not volume production.

Q48. A PM uses AI to prepare for a board meeting. The highest-leverage AI use in this scenario is:

- A. Using AI to take notes during the meeting
- B. Using AI to email board members after the meeting
- C. Generating the first draft of the board deck and narrative, which the PM then refines with insider context**
- D. Using AI to research board members on LinkedIn

Correct Answer: C

Explanation

AI can structure and populate board decks from raw data and objectives. The PM role is to inject the context and strategic framing AI cannot access.

Q49. A sales team uses AI to personalise outreach at scale. The most important human oversight in this workflow is:

- A. Reviewing a sample of AI-generated messages before sending to ensure accuracy and appropriate tone**
- B. Reviewing every message before sending
- C. Letting AI send all messages without review
- D. Having the AI send messages only to warm leads

Correct Answer: A

Explanation

Full review defeats the scale benefit; zero review creates risk. Sampling provides a quality check that catches systematic errors without eliminating efficiency.

Q50. A team uses AI to generate weekly status reports from raw data. After two months, leadership stops reading them. The most likely cause is:

- A. Leadership prefers verbal updates
- B. The reports are too long
- C. The AI is using the wrong template
- D. AI-generated reports lack the editorial judgment to surface what actually matters to leadership**

Correct Answer: D

Explanation

AI can aggregate and structure data but struggles to prioritise what matters to a specific audience. Reports without human editorial judgment lose relevance quickly.

Q51. A company uses AI to summarise all customer support tickets daily. The correct use of these summaries is:

- A. As a replacement for reading any individual tickets
- B. As the sole basis for product roadmap decisions
- C. As an input to the PM and support lead for trend spotting, not as a replacement for reading individual tickets on critical issues**
- D. As a public-facing customer communication

Correct Answer: C

Explanation

AI summaries are useful for trend detection at volume. Critical issues require reading the original ticket to preserve nuance that summaries compress away.

Q52. An operations team uses AI to automate weekly report generation. The process saves 8 hours per week. Which risk must be actively managed?

- A. The AI silently propagating an error in source data across multiple weeks of reports before detection**
- B. The reports being generated too quickly
- C. The AI choosing the wrong font in the report
- D. The AI generating more reports than requested

Correct Answer: A

Explanation

Automation amplifies both correct and incorrect inputs. A systematic data error can propagate undetected for weeks if no human spot-checks the source data periodically.

Q53. A PM uses AI to analyse qualitative user interview transcripts. What should they validate in the AI output?

- A. Whether the AI completed the analysis in under 2 minutes
- B. Whether the themes the AI identified match what stood out to the human interviewer during the sessions**
- C. Whether the AI used the correct summary format
- D. Whether the AI mentioned all interviewees by name

Correct Answer: B

Explanation

AI thematic analysis can miss emotional tone, implicit signals, and interviewer context. Cross-referencing AI themes against human recall surfaces what the AI flattened.

Q54. An AI is used to triage and categorise 500 inbound leads per week. What is the highest-risk failure mode to monitor?

- A. The AI taking too long to categorise each lead
- B. The AI generating too many categories
- C. High-value leads being systematically miscategorised into a low-priority bucket**
- D. The AI sending automatic responses to leads

Correct Answer: C

Explanation

Systematic misclassification of high-value leads has direct revenue impact. Monitoring should focus on false negatives in high-priority categories, not just overall accuracy.

Q55. A team uses AI to generate 30 social media posts per week. After 4 weeks, engagement drops significantly. The most productive diagnostic is:

- A. Increasing the number of posts per week
- B. Switching to a different AI tool
- C. Changing the posting time
- D. Comparing engagement on AI-generated posts vs posts that had significant human editing**

Correct Answer: D

Explanation

Engagement drop after AI content adoption often reflects a loss of authentic voice. The differential analysis reveals whether AI volume is substituting for human quality.

Q56. A company rolls out AI tools to all employees. Productivity gains are strong in some teams and absent in others. The most likely structural cause is:

- A. Different task types - AI accelerates execution tasks but adds friction to judgment-heavy tasks**
- B. Some teams are less intelligent
- C. Some teams received better laptops
- D. AI tools were not available in all office locations

Correct Answer: A

Explanation

AI productivity gains are not uniform. Execution-heavy roles see the most acceleration. Roles requiring deep judgment, relationships, or novel problem-solving see less gain.

Q57. A product team uses AI to write sprint tickets. The main risk of this workflow is:

- A. Tickets that are too long
- B. Tickets that are complete and well-formatted but describe the wrong work if the PM brief was imprecise**
- C. Tickets that do not use the correct Jira format
- D. The AI assigning tickets to the wrong engineers

Correct Answer: B

Explanation

AI amplifies the quality of inputs. A vague brief produces polished but misdirected tickets. AI-generated ticket quality is bounded by the PM clarity of thinking.

Q58. A team uses AI to produce meeting summaries with action items. The most common failure mode is:

- A. Summaries that are too long
- B. Summaries that miss the date of the meeting
- C. AI including off-topic comments in the summary
- D. Action items assigned to the wrong owner or with the wrong due date because the AI misidentified who said what**

Correct Answer: D

Explanation

Speaker identification errors cause action item misassignment. Reviewing action item ownership before distributing summaries is the critical human check in this workflow.

Q59. Which AI workflow acceleration practise has the highest compounding return over time?

- A. Building reusable prompt templates for recurring tasks that encode organisational context and standards**
- B. Using AI for as many different tasks as possible each week
- C. Switching to a new AI tool every quarter
- D. Limiting AI use to one task type to build depth

Correct Answer: A

Explanation

Reusable prompts compound because every improvement benefits every future run of that workflow. They also transfer institutional knowledge into repeatable AI behaviour.

Q60. A founder uses AI to run their company for 3 months with no employees. Which function is hardest for AI to fully replace?

- A. Drafting contracts
- B. Writing marketing copy
- C. Relationship-building with customers, partners, and investors**
- D. Generating financial projections

Correct Answer: C

Explanation

AI compresses most execution tasks but cannot build authentic human relationships. Trust, rapport, and influence require human presence - the durable limit of AI automation.

SET

5

AI-First Operations

Set 5 of 5 · 15 Questions · AI-Native Execution

Q61. What distinguishes an AI-first company from a company that uses AI tools?

- A. AI-first companies design their workflows, org structure, and decision-making around AI as the primary execution layer
- B. AI-first companies use more AI tools than others
- C. AI-first companies hire only AI engineers
- D. AI-first companies have no human employees

Correct Answer: A**Explanation**

AI-first is an architectural posture, not a tooling quantity. It means rebuilding how work gets done with AI as the default execution layer, not an add-on.

Q62. A startup is designing its operations from scratch with AI-first principles. What is the first decision to make?

- A. Which AI tools are cheapest
- B. How many AI tools to subscribe to
- C. Whether to use open-source or proprietary AI
- D. Which workflows will remain human-led and why, before designing AI automation for everything else

Correct Answer: D**Explanation**

AI-first design starts with judgment about where humans must lead - then automates everything else. This prevents both under-automation and inappropriate automation.

Q63. A company adopts AI-first operations and finds that middle management roles are changing significantly. The most accurate description of this change is:

- A. Middle managers are fully replaced by AI
- B. Middle managers shift from coordinating information flow to curating AI outputs and making judgment calls**
- C. Middle managers spend more time in meetings
- D. Middle managers become AI engineers

Correct Answer: B

Explanation

AI removes the coordination and information-synthesis tasks that defined middle management. The residual human role is judgment, context, and accountability.

Q64. An AI-first operations team uses AI agents to execute multi-step workflows. The critical human responsibility is:

- A. Monitoring token usage in real time
- B. Ensuring the AI uses the correct output format
- C. Defining the scope, reviewing outputs, and intervening when the agent exceeds its intended remit**
- D. Rotating which AI tool the agent uses each week

Correct Answer: C

Explanation

AI agents can drift or take unintended actions in multi-step workflows. Human oversight defines the boundary of acceptable action and catches scope creep before impact.

Q65. A company runs its customer onboarding entirely on AI. A new enterprise client has a complex, non-standard setup. The correct operational response is:

- A. Escalating to a human CS team immediately rather than forcing the non-standard case through an AI-only flow**
- B. Asking the AI to handle the complex setup anyway
- C. Declining the enterprise client
- D. Redesigning the entire onboarding flow for this one client

Correct Answer: A

Explanation

AI-first operations must have human escalation paths for edge cases. Forcing non-standard inputs through AI-only flows produces poor experiences and churn.

Q66. A company measures operational efficiency as cost per task completed. After deploying AI, cost per task drops 60%. What secondary metric must be tracked?

- A. Number of tasks completed per week
- B. Number of AI tools used
- C. Quality per task, to ensure cost reduction is not coming from lower-quality outputs**
- D. Time the AI takes per task

Correct Answer: C

Explanation

Cost efficiency gains are only real if quality is maintained. AI can reduce cost by reducing quality - the secondary quality metric catches this substitution.

Q67. A company uses AI to handle all internal knowledge queries. Employees stop documenting processes because the AI knows everything. The long-term risk is:

- A. Employees becoming too fast at finding information
- B. The AI becoming too powerful
- C. Employees spending less time searching for information
- D. Institutional knowledge degrading as undocumented processes exist only in conversation history that may not be retained**

Correct Answer: D

Explanation

AI as a knowledge repository does not replace written documentation. Conversation histories are ephemeral and unsearchable. Documentation degrades silently in an AI-only model.

Q68. An AI-first company wants to hire its first employee. The role that generates the highest marginal value is:

- A. A data entry specialist
- B. Someone who can direct and improve AI systems - a judgment and orchestration role, not a pure execution role**
- C. A social media manager
- D. A generalist who can do everything manually

Correct Answer: B

Explanation

In an AI-first company, execution is already handled. The first hire marginal value is in improving AI direction, quality, and strategy - not in adding execution capacity.

Q69. A company runs all its marketing on AI-generated content for 6 months. Organic reach declines. The most likely cause is:

- A. The AI is generating too much content
- B. The AI is using outdated social media algorithms
- C. AI-generated content lacks the specificity, authenticity, and novelty that audiences respond to over time**
- D. Competitors have larger budgets

Correct Answer: C

Explanation

AI content often converges on averaged, generic outputs. Audiences habituate to AI patterns quickly. Human editorial distinctiveness is the differentiator AI cannot replicate at scale.

Q70. A PM wants to build an AI-first onboarding workflow. The correct starting point is:

- A. Mapping the current manual onboarding steps and identifying which are highest-volume and most rule-based**
- B. Choosing which AI tool to use first
- C. Building the AI onboarding before the manual version exists
- D. Outsourcing onboarding to a third-party AI vendor immediately

Correct Answer: A

Explanation

AI automation should follow a documented manual process. Without mapping the current workflow, AI onboarding automates the wrong steps or misses critical ones.

Q71. An AI-first company tracks AI coverage - the percentage of workflow steps handled by AI. At 90% AI coverage, the most important human focus is:

- A. Increasing AI coverage to 100%
- B. The 10% of steps that remain human - which are there because they require judgment, relationships, or accountability**
- C. Reducing the number of workflow steps
- D. Documenting the AI coverage metric for investors

Correct Answer: B

Explanation

The residual human steps in an AI-first operation are not inefficiencies - they are the highest-value human contributions. They deserve the most human attention, not elimination.

Q72. A company uses AI to draft all legal agreements. The non-negotiable human role in this workflow is:

- A. Having the AI verify its own draft
- B. Using AI to negotiate the terms after drafting
- C. Sending AI drafts directly to counterparties without legal review
- D. A qualified lawyer reviewing every agreement before execution, regardless of how good the AI draft appears**

Correct Answer: D

Explanation

Legal agreements carry binding obligations. AI drafts can contain errors with significant legal consequences. Qualified human review is a non-negotiable control regardless of AI quality.

Q73. An AI-first ops team is building a new internal tool. The correct build sequence is:

- A. Choose an AI tool, give it instructions, then launch
- B. Hire an engineer, write code, then add AI later
- C. Define the workflow in plain language, then identify AI automation points, then build the human exception path, then deploy**
- D. Launch with full AI automation, then add human steps if problems emerge

Correct Answer: C

Explanation

AI-first tooling requires human exception paths to be designed in, not bolted on. Defining the workflow first ensures AI is applied to the right steps.

Q74. A company has replaced most of its customer support with AI. Customer satisfaction scores are up but escalation volume is also up. What does this signal?

- A. The escalation team is too small
- B. AI handles routine queries well but is creating friction on complex or sensitive issues that quickly exhaust users**
- C. Customers prefer to speak to humans regardless of AI quality
- D. The AI is answering too quickly

Correct Answer: B

Explanation

Rising escalation alongside rising CSAT means AI is succeeding on easy queries but failing on hard ones. The fix is better routing, not more AI in complex cases.

Q75. A founder builds a company that reaches 1000 customers with a team of 3 people using AI. The most accurate description of this business model is:

- A. An understaffed company that will collapse at scale
- B. A company that has replaced all human judgment with AI
- C. A company that uses AI for marketing only
- D. AI-leveraged operations where automation handles a volume of work that would traditionally require 20 to 50 headcount**

Correct Answer: D

Explanation

AI-native companies operate with headcount leverage ratios that were previously impossible. Automation handles execution volume; a small human team handles judgment and direction.

END OF QUESTION BANK

*Your turning
point starts
here.*