menler

Your turning point in the AI era.

AI APTITUDE QUESTION BANK

AI Product Thinking

About This Question Bank

These questions cover the skills required to build, prioritise, and operate AI-powered products: UX trust design, product decision-making, feature prioritisation, workflow integration, and measuring operational value.

5	15	75
Sets	Questions per Set	Total Questions

SET GUIDE

Set 1 · AI UX and Trust
AI UX and trust — explainability, agency transparency, appropriate reliance

Set 2 · Product Decision-Making
Product decision-making — build vs buy, failure mode planning, model versioning

Set 3 · Feature Prioritisation
Feature prioritisation — JTBD, binding constraints, kill criteria, retention signals

Set 4 · User Workflows
User workflows — AI insertion points, delegation design, feedback loops

Set 5 · Operational Usefulness
Operational usefulness — ROI measurement, escalation design, governance

Each question has one correct answer. Select the best answer from the options provided.

Correct answers and explanations follow each question.

SET
1

AI UX and Trust

Set 1 of 5 · 15 Questions · AI Product Thinking

Q1. A product team ships an AI feature that gives confidently worded answers but is wrong 20% of the time. What UX pattern best mitigates trust erosion?

A. Remove the feature until accuracy reaches 100%

B. Show confidence scores or hedged language so users calibrate trust appropriately

C. Add a disclaimer in the terms of service

D. Increase font size of outputs to appear more authoritative

Correct Answer: B

Explanation

Users calibrate trust based on cues. Explicit uncertainty signals prevent over-reliance and protect the product when errors occur.

Q2. Which design principle best describes showing users why an AI made a specific recommendation?

A. Latency optimisation

B. A/B testing

C. Feature gating

D. Explainability

Correct Answer: D

Explanation

Explainability surfaces the reasoning behind AI outputs, enabling users to verify, trust, or correct the system.

Q3. A user completes a task using an AI assistant but cannot remember if the AI did it or they did. This is a failure of:

A. Agency transparency

B. Latency management

C. Prompt injection prevention

D. Output formatting

Correct Answer: A

Explanation

Agency transparency means the product clearly communicates when AI acted on behalf of the user vs when the user acted themselves.

Q4. What is the primary purpose of an undo or revert affordance in an AI-powered product?

A. Reducing API costs

B. Improving model accuracy

C. Restoring user control after AI-initiated actions

D. Satisfying GDPR requirements

Correct Answer: C

Explanation

AI actions can be irreversible or unexpected. An undo mechanism preserves user agency and reduces fear of delegation.

Q5. A language learning app uses AI to grade essays. Students report feeling judged rather than helped. The most effective UX fix is:

A. Remove AI grading entirely

B. Frame AI feedback as suggestions with a growth-oriented tone

C. Switch to a different model

D. Add more correction categories

Correct Answer: B

Explanation

Tone and framing shape how users receive AI feedback. Coaching language reduces defensiveness and increases engagement.

Q6. When is progressive disclosure most valuable in an AI product interface?

A. When AI outputs are complex and users may only need a summary initially

B. When the AI has low latency

C. When the model is fine-tuned

D. When the product is in beta

Correct Answer: A

Explanation

Progressive disclosure prevents overwhelm by surfacing the most relevant information first, with detail available on demand.

Q7. A PM notices users frequently edit AI-generated content before using it. This behavior most likely indicates:

A. The model needs to be replaced

B. Users do not trust AI in principle

C. The feature should be removed

D. The AI output is directionally useful but not precise enough for direct use

Correct Answer: D

Explanation

High edit rates signal partial value. Users find the output worth refining rather than discarding, pointing to a gap in specificity or context.

Q8. What does appropriate reliance mean in the context of AI-assisted decision-making?

A. Users always follow AI recommendations

B. Users ignore AI and make independent decisions

C. Users trust AI when it is correct and override it when it is wrong

D. Users rate AI outputs before acting on them

Correct Answer: C

Explanation

Appropriate reliance is the calibrated use of AI - neither over-reliance (automation bias) nor under-reliance (ignoring correct suggestions).

Q9. An AI product shows users a generated by AI label on all outputs. The main benefit of this label is:

A. Setting accurate expectations and enabling informed verification

B. Increasing click-through rates

C. Reducing model inference costs

D. Satisfying animation guidelines

Correct Answer: A

Explanation

Transparency labels help users know what to verify and adjust their scrutiny appropriately, reducing over-trust.

Q10. Which user segment is most likely to experience automation bias when using an AI product?

A. Expert users with deep domain knowledge

B. Users who dislike technology

C. Users who frequently edit AI outputs

D. Users under time pressure with limited domain expertise

Correct Answer: D

Explanation

Automation bias is strongest when users lack the knowledge or time to critically evaluate AI outputs - they default to accepting suggestions.

Q11. A product team wants to measure user trust in their AI feature. The most direct behavioral metric is:

A. Daily active users

B. Rate at which users act on AI suggestions without modification

C. Time to first interaction

D. Net Promoter Score

Correct Answer: B

Explanation

Unmodified acceptance rate is a behavioral proxy for trust - it captures actual reliance rather than stated preference.

Q12. What is the risk of making AI features invisible or seamlessly integrated with no indication to users?

A. Latency increases significantly

B. Model fine-tuning becomes more expensive

C. Users cannot calibrate trust or know when to verify outputs

D. API rate limits are more likely to be hit

Correct Answer: C

Explanation

Invisible AI removes users from the verification loop. When errors occur they are less likely to be caught.

Q13. A fintech product uses AI to flag suspicious transactions. False positives frustrate users. The best UX approach is:

A. Remove the flagging feature

B. Increase the flagging threshold to reduce all flags

C. Add more flags to balance the false positives

D. Show the reason for the flag and provide a clear dispute path

Correct Answer: D

Explanation

Explainability plus recourse transforms a frustrating AI error into a manageable interaction. Users can act on and resolve the issue.

Q14. A product manager wants to build user confidence in a new AI writing assistant. The most effective early onboarding move is:

A. Displaying model benchmarks on the landing page

B. Showing users a before and after example of a task the AI improved

C. Requiring users to complete a tutorial before using the product

D. Adding a disclaimer about AI limitations

Correct Answer: B

Explanation

Concrete demonstrations of value build trust faster than abstract claims. Before and after examples make the benefit tangible.

Q15. Which of the following is a signal that users have developed appropriate mental models of an AI product?

A. Users choose when to use the AI feature and when to handle tasks manually

B. Users always use the AI feature for every task

C. Users never override AI suggestions

D. Users spend more time on the product

Correct Answer: A

Explanation

Selective, deliberate use shows users understand the AI strengths and limits - the hallmark of an accurate mental model.

SET
2

Product Decision-Making

Set 2 of 5 · 15 Questions · AI Product Thinking

Q16. A PM must decide whether to build a custom model or use a third-party LLM API. The most critical factor in this decision is:

A. Whether the CEO prefers open-source or proprietary models

B. Whether the team has Python skills

C. Whether the required capability gap justifies the cost and maintenance overhead of custom development

D. Whether competitors are using custom models

Correct Answer: C

Explanation

Build vs buy for models is a capability, cost, and maintenance decision. The threshold is whether existing APIs can meet the requirement acceptably.

Q17. When evaluating AI features for a product roadmap, which framework best captures the unique risks of AI?

A. Impact-effort matrix weighted for failure mode severity

B. Story point estimation

C. Velocity tracking

D. RICE scoring without modification

Correct Answer: A

Explanation

Standard prioritisation frameworks do not account for AI-specific risks like hallucination or bias. Weighting for failure severity makes them appropriate.

Q18. A product team receives a request to add a generative AI chatbot to a healthcare app. The first decision gate should be:

A. Choosing between GPT-4 and Claude

B. Designing the chat UI

C. Estimating development time

D. Assessing regulatory requirements and liability for AI-generated health information

Correct Answer: D

Explanation

In healthcare, regulatory compliance and liability are pre-conditions for any AI feature. Technical choices come after viability is established.

Q19. A PM sees that an AI feature has high usage but low satisfaction scores. The most productive next step is:

A. Increasing inference speed

B. Qualitative research to understand the gap between usage intent and output quality

C. Switching the underlying model

D. Adding more features to the same surface

Correct Answer: B

Explanation

High usage with low satisfaction means users are trying the feature but not getting what they need. Qualitative research reveals the specific unmet need.

Q20. Which metric best signals that an AI feature is creating genuine product value rather than just generating engagement?

A. Session duration

B. Number of AI queries per session

C. Task completion rate for outcomes users could not achieve before the feature

D. Feature discovery rate

Correct Answer: C

Explanation

Real value is measured by whether the AI enables outcomes that were previously impossible or impractical for users - not just activity metrics.

Q21. A PM is deciding whether to launch an AI feature in beta or full release. The strongest argument for beta is:

A. Beta features are free to build

B. Beta reduces API costs

C. Full release requires board approval

D. The failure mode of AI errors in production requires real user feedback before full exposure

Correct Answer: D

Explanation

AI features can fail in unexpected ways that are hard to anticipate in staging. Beta allows discovery and remediation before wide impact.

Q22. A startup wants to use AI to automate customer support. The most important pre-launch question is:

A. What happens when the AI gives an incorrect or harmful response to a customer?

B. Which LLM has the lowest per-token cost?

C. How fast can we ship the chatbot?

D. Should the chatbot have a name?

Correct Answer: A

Explanation

Failure mode planning is the first question for customer-facing AI. The answer determines guardrails, escalation paths, and liability exposure.

Q23. A product decision to remove human review from an AI-powered loan decisioning system should be gated on:

A. Speed improvement metrics alone

B. Demonstrated accuracy across demographic segments and regulatory clearance

C. Cost savings projections

D. Model benchmark scores on public datasets

Correct Answer: B

Explanation

Removing human oversight in high-stakes domains requires fairness validation and regulatory sign-off - not just aggregate accuracy.

Q24. Which of the following is the strongest signal that a product team has underinvested in AI product quality?

A. The team uses a third-party API

B. The product has a free tier

C. The team ships weekly

D. Users frequently screenshot AI outputs and manually share them to get corrections

Correct Answer: D

Explanation

Manual workarounds to fix AI errors are a strong signal that the error rate or recourse path inside the product is inadequate.

Q25. When should a PM introduce model versioning as a product requirement?

A. Only when enterprise customers request it

B. After reaching 1 million users

C. When the product core experience depends on consistent AI behaviour across releases

D. When switching from one provider to another

Correct Answer: C

Explanation

Model updates can silently change product behaviour. Versioning is a product requirement whenever consistency is part of the value proposition.

Q26. A team ships an AI feature and notices a drop in retention among a specific user cohort. The most likely AI-specific cause is:

A. The feature is too slow

B. The feature produces outputs that do not match that cohort context, language, or needs

C. The feature is too cheap

D. The feature was not announced in release notes

Correct Answer: B

Explanation

AI products often underperform for specific cohorts because training data or prompts do not represent their context. Cohort analysis reveals this gap.

Q27. A PM wants to reduce AI hallucinations without waiting for model improvements. The best product-level intervention is:

A. Constraining the AI to only answer within a verified knowledge base and surfacing citations

B. Increasing the temperature parameter

C. Adding more tokens to the context window

D. Switching to a larger model

Correct Answer: A

Explanation

RAG with citation is the primary product-level tool for reducing hallucination impact - it grounds outputs in verified sources users can check.

Q28. A product team is debating whether to show AI-generated content inline or in a separate review panel. The main variable determining the right answer is:

A. Which design pattern is more common in the market

B. The team design preferences

C. How consequential errors would be if users act on them without reviewing

D. The number of tokens in each output

Correct Answer: C

Explanation

High-stakes outputs warrant friction and separation to encourage review. Low-stakes outputs benefit from inline delivery for speed and flow.

Q29. Which of the following is a product-level solution to AI bias in a hiring tool?

A. Using a larger model

B. Adding more training data

C. Reducing the temperature

D. Auditing output distributions across candidate demographic groups before launch

Correct Answer: D

Explanation

Bias audits at the product level catch discriminatory patterns before they affect real candidates - regardless of which model is used.

Q30. A PM must decide between two AI features: one that delights 10% of users intensely and one that mildly improves the experience for 80%. The decision framework should prioritise:

A. The feature with the larger user count always

B. Strategic fit with the product core value proposition and retention impact

C. The feature with the higher NPS lift always

D. The feature that is cheaper to build

Correct Answer: B

Explanation

Neither intensity nor breadth is universally correct. The right choice depends on which feature advances core product strategy and drives retention.

SET
3

Feature Prioritisation

Set 3 of 5 · 15 Questions · AI Product Thinking

Q31. A product team has three AI feature requests: faster inference, better accuracy, and an explainability panel. How should they prioritise?

A. By mapping each to the top user pain point identified in research

B. Accuracy first, always

C. Ship all three simultaneously

D. Start with explainability because it is cheapest

Correct Answer: A

Explanation

Prioritisation without user pain mapping produces technically impressive but adoption-poor features. Research grounds the decision.

Q32. Which prioritisation signal most directly indicates that an AI feature should be deprioritised?

A. The feature takes more than two sprints to build

B. The feature requires a new model

C. Users who discover the feature do not return to use it a second time

D. The feature has no competitor equivalent

Correct Answer: C

Explanation

Zero repeat usage after discovery signals that the feature solved no real problem or delivered no perceived value - the clearest deprioritisation signal.

Q33. A PM has limited engineering capacity. An AI feature is requested by the top 5 enterprise customers but not by any self-serve users. The best response is:

A. Immediately prioritise it because enterprise revenue is higher

B. Evaluate whether the enterprise use case is a leading indicator of broader need

C. Ignore it because it has no self-serve demand

D. Ask the customers to pay for custom development

Correct Answer: B

Explanation

Enterprise requests often preview emerging needs. The PM should assess whether this is niche or a signal of wider unmet demand before committing resources.

Q34. A product team is using a jobs-to-be-done framework to evaluate AI features. What does this approach prioritise?

A. Technical capability benchmarks

B. The frequency with which a feature is used

C. Revenue per feature

D. The functional progress users are trying to make, not the features themselves

Correct Answer: D

Explanation

Jobs-to-be-done centres on user motivation. AI features that address real jobs get adopted; features built around capability showcasing do not.

Q35. When should AI personalisation be prioritised over a general-purpose AI feature?

A. When there is strong evidence that user context meaningfully changes the optimal output

B. Always, because personalisation is always better

C. Only for enterprise tiers

D. When the team has idle ML engineers

Correct Answer: A

Explanation

Personalisation adds complexity and maintenance. It is worth building only when user-specific context provably improves outcomes.

Q36. A team is building an AI product. They can invest in either model accuracy improvements or UI improvements for the same cost. Research shows users blame the UI when they distrust outputs. What should they prioritise?

A. Model accuracy, because ground truth matters more

B. UI improvements, because trust perception is currently the binding constraint

C. Neither, wait for user feedback

D. Build both in parallel

Correct Answer: B

Explanation

The binding constraint is what limits adoption. If trust perception is the bottleneck, UI investment yields higher returns than model accuracy.

Q37. A product manager receives conflicting feature requests from power users and new users for the same AI surface. The right approach is:

A. Prioritise power users because they generate more revenue

B. Prioritise new users because they represent growth

C. Build two separate products

D. Design for progressive complexity - simple defaults with power user controls accessible on demand

Correct Answer: D

Explanation

Progressive complexity resolves the power user vs new user tension by making the product approachable without limiting advanced usage.

Q38. Which type of AI feature should be given the highest priority for early kill decision-making?

A. Features with low usage rates

B. Features that are technically complex to maintain

C. Features that produce harmful or misleading outputs even in edge cases

D. Features that exceed budget

Correct Answer: C

Explanation

Harmful outputs in edge cases represent liability and trust damage that can outweigh the value of the entire feature. Kill criteria must include safety thresholds.

Q39. A startup AI product has 10 feature ideas and 3 months of runway. The correct prioritisation lens is:

A. Which feature users mentioned most in surveys

B. Which single feature most increases the probability of survival through retention or revenue

C. Which feature is easiest to build

D. Which feature the founder finds most interesting

Correct Answer: B

Explanation

Under resource constraint, survival prioritisation means concentrating on the one feature with the highest leverage on key business metrics.

Q40. A PM is evaluating two AI features: one with high ceiling impact but high variance, and one with moderate but reliable impact. Which is preferable in a mature product with a large user base?

A. The reliable impact feature, because variance at scale creates unpredictable user experience

B. The high ceiling feature, because ambition signals product leadership

C. The feature that ships first

D. The feature preferred by the design team

Correct Answer: A

Explanation

At scale, variance in AI output quality creates inconsistent experiences across millions of users. Reliability compounds into retention; ceiling impact rarely does.

Q41. A PM must decide when to sunset an AI feature. The strongest signal is:

A. The feature has been live for more than one year

B. A competitor has shipped a similar feature

C. Maintenance cost exceeds the value delivered relative to alternative uses of the same engineering time

D. The original PM who built it has left

Correct Answer: C

Explanation

Feature sunset is an opportunity cost decision. A feature should be removed when it consumes more resources than the value it generates compared to alternatives.

Q42. A team has data showing that 90% of users only use 20% of the AI features. What should the roadmap prioritise?

A. Marketing the unused 80% to increase adoption

B. Removing the 20% and rebuilding with new features

C. Adding even more features to increase surface area

D. Deepening the 20% of features most used, not expanding the unused 80%

Correct Answer: D

Explanation

Usage concentration is a signal to invest in depth, not breadth. The used features represent the product actual value - deepen them.

Q43. Which of the following is the best leading indicator for an AI feature long-term retention impact?

A. Feature launch day traffic

B. Whether users who adopt the feature have higher 30-day retention than those who do not

C. Number of press mentions at launch

D. Internal team excitement about the feature

Correct Answer: B

Explanation

Differential retention between adopters and non-adopters is the strongest early signal that a feature creates durable value.

Q44. A product team is prioritising between an AI feature that improves the core workflow and one that is novel but peripheral. The correct default is:

A. Novel feature, because it generates press coverage

B. Peripheral feature, because it differentiates from competitors

C. Core workflow improvement, because retention impact is higher for features on the critical path

D. Whichever is cheaper to build

Correct Answer: C

Explanation

Features on the critical path have disproportionate retention impact. Peripheral novelty rarely compounds into retention the way core workflow improvements do.

Q45. A PM receives a request to add an AI feature that would require collecting additional user data. The first question is:

A. Whether the value the feature delivers justifies the privacy cost and user consent requirement

B. Whether the data is available to purchase

C. Whether the feature has a competitor equivalent

D. Whether the model supports the data type

Correct Answer: A

Explanation

Privacy cost must be weighed against feature value before any data collection decision. Collecting data without this analysis creates regulatory and trust risk.

SET
4

User Workflows

Set 4 of 5 · 15 Questions · AI Product Thinking

Q46. A PM is mapping user workflows to identify where AI can add the most value. The highest-value AI insertion point is typically:

A. Steps that users enjoy the most

B. Steps that happen at the end of the workflow

C. Steps that require no data

D. Steps that are high-frequency, cognitively demanding, and have predictable structure

Correct Answer: D

Explanation

AI delivers the most leverage on tasks that repeat often, drain cognitive energy, and have patterns the model can learn - these are the best automation targets.

Q47. A user research session reveals that users complete a workflow faster with AI but feel less ownership over the output. This is a signal to:

A. Remove the AI feature

B. Redesign the AI interaction to give users meaningful editorial control at key steps

C. Add more AI automation to reduce effort further

D. Change the underlying model

Correct Answer: B

Explanation

Reduced ownership erodes satisfaction and retention even when speed improves. Meaningful control restores agency without sacrificing efficiency.

Q48. Which workflow insertion pattern is most appropriate when AI errors are costly and hard to reverse?

A. Full automation with post-hoc audit

B. AI acting autonomously and logging decisions

C. AI as a drafting assistant with mandatory human review before any action is taken

D. Removing AI from the workflow until errors reach zero

Correct Answer: C

Explanation

High-stakes, hard-to-reverse workflows require human-in-the-loop designs. AI drafts reduce effort while mandatory review preserves oversight.

Q49. A product team notices that users skip the AI suggestion step in a workflow and complete the task manually. The most likely reason is:

A. The AI suggestion is too generic and does not save meaningful time for this workflow

B. Users prefer typing

C. The AI is too slow

D. The feature is not visible enough

Correct Answer: A

Explanation

Skip behavior reveals that the AI output is not sufficiently useful relative to the effort of reviewing it. Context specificity is usually the gap.

Q50. In a document editing product, AI suggestions are shown inline but users report feeling distracted. The best UX fix is:

A. Increasing the font size of suggestions

B. Making suggestions appear faster

C. Requiring users to accept all suggestions before proceeding

D. Moving AI suggestions to a sidebar or making them opt-in rather than always-on

Correct Answer: D

Explanation

Inline AI suggestions can interrupt flow for experienced users. Sidebar or opt-in delivery respects user attention and workflow rhythm.

Q51. A PM is designing an AI onboarding flow. The most effective early moment is:

A. Showing a feature list video

B. Asking users to rate the AI before using it

C. Having the AI complete a real task for the user using their actual data

D. Requiring users to complete a quiz about AI

Correct Answer: C

Explanation

First-hand value experience with real data converts more users than demonstrations. The AI should prove itself immediately in the user actual context.

Q52. A workflow automation product allows AI to take multi-step actions on behalf of users. The most important design requirement is:

A. A clear audit trail showing every action the AI took and why

B. Minimising the number of steps shown to the user

C. Hiding AI decision points to avoid confusion

D. Allowing the AI to modify its own instructions

Correct Answer: A

Explanation

Multi-step AI actions require full traceability. Users need to see and verify what was done before trusting the system with consequential tasks.

Q53. When designing an AI-assisted research workflow, which principle most improves output quality?

A. Removing user control to reduce decision fatigue

B. Giving users control over what sources the AI searches and how results are filtered

C. Maximising the number of results returned

D. Hiding source information to simplify the display

Correct Answer: B

Explanation

Research quality depends on source quality and relevance. User control over scope produces more targeted, trustworthy results.

Q54. A team notices that AI-generated summaries are used directly without review in 70% of cases. This observation most warrants:

A. Improving the summary model

B. Reducing the length of summaries

C. Adding inline quality signals and encouraging review for high-stakes summaries

D. Removing the human review option

Correct Answer: C

Explanation

70% unreviewed acceptance is high automation bias. Quality signals and review nudges help users calibrate without making review mandatory.

Q55. A support tool uses AI to draft responses for agents. Agents report copying AI drafts verbatim without reading them. The product risk is:

A. Agents becoming too fast

B. Support costs dropping too quickly

C. AI drafts becoming longer over time

D. Harmful or incorrect responses reaching customers without human verification

Correct Answer: D

Explanation

Verbatim copy-paste without review defeats the human-in-the-loop design. The product must create meaningful friction that prompts at least scanning.

Q56. A PM is designing a feedback mechanism for AI workflow outputs. The most useful feedback type for model improvement is:

A. Specific inline edits made by users, captured as diff data

B. Star ratings after task completion

C. Binary thumbs up or down without context

D. NPS surveys sent monthly

Correct Answer: A

Explanation

Inline edit diffs reveal exactly what the AI got wrong and how users corrected it - the richest training signal for improving model outputs.

Q57. A product manager is evaluating whether to add AI to a simple, well-understood workflow. The key question is:

A. Whether the team has the AI budget

B. Whether AI reduces friction enough to justify the added complexity and failure risk

C. Whether competitors have AI in this workflow

D. Whether the design team prefers AI interfaces

Correct Answer: B

Explanation

Not every workflow benefits from AI. The bar is whether the friction reduction exceeds the complexity and failure modes introduced.

Q58. A product allows users to delegate tasks to an AI agent. Users hesitate to delegate high-stakes tasks. The best product response is:

A. Increasing the AI agent autonomy

B. Marketing the agent accuracy more aggressively

C. Restricting agent use to low-stakes tasks only

D. Providing transparency into agent actions and an easy recall or cancel mechanism

Correct Answer: D

Explanation

Hesitation around delegation signals insufficient trust. Transparency into what the agent is doing and the ability to stop it are the key trust builders.

Q59. A workflow product tracks that AI-assisted tasks take 40% less time but user satisfaction is flat. The most likely explanation is:

A. Speed improvement alone does not address the quality or confidence gap users experience with AI outputs

B. Users prefer slower workflows

C. The feature has not been marketed enough

D. The AI is using the wrong model

Correct Answer: A

Explanation

Time savings are necessary but not sufficient. If users are not satisfied with the output quality or feel uncertain about results, satisfaction stays flat.

Q60. A PM wants to design an AI feature that helps users make better decisions rather than making decisions for them. The correct design pattern is:

A. Full automation - letting the AI decide and notifying the user

B. Removing the decision step entirely from the workflow

C. Decision support - surfacing relevant information, options, and tradeoffs without prescribing a choice

D. Adding more data fields for users to fill in before the AI acts

Correct Answer: C

Explanation

Decision support augments human judgment. It preserves agency while reducing cognitive load - appropriate when users need to own the outcome.

SET

5

Operational Usefulness

Set 5 of 5 · 15 Questions · AI Product Thinking

Q61. A PM wants to measure the ROI of an AI feature for an internal operations team. The most direct metric is:

A. Time saved per task multiplied by the number of tasks and loaded staff cost

B. Number of AI requests made per day

C. Model accuracy on internal benchmarks

D. Number of features shipped using AI

Correct Answer: A

Explanation

Operational ROI is time x volume x cost. This gives a financial figure the business can compare directly against the cost of building and running the feature.

Q62. An operations team adopts an AI tool but usage drops after week 2. The most likely cause is:

A. The tool is too expensive

B. The team lacks technical skills

C. The model version changed

D. Initial novelty wore off and the tool does not save enough time to justify the behaviour change

Correct Answer: D

Explanation

Week-2 drop-off is a classic signal of novelty-driven initial adoption without a habit-forming value proposition. The core workflow benefit is insufficient.

Q63. Which operational AI use case has the highest value threshold for accuracy?

A. Email subject line generation

B. Medical diagnosis support, where errors directly affect patient outcomes

C. Blog post summarisation

D. Social media caption writing

Correct Answer: B

Explanation

Accuracy thresholds must be proportional to consequence. Medical decisions have life-altering stakes, making them the highest bar for AI accuracy requirements.

Q64. A team deploys an AI tool for internal knowledge retrieval. After launch, employees report using it less than the old search system. The most productive diagnostic question is:

A. Is the AI model large enough?

B. Was there sufficient marketing for the new tool?

C. Are users finding what they need faster with the old system, and why does the AI fail to match that?

D. Is the system deployed in the cloud?

Correct Answer: C

Explanation

Reverting to the old tool is a direct comparison signal. Understanding why the AI fails on the same tasks is the most useful diagnostic for improvement.

Q65. An AI coding assistant is deployed to a 50-person engineering team. What is the most important operational metric to track after 30 days?

A. Change in pull request review cycles and defect rates, not just lines of code generated

B. Number of AI code suggestions accepted

C. Total tokens consumed by the team

D. How often engineers open the tool

Correct Answer: A

Explanation

Code quality metrics reveal whether AI assistance improves actual engineering outcomes. Acceptance rate and usage are activity metrics, not outcome metrics.

Q66. A PM is building the business case for an internal AI tool. Stakeholders ask for the payback period. What inputs are required?

A. Model benchmark scores and API latency

B. Number of AI features available in the tool

C. Total implementation cost, ongoing API cost, and measurable time saved per user per week

D. Competitor adoption rates

Correct Answer: C

Explanation

Payback period equals total cost divided by value generated per period. Time saved per user converted to labour cost is the primary value numerator for internal tools.

Q67. An internal AI assistant is deployed across departments with different workflows. Adoption is high in marketing but low in finance. The most likely reason is:

A. Finance employees are less technically sophisticated

B. Marketing has a larger team

C. Finance did not receive the announcement email

D. The AI outputs are better calibrated for unstructured creative tasks than for the precise, regulated outputs finance requires

Correct Answer: D

Explanation

LLMs perform better on open-ended creative tasks than on structured, accuracy-critical outputs like financial models. Adoption reflects fit to task type.

Q68. A company deploys AI to handle tier-1 customer support. The correct escalation design is:

A. AI handles all queries including complaints and legal threats

B. AI handles routine queries and escalates to humans when confidence is low or the issue is sensitive

C. AI only handles queries when human agents are unavailable

D. Humans handle all queries and AI only logs them

Correct Answer: B

Explanation

Escalation design is the critical safety mechanism for customer-facing AI. Low confidence and sensitive topics are the two primary triggers for human handoff.

Q69. A product manager wants to track whether an AI tool improves employee productivity without creating new risks. Which metric pair is most complete?

A. Usage frequency plus session length

B. Tokens consumed plus model version

C. Output quality maintained plus time per task reduced

D. Number of prompts submitted plus acceptance rate

Correct Answer: C

Explanation

Productivity gain without quality maintenance is false gain. The correct measurement pair captures both the efficiency improvement and the quality safeguard.

Q70. A PM launches an AI feature that automates a manual data entry workflow. Six months later the team reports the manual workflow has atrophied and no one remembers how to do it. This is a risk of:

A. Over-automation creating single points of failure and skill degradation

B. Successful adoption

C. High user satisfaction

D. Strong retention metrics

Correct Answer: A

Explanation

Skill atrophy through automation is an operational risk. When AI fails, the fallback capability no longer exists - a dependency that must be planned for.

Q71. A PM is evaluating three AI vendors for an internal HR tool. Beyond accuracy and cost, the most important evaluation criterion is:

A. The vendor market share

B. Data privacy terms, including who owns training rights over submitted employee data

C. Whether the vendor offers a free trial

D. Whether the vendor has a mobile app

Correct Answer: B

Explanation

Employee data submitted to AI tools can be used for model training unless explicitly prohibited. Privacy terms are a compliance and trust-critical evaluation factor.

Q72. A PM notices that AI-generated reports are used in board presentations without any human editing. The correct product response is:

A. Improving model quality to match human-level writing

B. Disabling the export function

C. Removing the report generation feature

D. Adding a mandatory review checklist and clearly marking content as AI-generated before export

Correct Answer: D

Explanation

Board-level documents with unreviewed AI content carry reputational and accuracy risk. Review friction and transparency labels are appropriate product guardrails.

Q73. Which of the following best describes the operational risk of prompt injection in an enterprise AI tool?

A. The model runs out of context window during a long prompt

B. Users ask the AI questions outside its intended domain

C. Malicious input in user-submitted data causes the AI to take unauthorised actions or leak information

D. The AI produces outputs that are too long for the UI to display

Correct Answer: C

Explanation

Prompt injection allows attackers to hijack AI behaviour through content the model processes as instructions - a critical enterprise security risk.

Q74. An internal AI tool is being used by employees to generate external-facing content. The biggest operational governance risk is:

A. Employees spending too much time using the tool

B. AI-generated content being published without review, creating accuracy or brand-voice issues

C. The tool generating content faster than the marketing team can review

D. The tool learning from employee writing style

Correct Answer: B

Explanation

External content represents the organisation publicly. AI-generated content published without review is an accuracy, brand, and legal risk that requires governance.

Q75. A PM is writing success criteria for an AI operational tool six months post-launch. Which set of criteria is most complete?

A. Adoption rate plus number of features shipped

B. API uptime plus token cost within budget

C. Number of support tickets raised about the tool

D. Adoption rate plus measurable time saved plus error rate maintained below threshold plus positive user feedback

Correct Answer: D

Explanation

Complete success criteria for an operational AI tool include adoption, efficiency, quality, and satisfaction - all four pillars.

END OF QUESTION BANK

Your turning point starts here.